

Computer Systems II

Li Lu

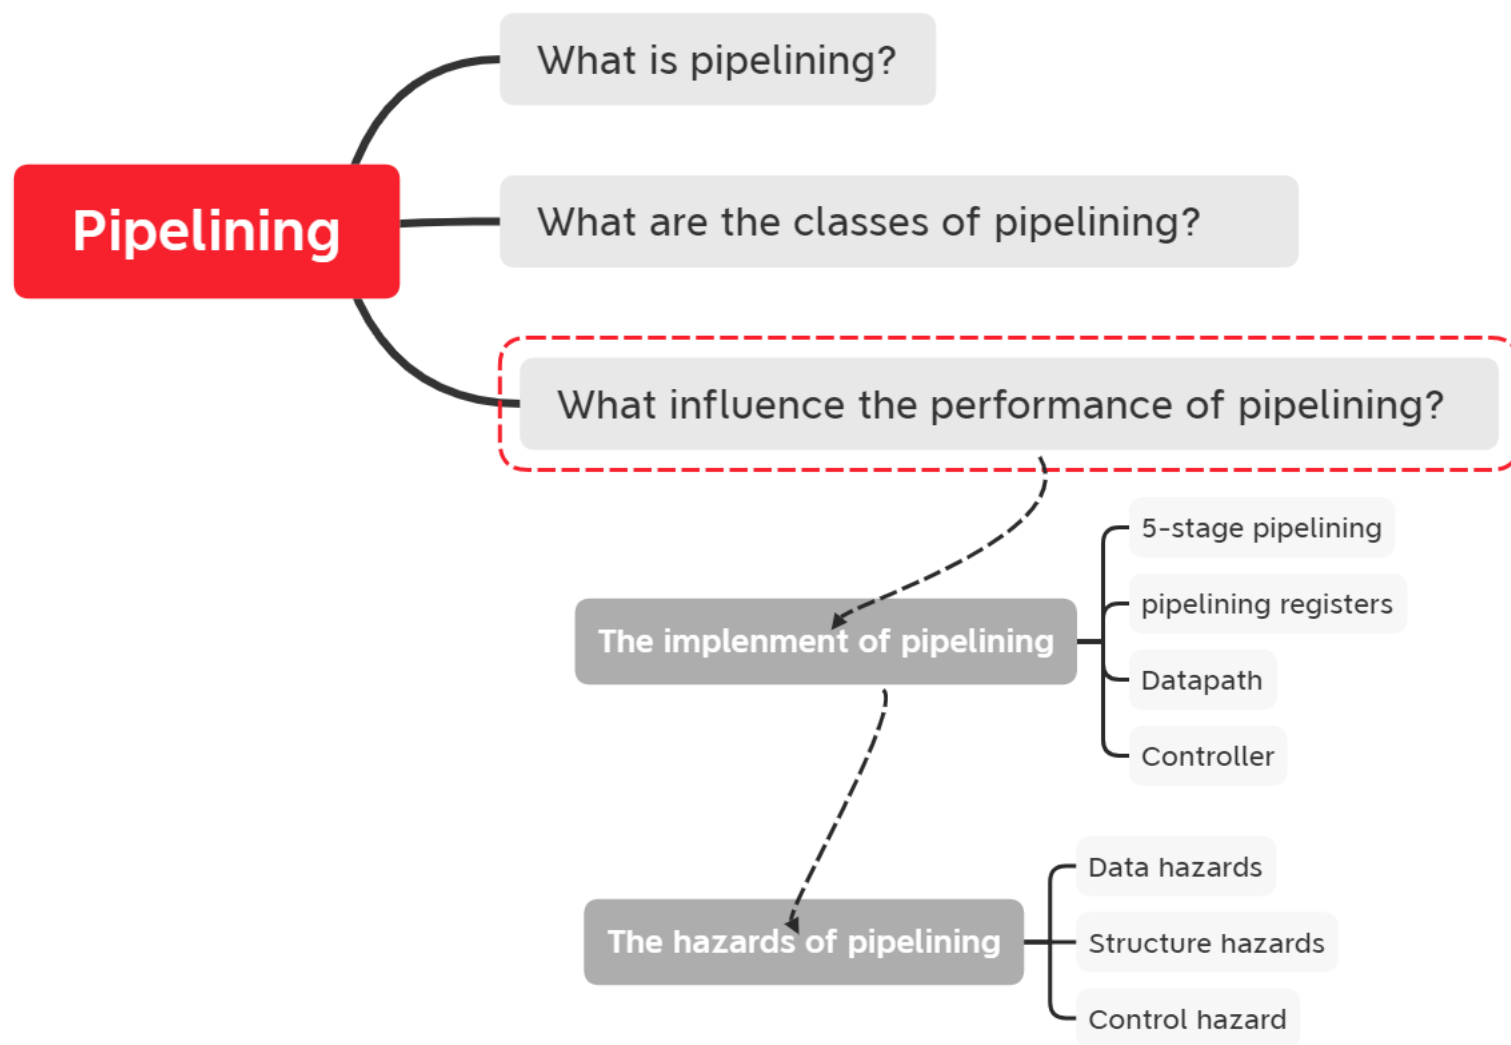
Room 319, Yifu Business and Management Building

Yuquan Campus

li.lu@zju.edu.cn

<https://person.zju.edu.cn/lynnluli>





How Pipelining Improves Performance?

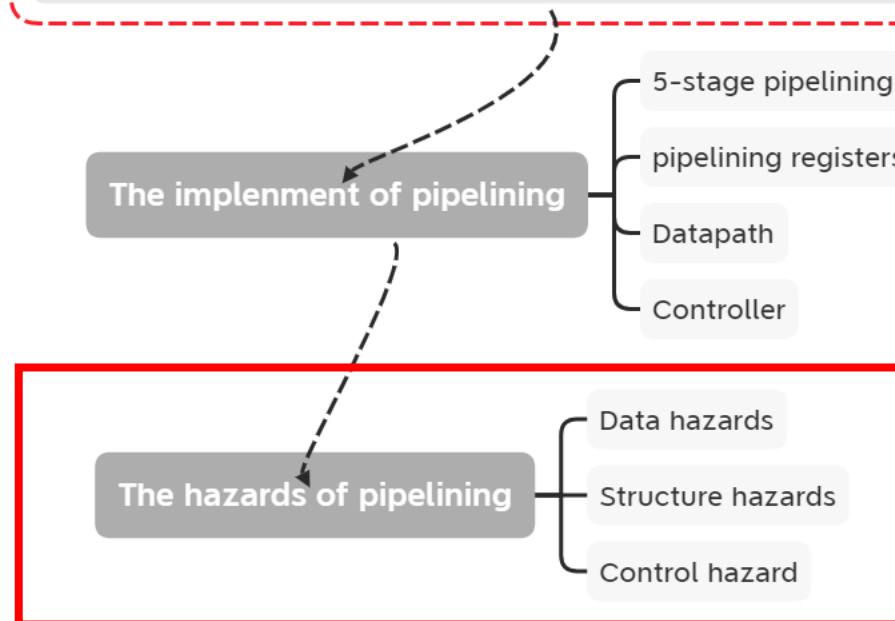
Decreasing the execution time of an individual instruction ✕

Increasing instruction throughput ✓



Pipelining

Are pipeline always execute commands correctly?



Pipeline Hazards


- Structural Hazard A required resource is busy
- Data Hazards
 - Data dependency between instructions
 - Need to wait for previous instruction to complete its data read/write
- Control Hazards Flow of execution depends on previous instruction



Structural Hazard

A required resource is busy

Example: Consider the situation while the pipeline only has a single memory



Question: Can the four instructions execute correctly?



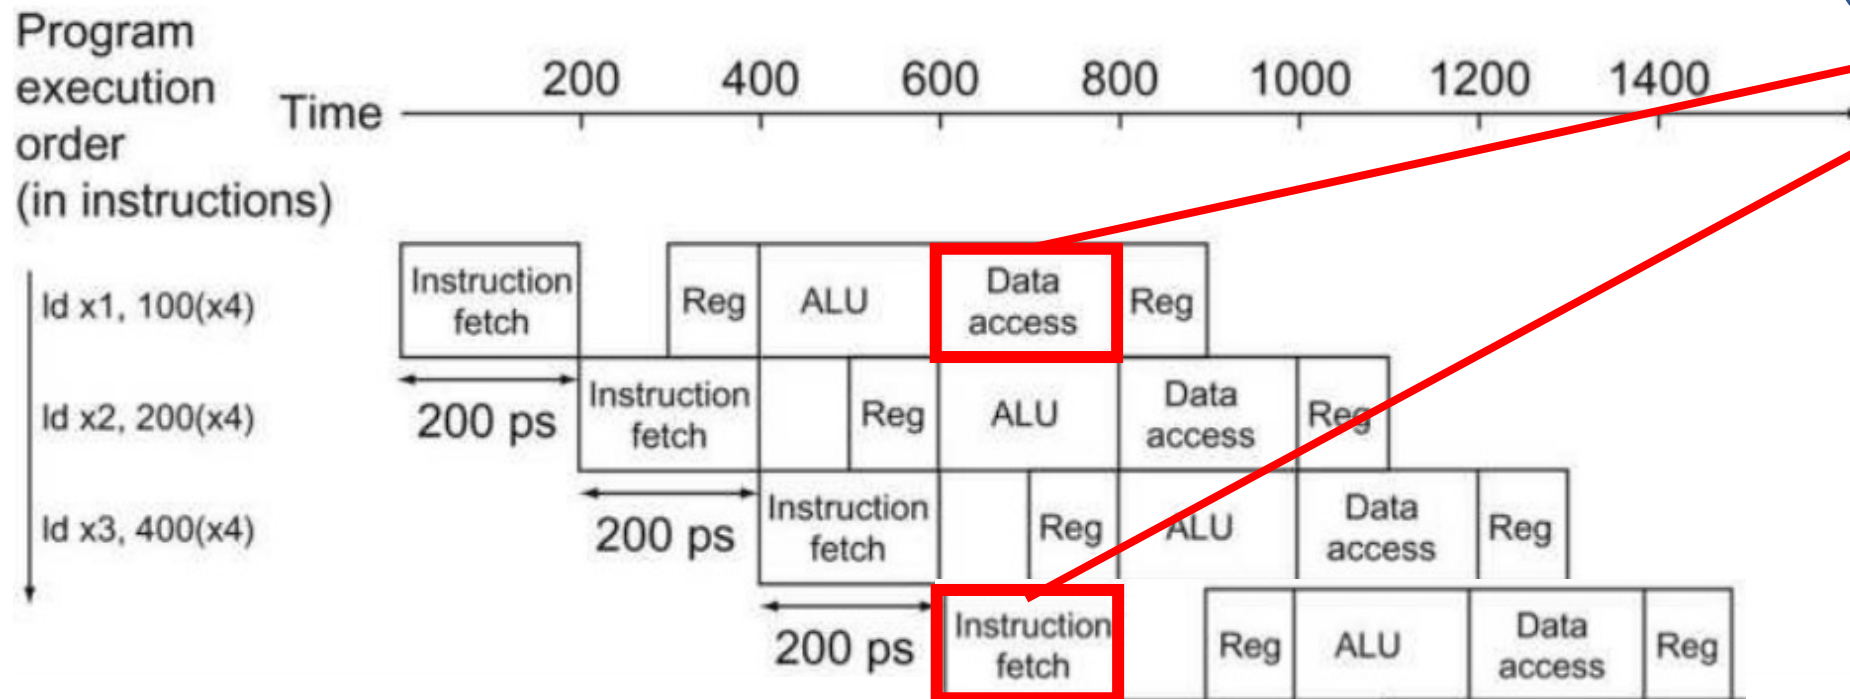
Structural Hazard

A required resource is busy

Solution:

Use Instruction and data memory simultaneously

Example: Consider the situation while the pipeline only has **a single memory**



How to Deal with Structural Hazard?

Problem: Two or more instructions in the pipeline compete for access to a single physical resource

- Solution 1: Instructions take it in turns to use resource, some instructions have to stall
- Solution 2: Add more hardware to machine


Can always solve a structural hazard by adding more hardware



How to Stall ?

NOP instruction

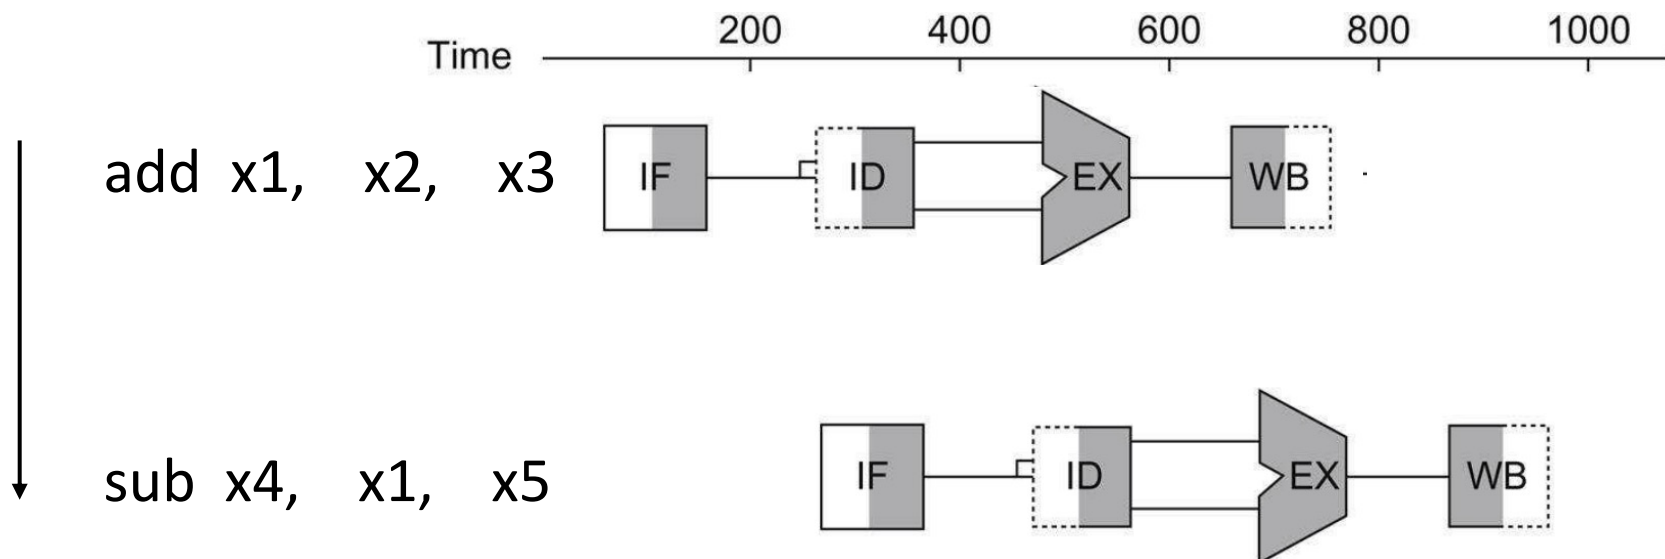
ADDI x0, x0, 0



Data Hazard

- Data dependency between instructions
- Need to wait for previous instruction to complete its data read/write

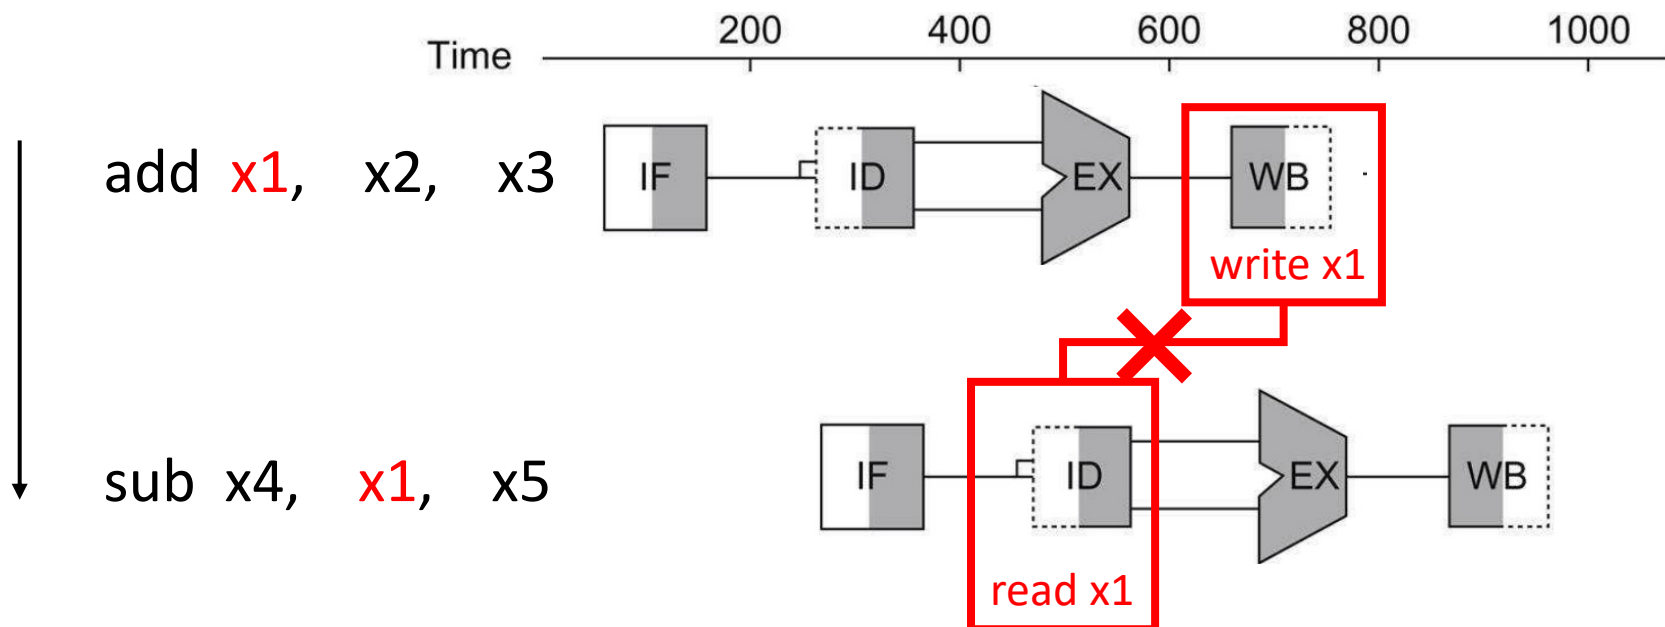
Problem: Instruction depends on result from previous



Data Hazard

- Data dependency between instructions
- Need to wait for previous instruction to complete its data read/write

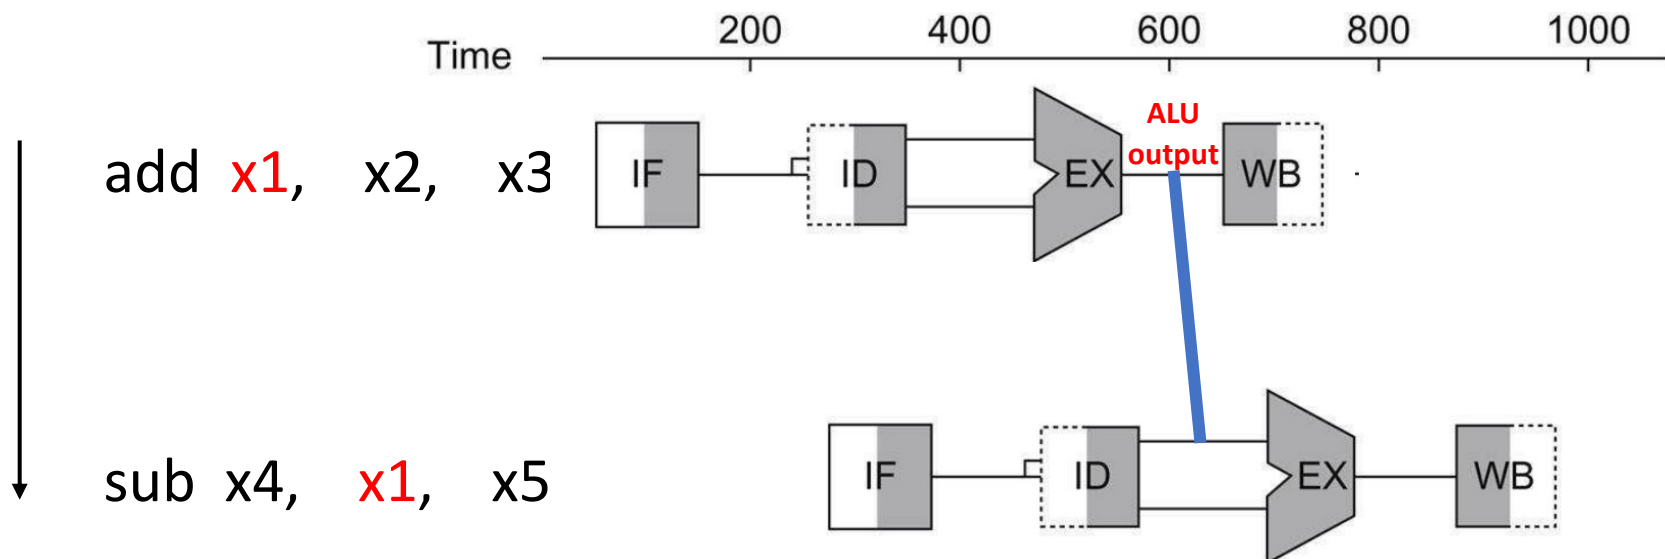
Problem: Instruction depends on result from previous



Data Hazard

- Data dependency between instructions
- Need to wait for previous instruction to complete its data read/write

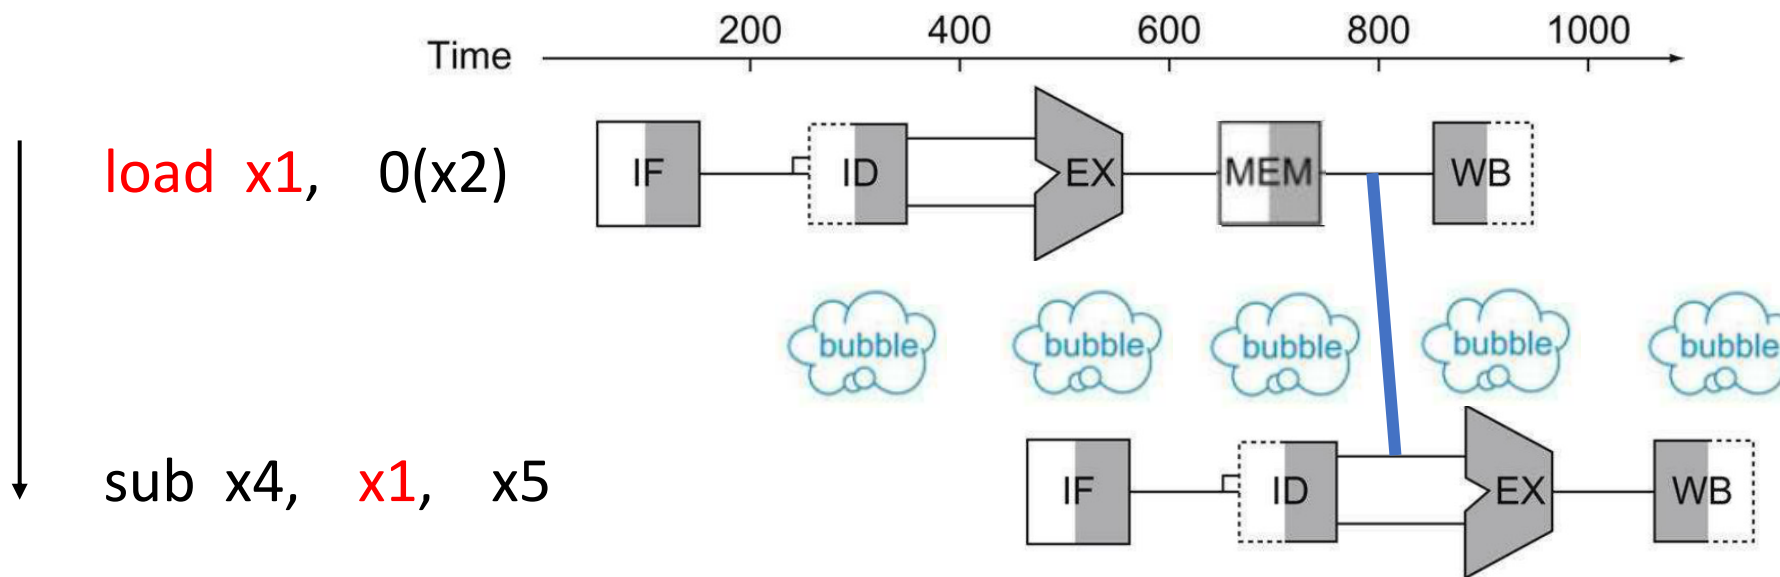
Solution “forwarding”: Adding extra hardware to retrieve the missing item early from the internal resources



Data Hazard

- Data dependency between instructions
- Need to wait for previous instruction to complete its data read/write

Solution “forwarding”: Could not avoid all pipeline stalls



Data Hazard

- Data dependency between instructions
- Need to wait for previous instruction to complete its data read/write

Example: Reordering code to avoid pipeline stalls

Consider the following code segment in C:

`a = b + e;`

`c = b + f;`

- Assuming all variables are in memory
- and are addressable as offsets from x31

The generated RISC-V code:

`ld x1, 0(x31) // Load b`

`ld x2, 8(x31) // Load e`

`add x3, x1, x2 // b + e`

`sd x3, 24(x31) // Store a`

`ld x4, 16(x31) // Load f`

`add x5, x1, x4 // b + f`

`sd x5, 32(x31) // Store c`

Question: Find the data hazard and reorder RISC-V code



Data Hazard

- Data dependency between instructions
- Need to wait for previous instruction to complete its data read/write

Example: Reordering code to avoid pipeline stalls

Consider the following code segment in C:

`a = b + e;`

`c = b + f;`

- Assuming all variables are in memory
- and are addressable as offsets from x31

The generated RISC-V code:

`ld x1, 0(x31) // Load b`

`ld x2, 8(x31) // Load e`

`add x3, x1, x2 // b + e`

`sd x3, 24(x31) // Store a`

`ld x4, 16(x31) // Load f`

`add x5, x1, x4 // b + f`

`sd x5, 32(x31) // Store c`



Data Hazard

- Data dependency between instructions
- Need to wait for previous instruction to complete its data read/write

Example: Reordering code to avoid pipeline stalls

Consider the following code segment in C:

`a = b + e;`

`c = b + f;`

- Assuming all variables are in memory
- and are addressable as offsets from x31

The generated RISC-V code:

`ld x1, 0(x31) // Load b`

`ld x2, 8(x31) // Load e`

`add x3, x1, x2 // b + e`

`sd x3, 24(x31) // Store a`

`ld x4, 16(x31) // Load f`

`add x5, x1, x4 // b + f`

`sd x5, 32(x31) // Store c`

eliminates
both hazards

`ld x1, 0(x31)`

`ld x2, 8(x31)`

`ld x4, 16(x31)`

`add x3, x1, x2`

`sd x3, 24(x31)`

`add x5, x1, x4`

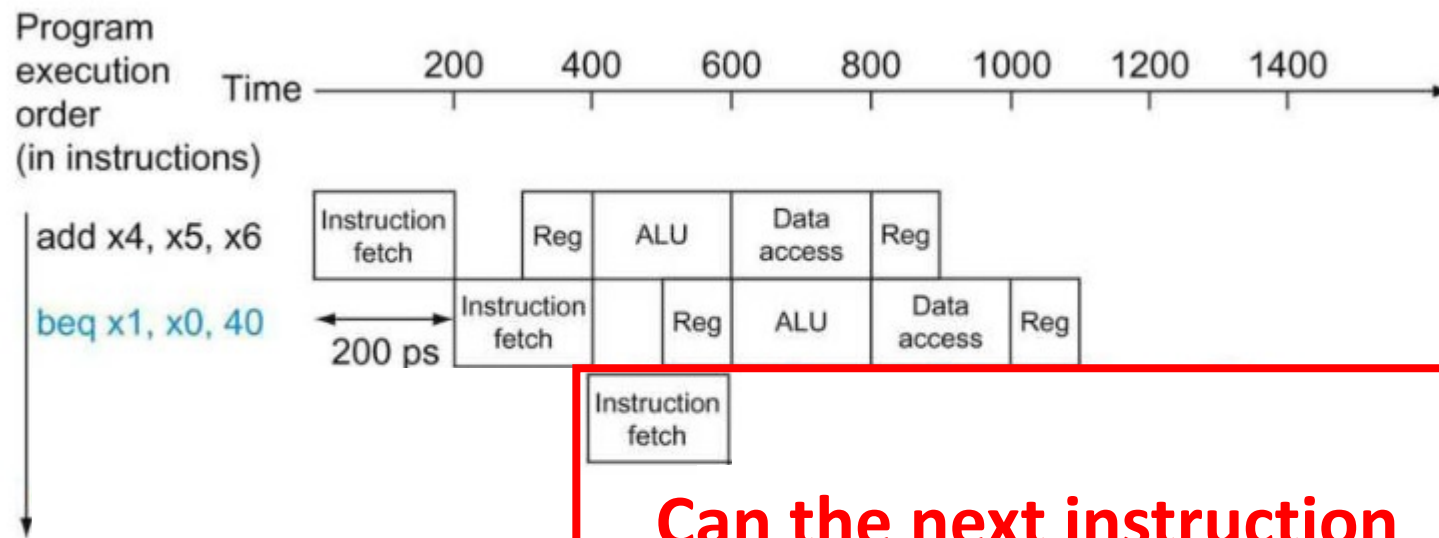
`sd x5, 32(x31)`



Control Hazard

Flow of execution depends on previous instruction

Problem: The conditional branch instruction




**Can the next instruction
be executed immediately?**



Control Hazard

Flow of execution depends on previous instruction

Problem: The conditional branch instruction




the pipeline cannot possibly know
what the next instruction should be



Control Hazard

Flow of execution depends on previous instruction

Solution 1: Stall



- test a register
- calculate the branch address
- update the PC

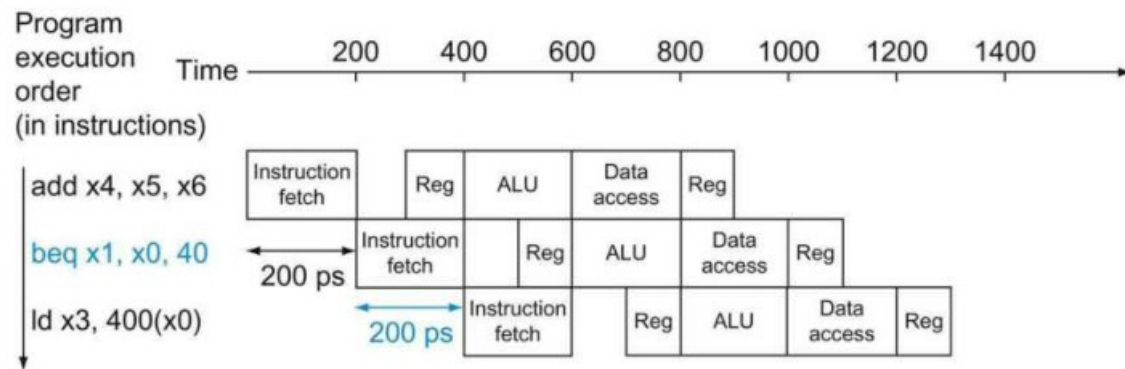
Control Hazard

Flow of execution depends on previous instruction

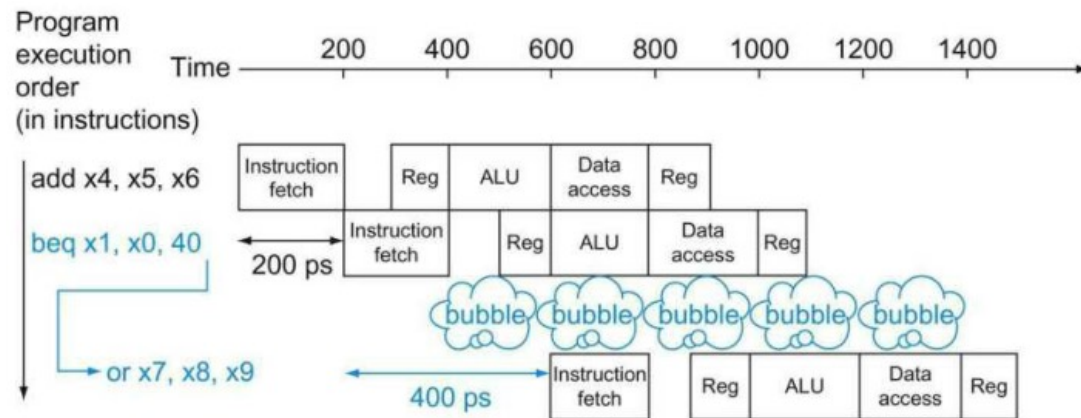
Solution 2: Prediction (simple version)

Predicts **always** that conditional branches will be untaken

branch is not taken



branch is taken




Control Hazard

Flow of execution depends on previous instruction

Solution 2: Prediction (sophisticated version)

Predicts that **some** conditional branches will be untaken

Example: While the Conditional Branches at the bottom of loops



**they are likely to be taken to
the top of the loops**

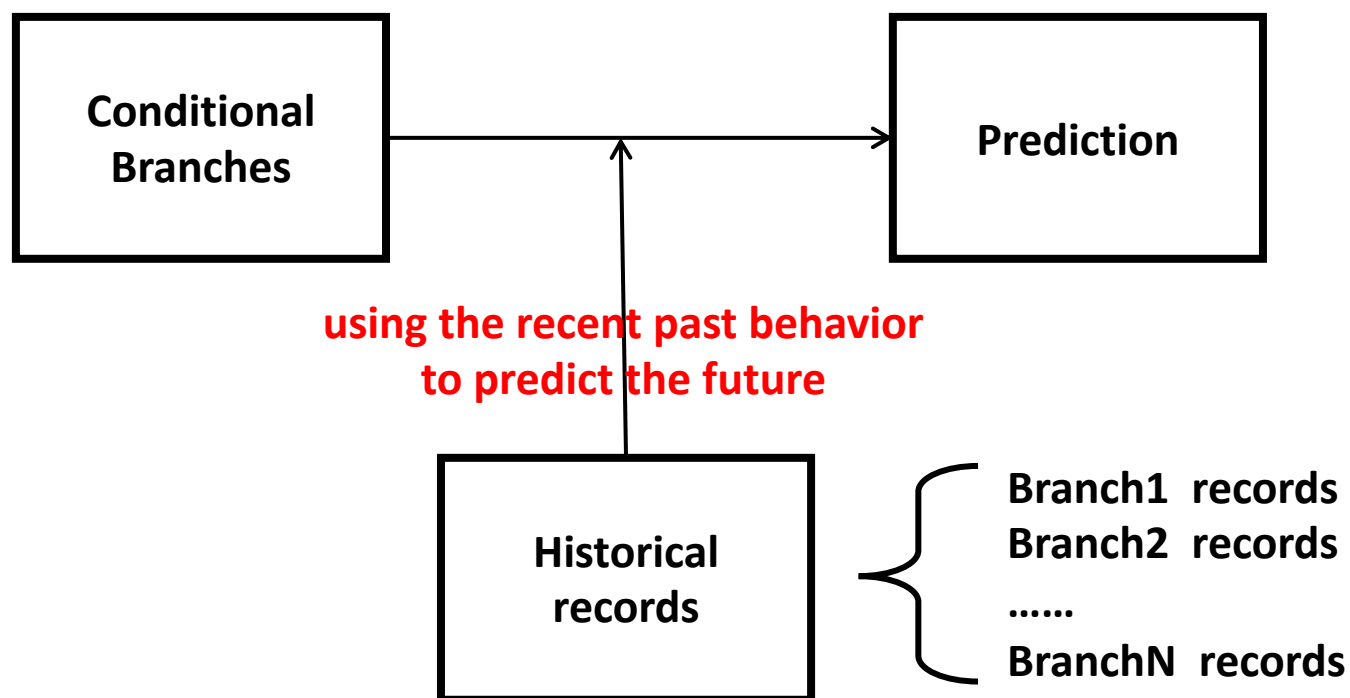


Control Hazard

Flow of execution depends on previous instruction

Solution 2: Prediction (dynamic prediction)

Predicts that **some** conditional branches will be untaken



Summary of Pipelining Design

- Pipeline is a technique that exploits **parallelism** between the instructions in a sequential instruction stream
- Pipeline increases the **number** of **simultaneously** executing instructions and the rate at which instructions are started and completed
- Pipeline designers need to cope with structural, control, and data **hazards**
- Pipeline is fundamentally **invisible** to the programmer

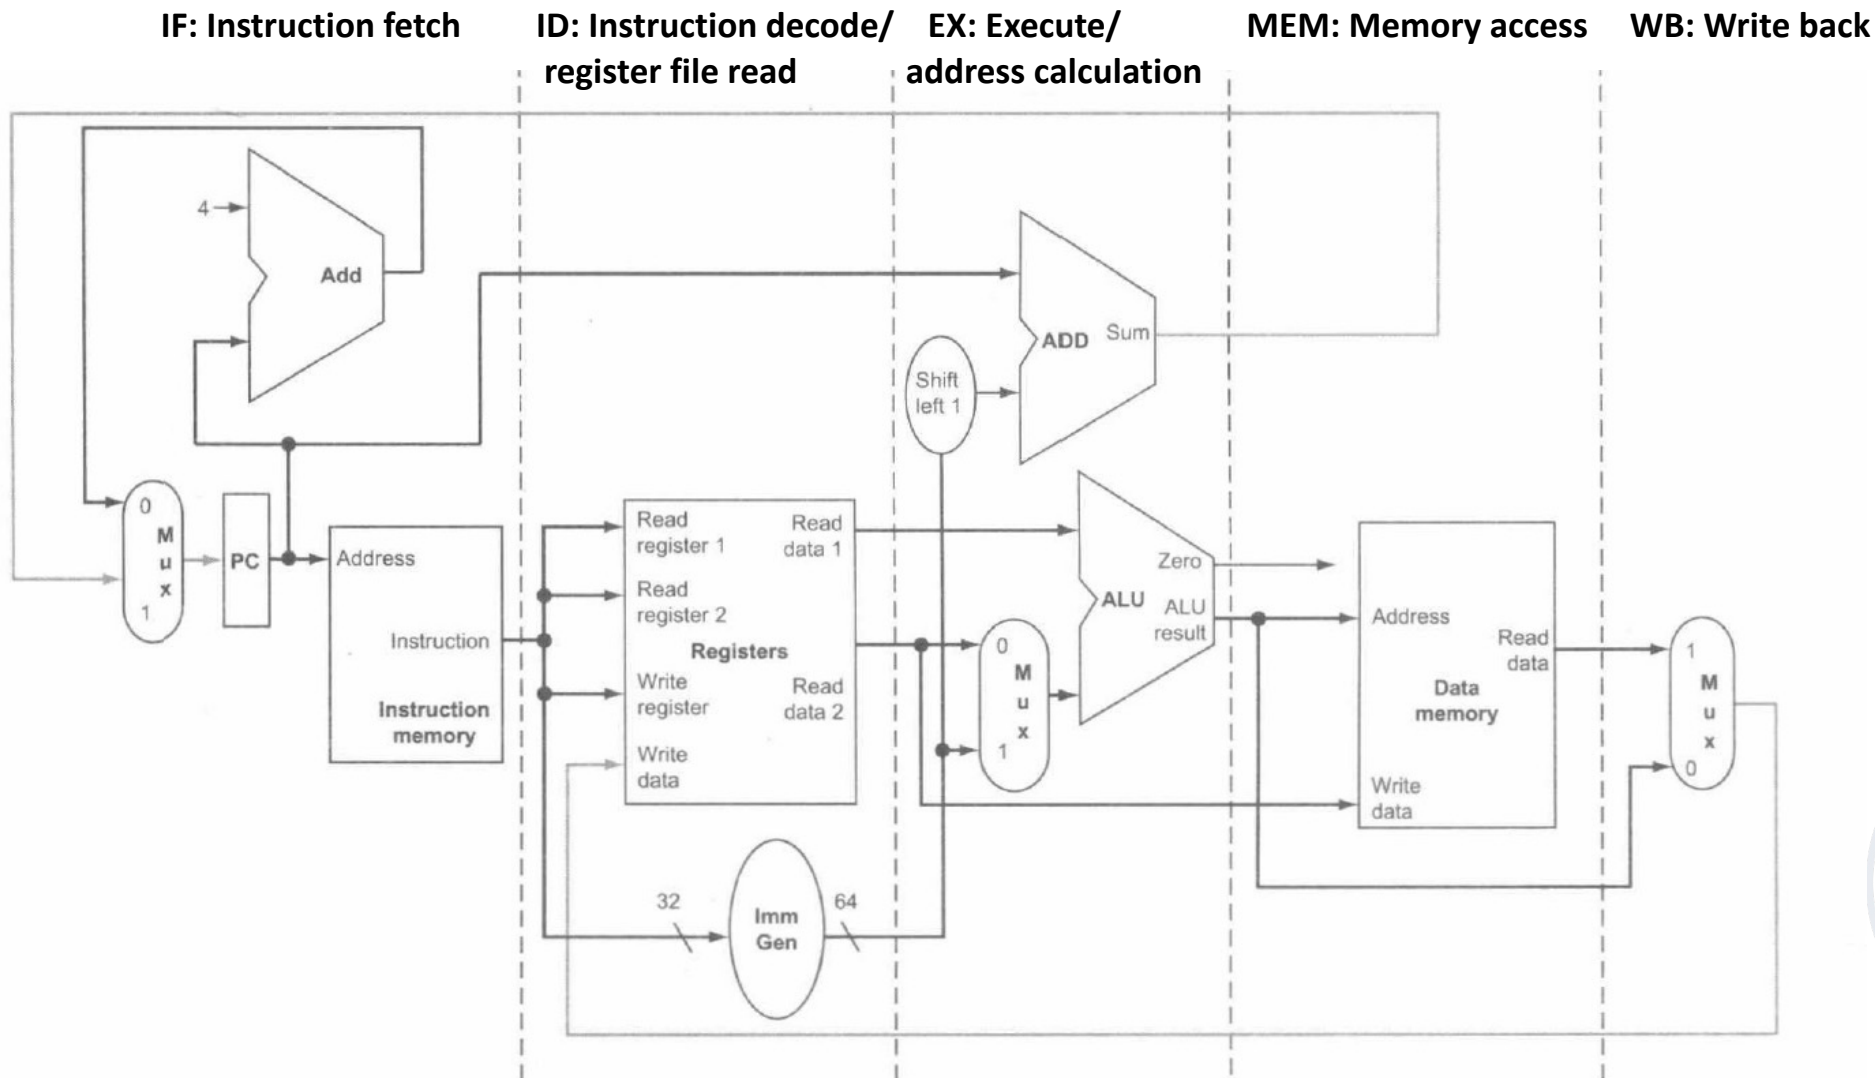


Pipelined Datapath



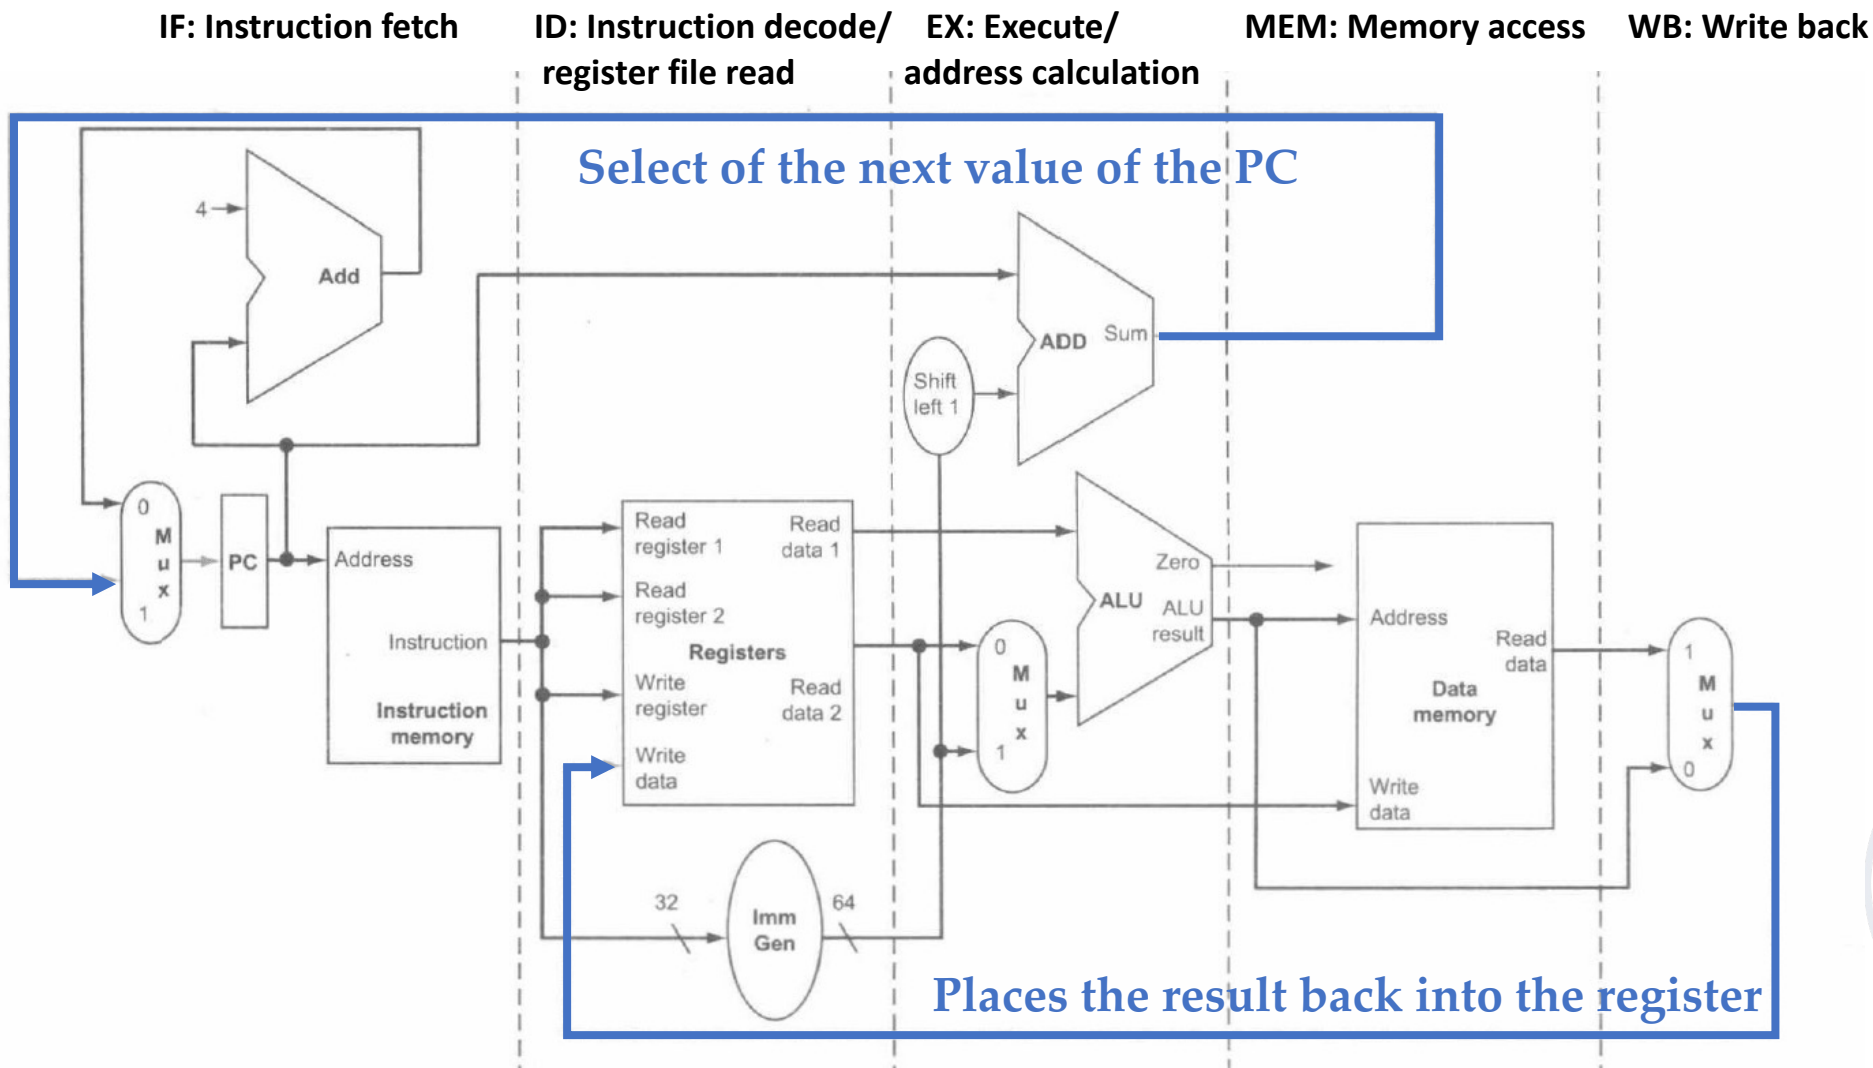
Single-Cycle Datapath

Let's find two exceptions
to this left-to-right flow of instructions



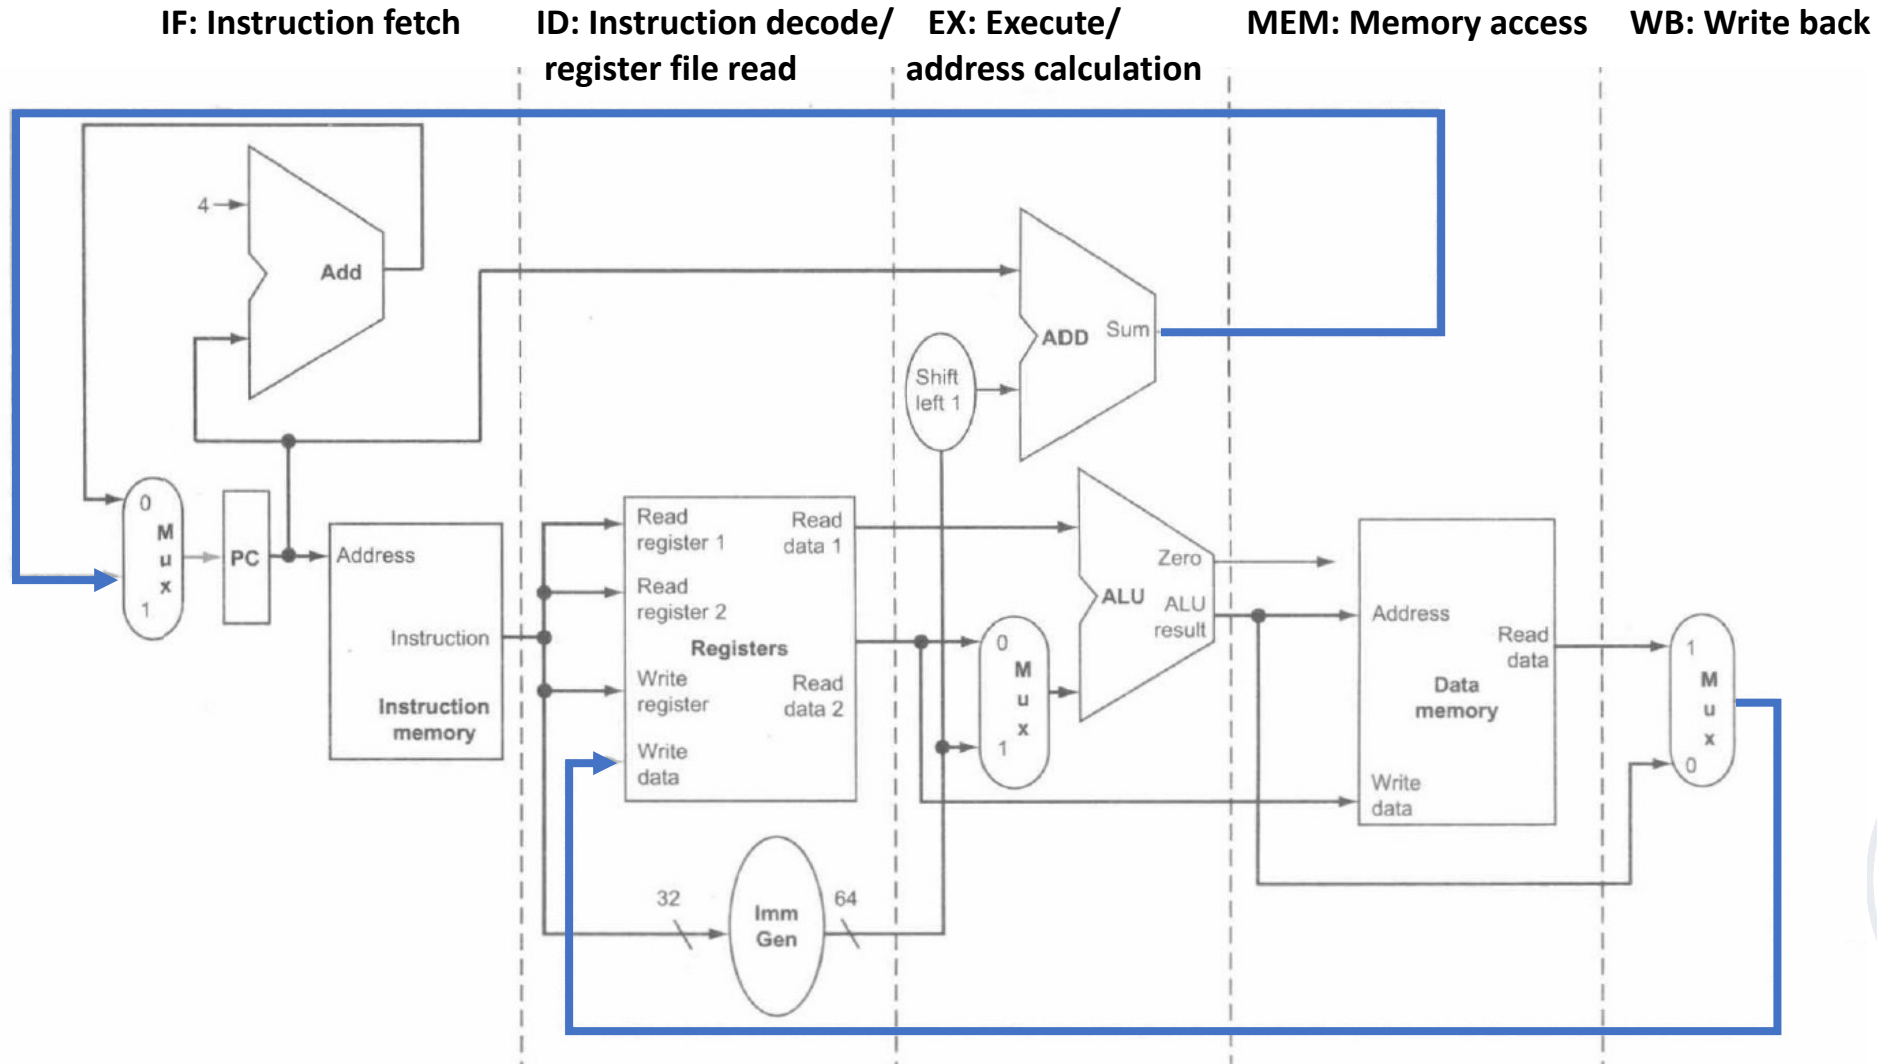
Single-Cycle Datapath

Let's find two exceptions
to this left-to-right flow of instructions



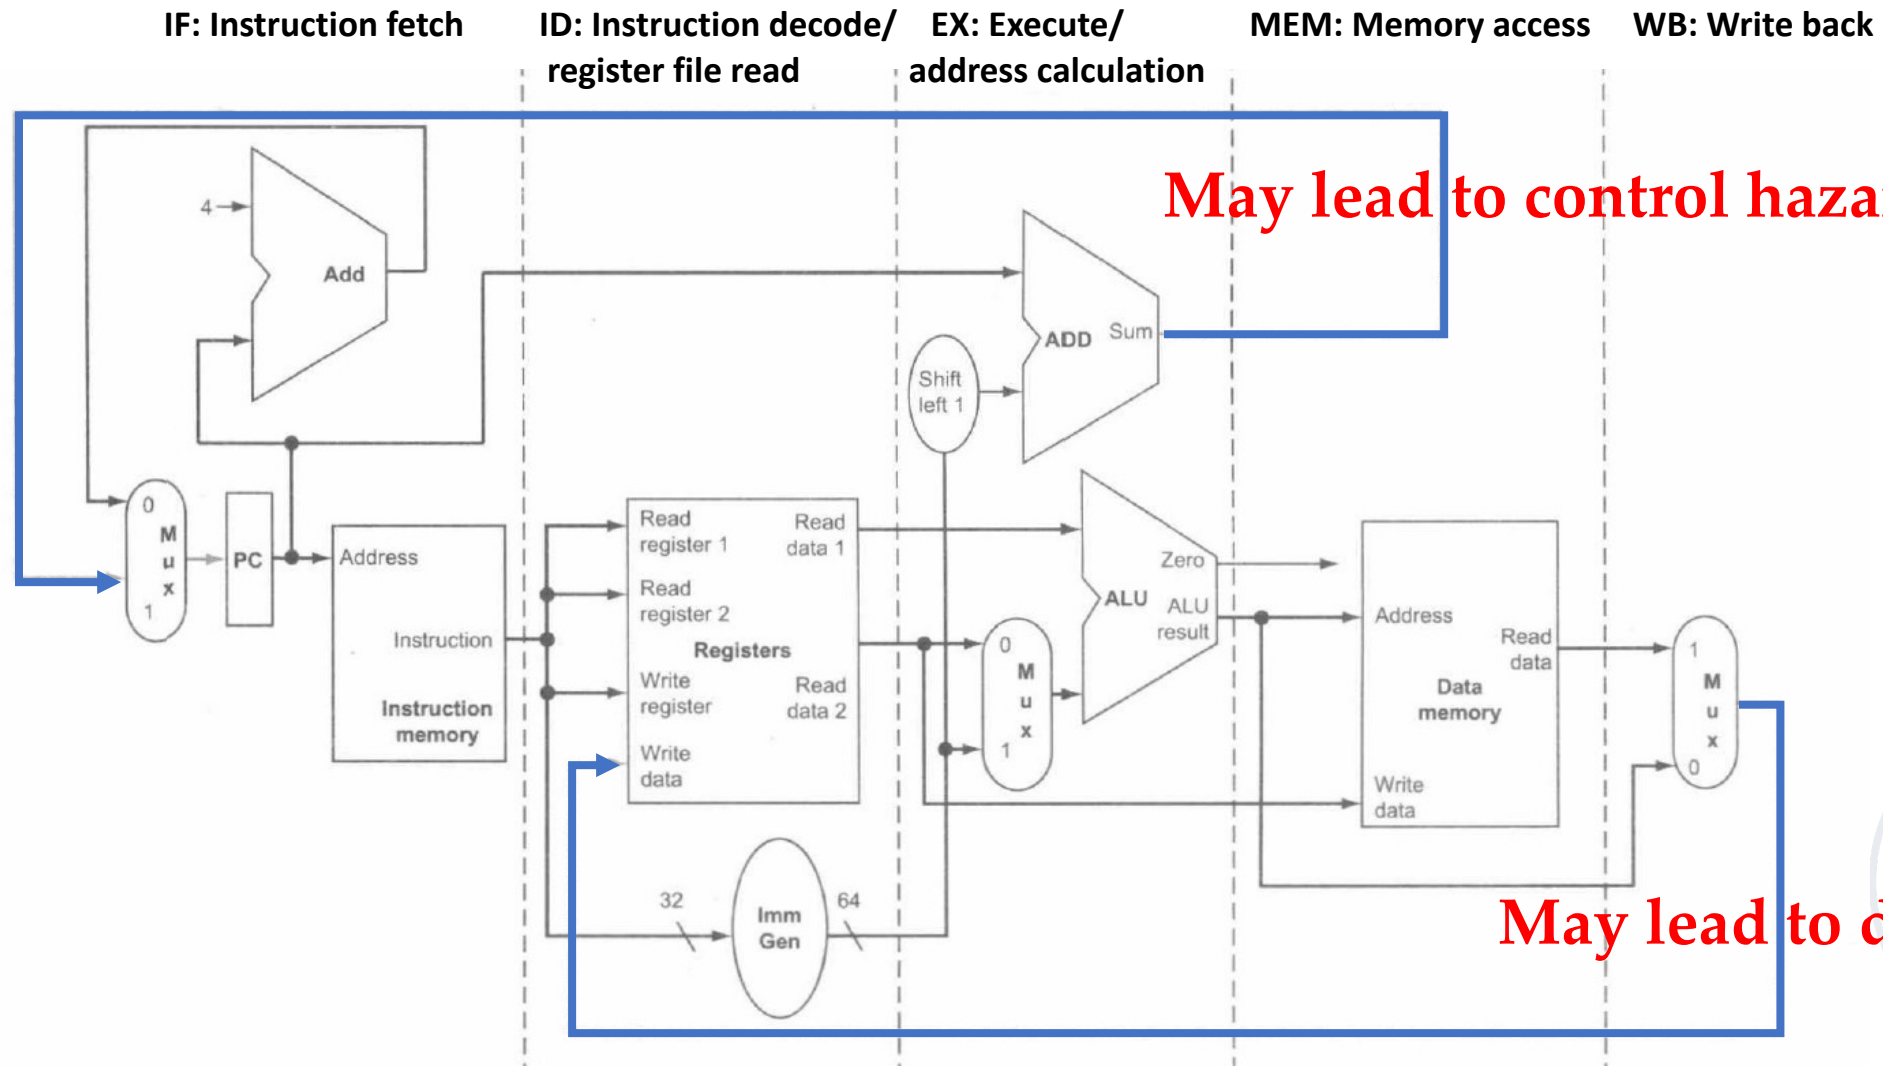
Single-Cycle Datapath

Question: The impacts of the two exceptions



Single-Cycle Datapath

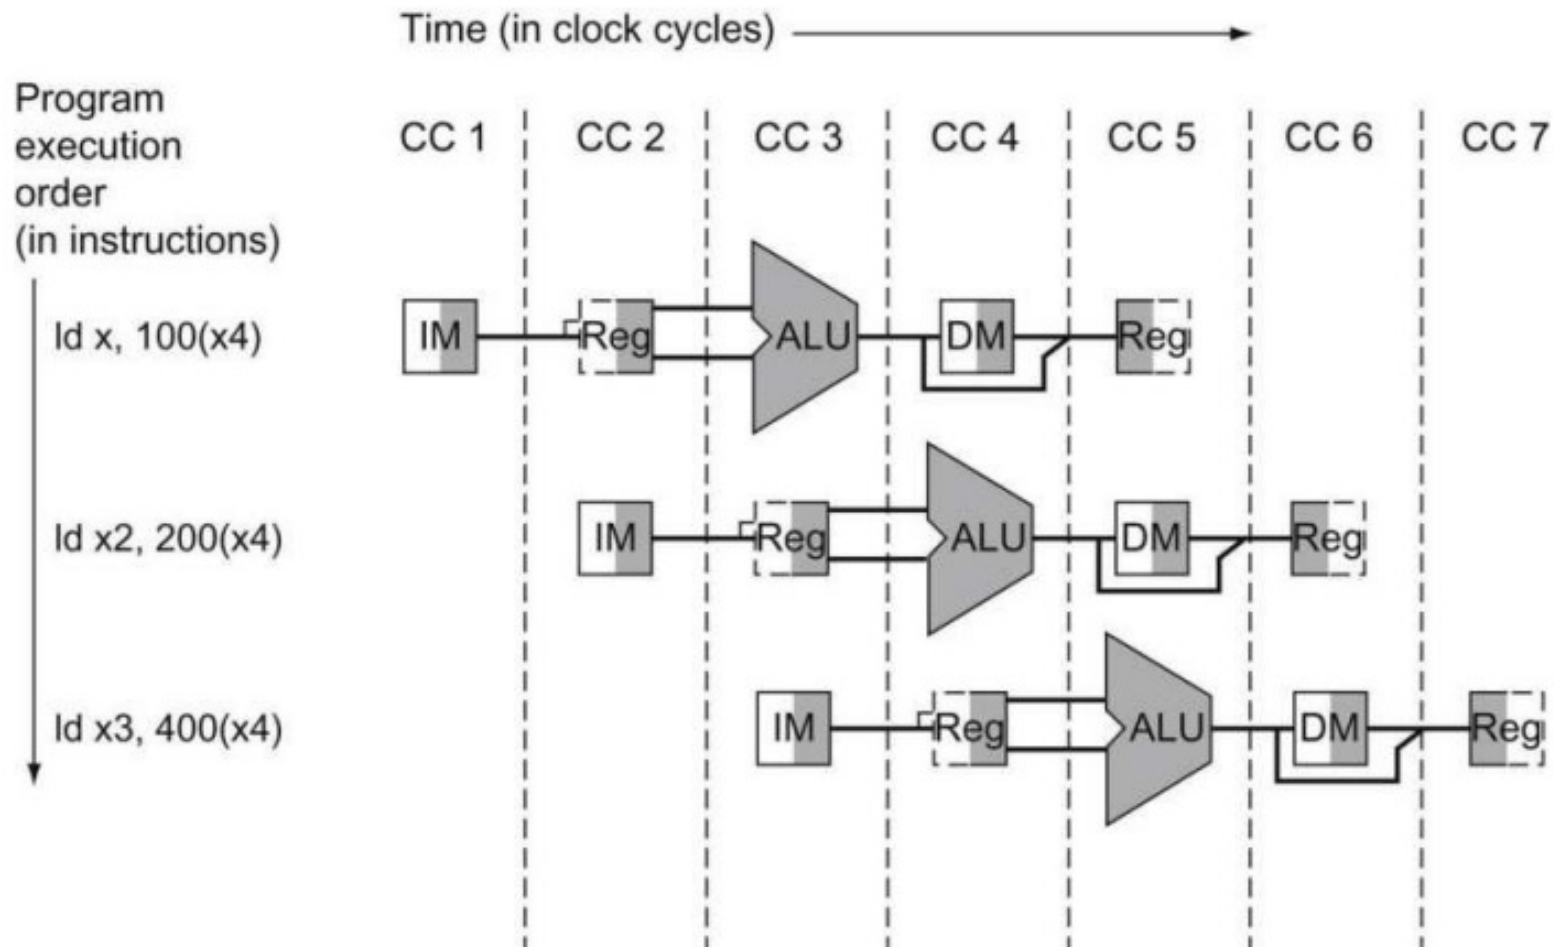
Question: The impacts of the two exceptions



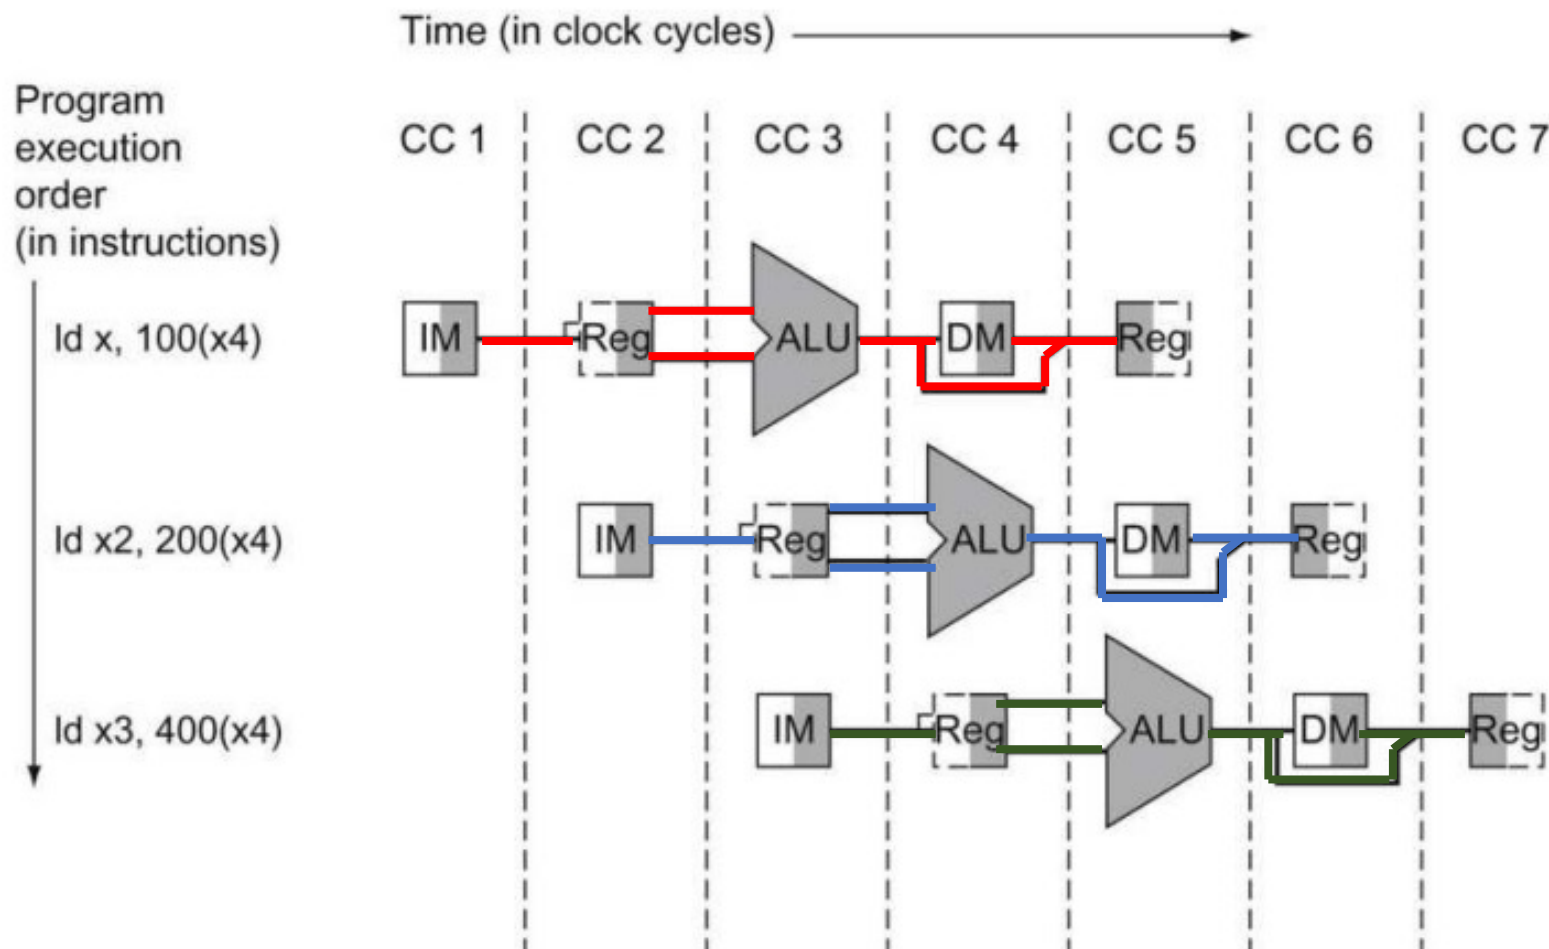
Single-Cycle Datapath to the Pipelined Version



Instructions Being Executed Using the Single-Cycle Datapath



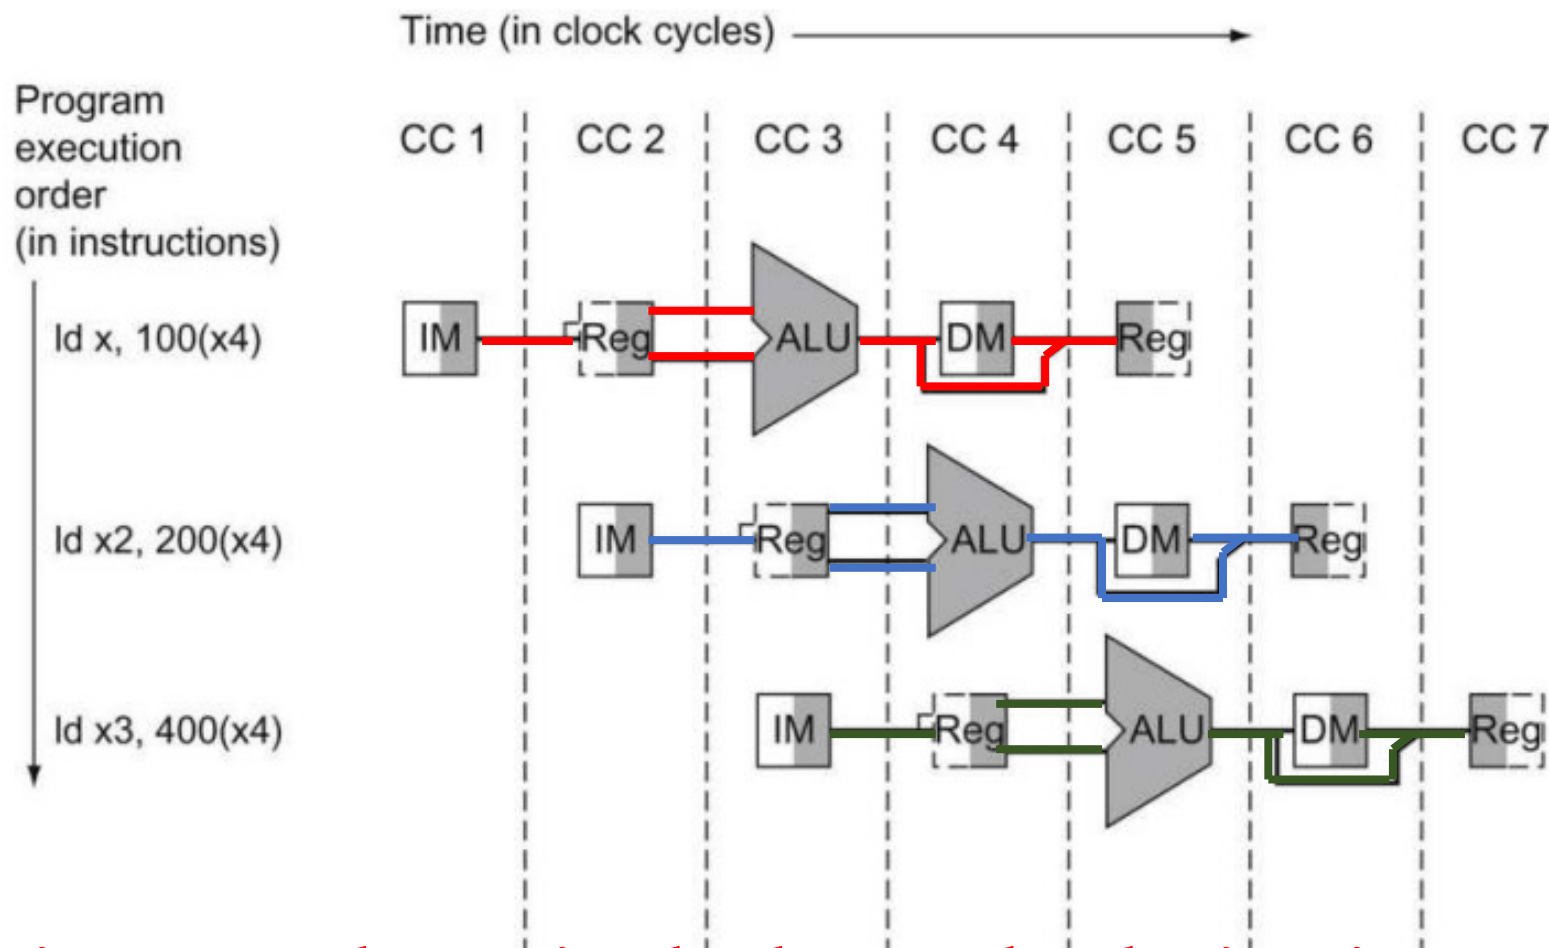
Instructions Being Executed Using the Single-Cycle Datapath



Three instructions need three datapaths

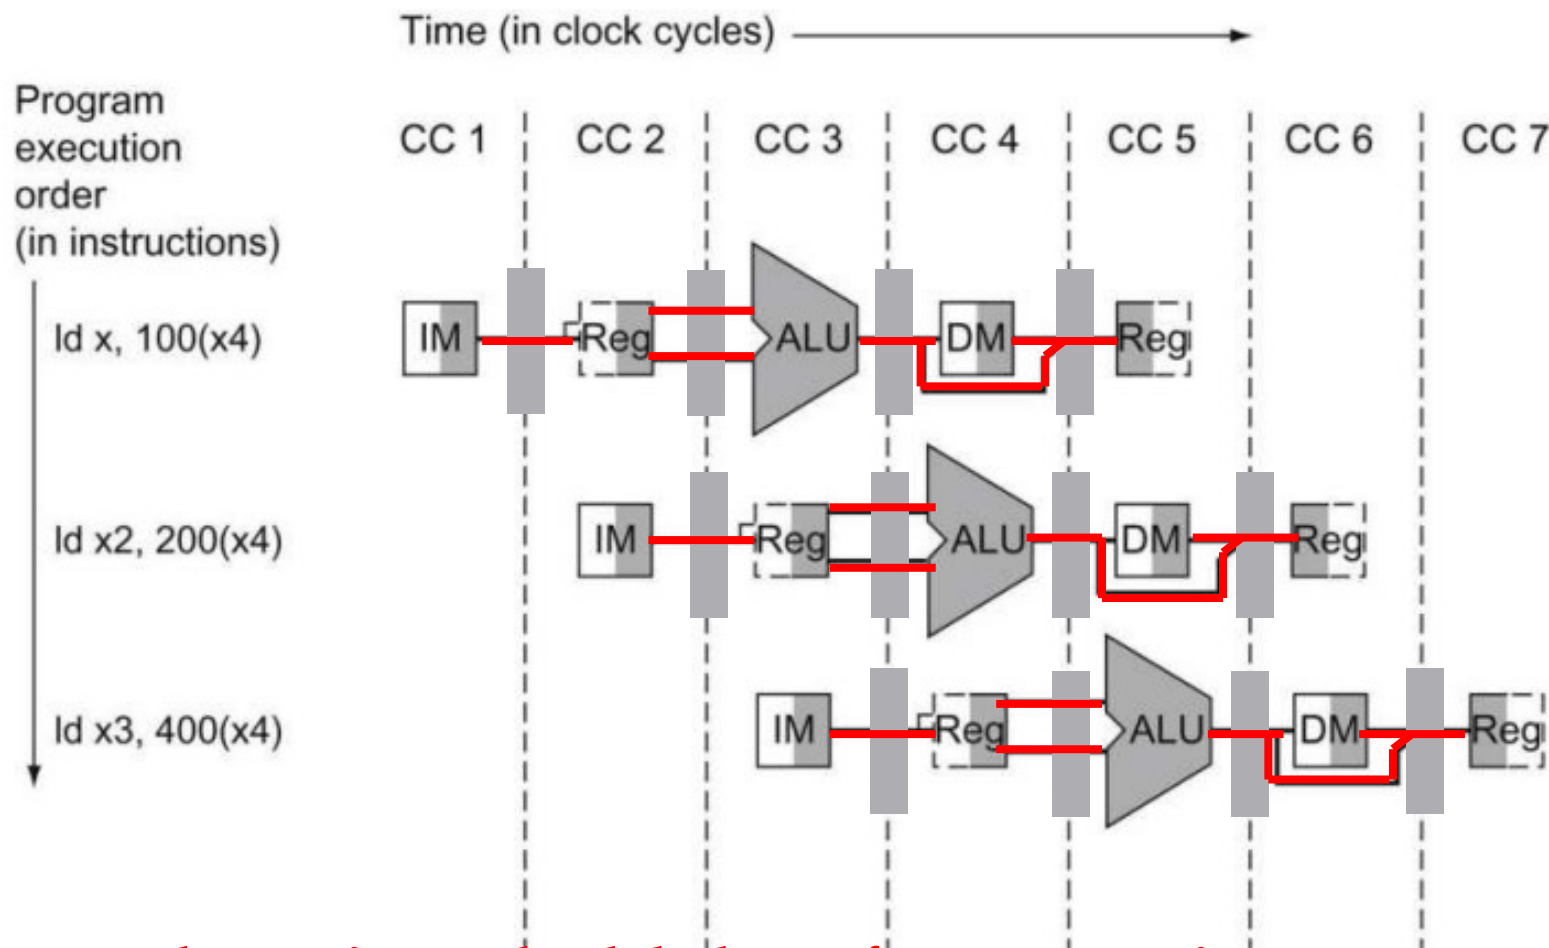


Instructions Being Executed Using the Single-Cycle Datapath



Let's add registers to share single datapaths during instruction execution

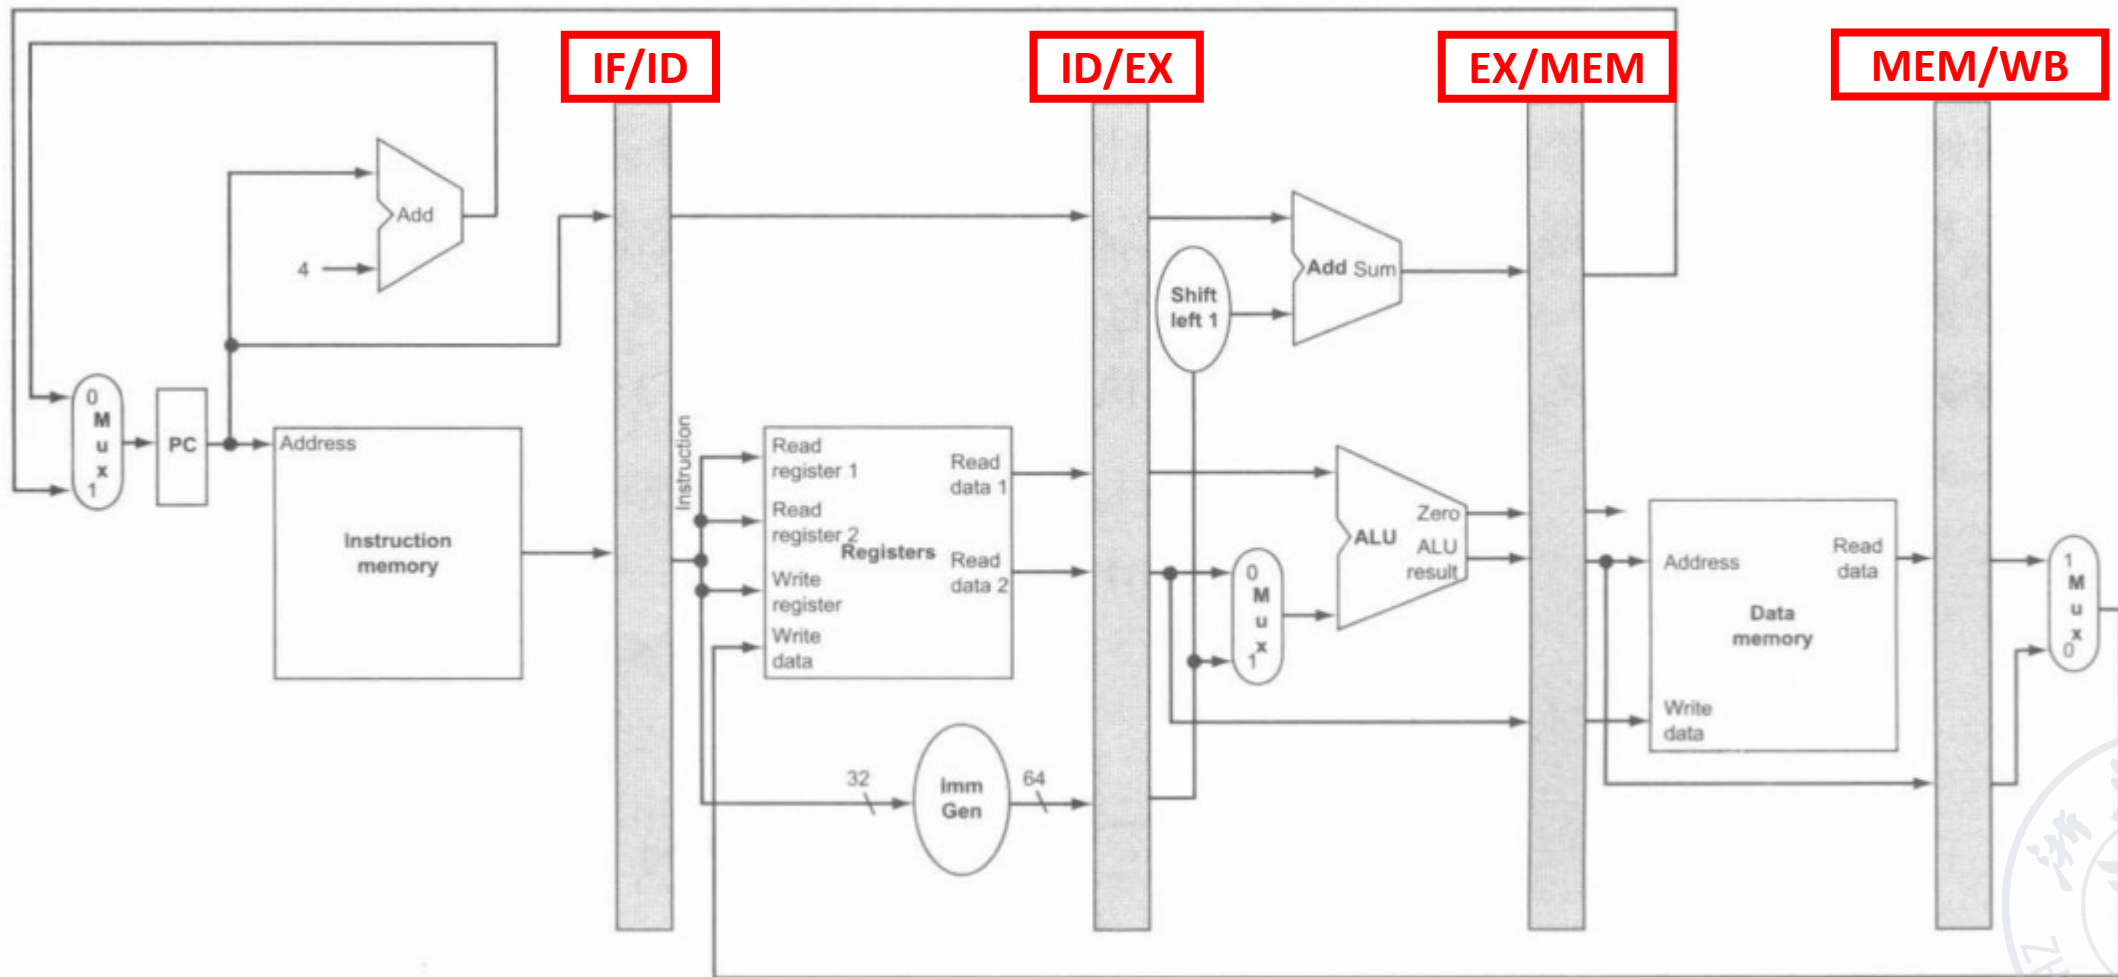
Instructions Being Executed Using the Single-Cycle Datapath



Each register hold data from previous stage



Pipelined Version of the Datapath



Understand How Pipelined Datapath Work

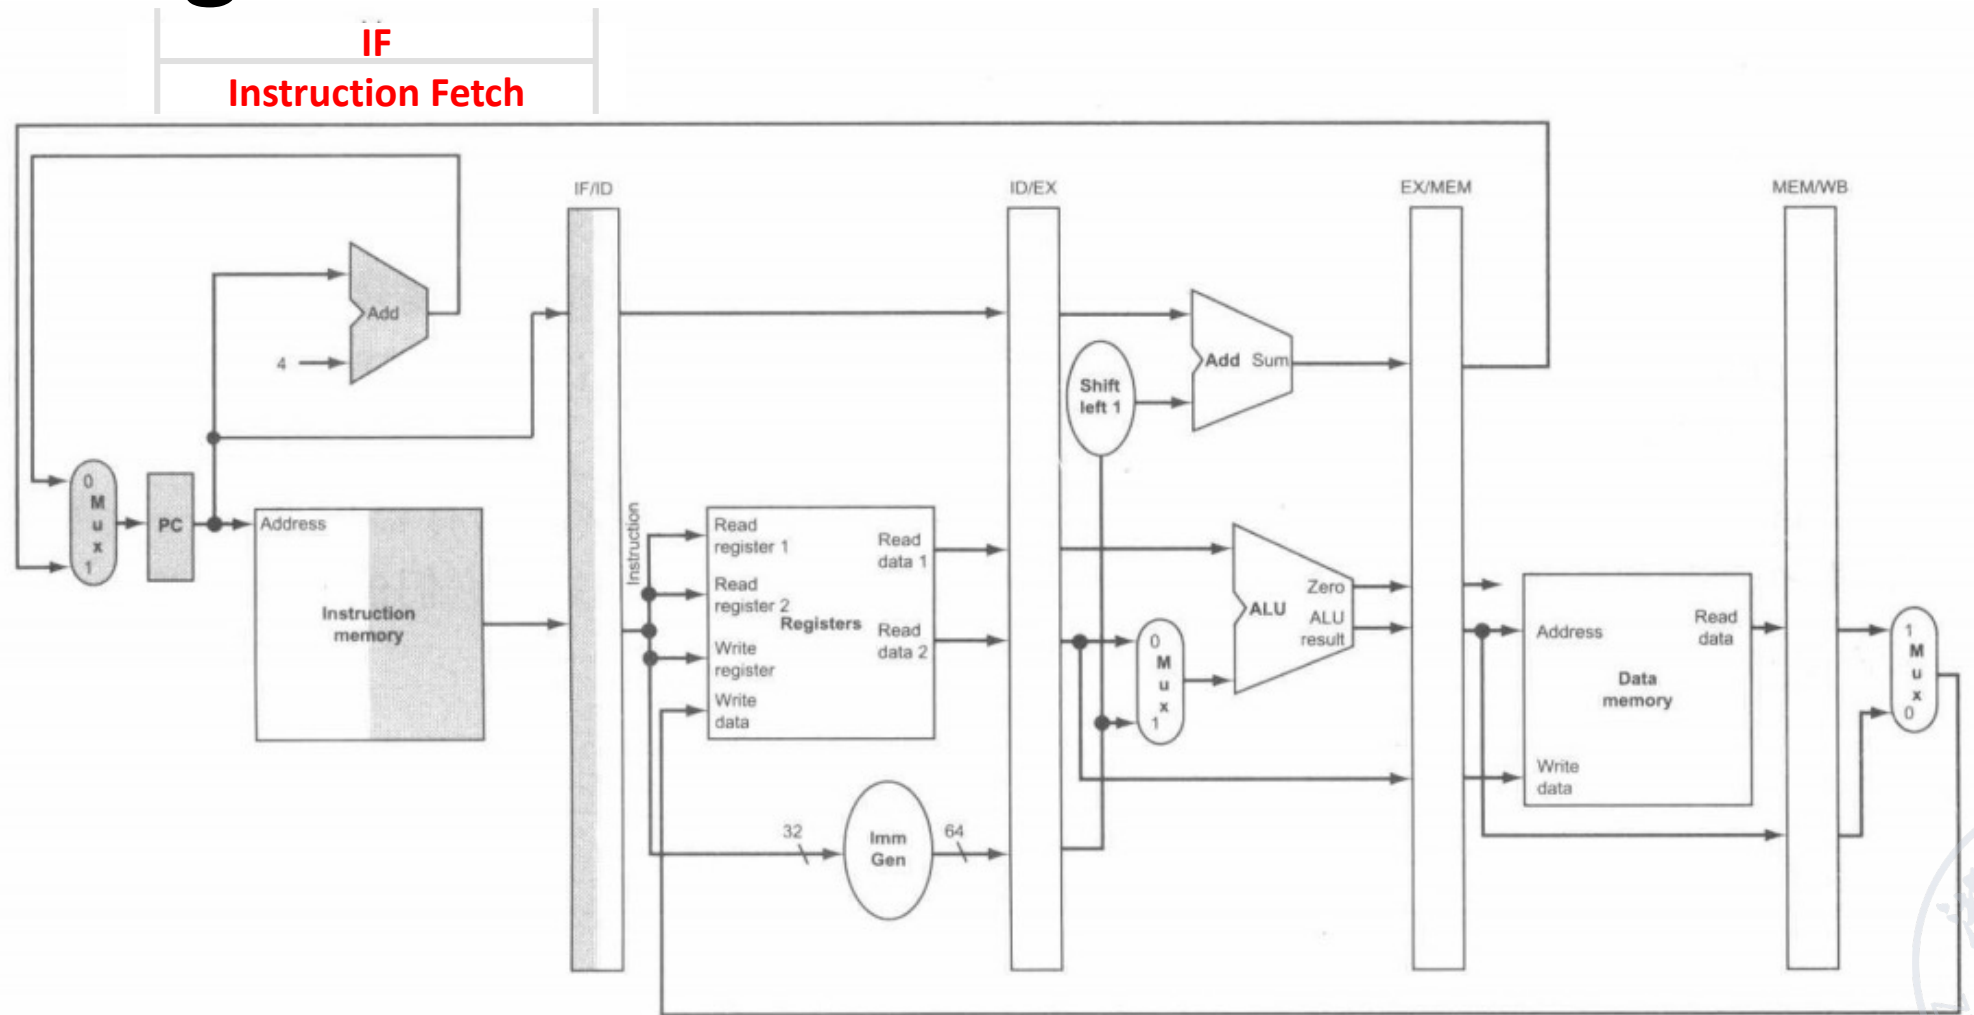
Through Load Instruction and Store Instruction



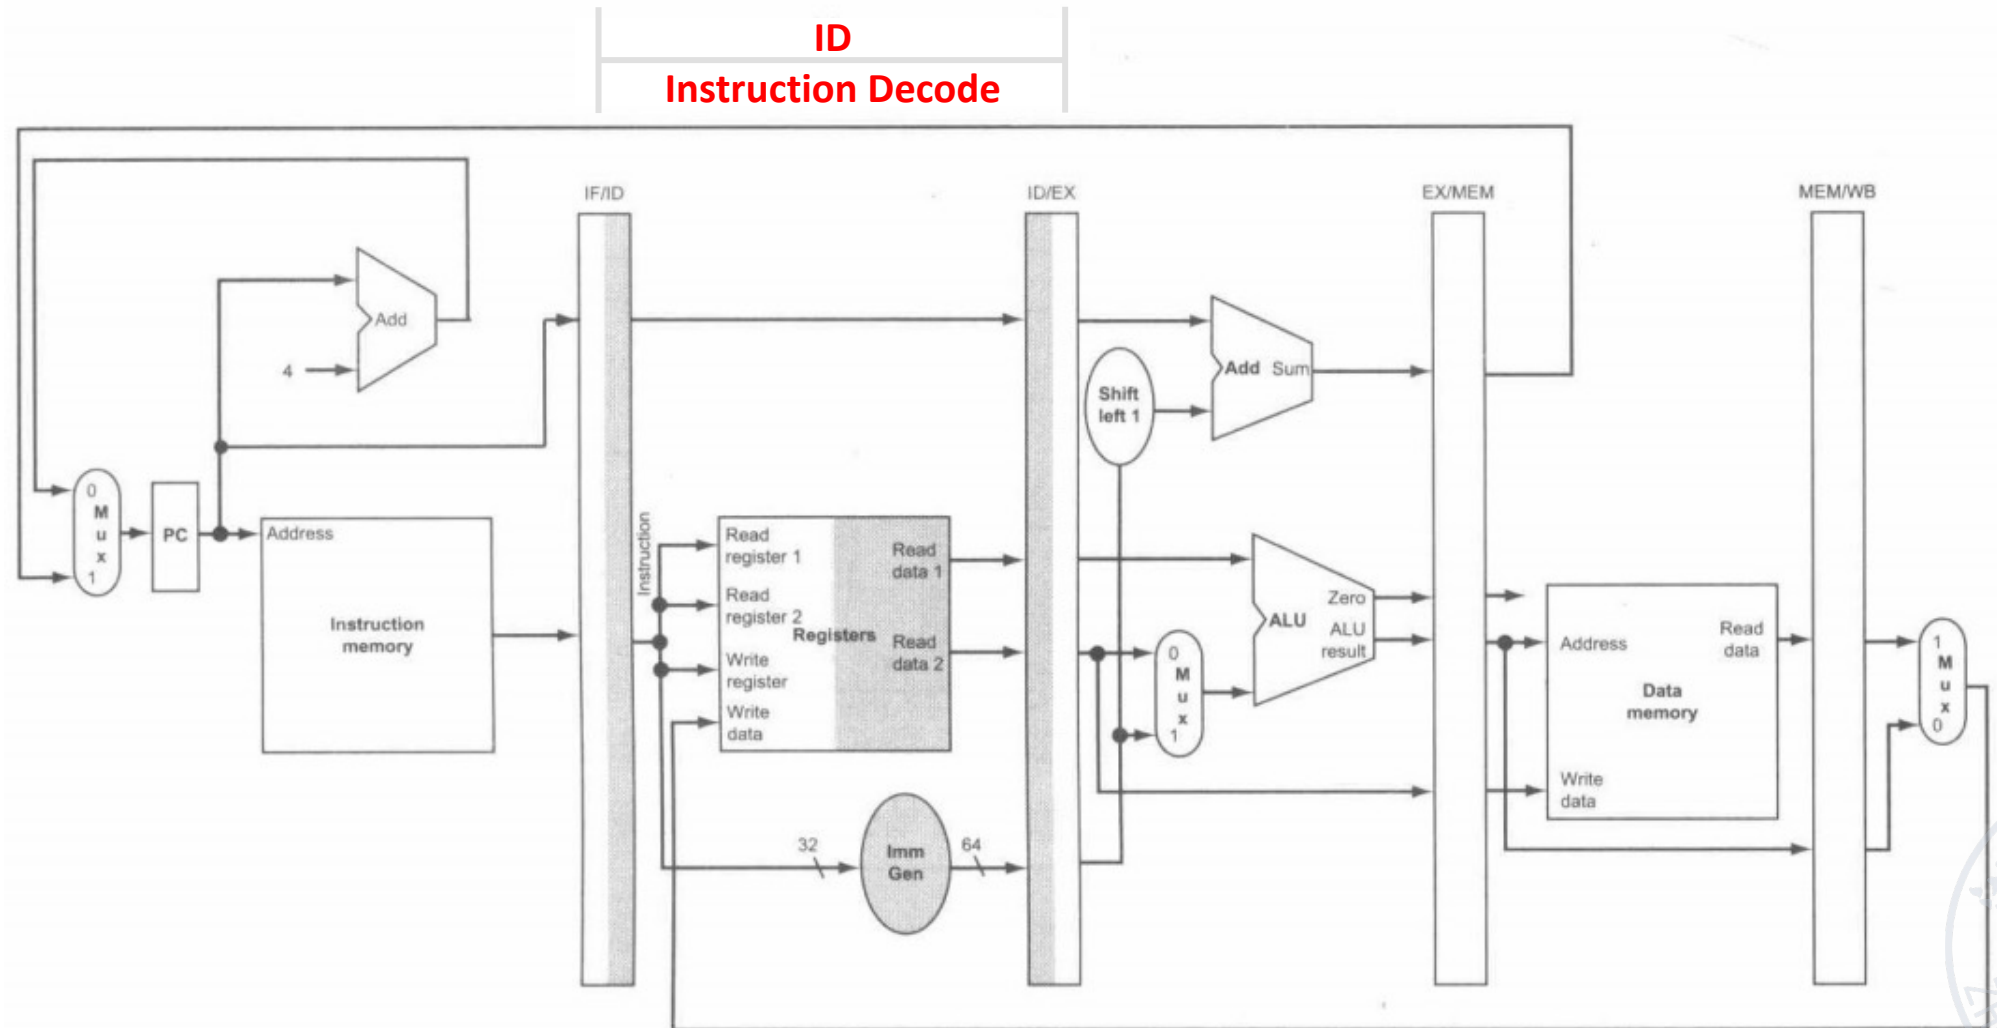
Load Instruction



IF Stage of **Load** Instruction




ID Stage of **Load** Instruction

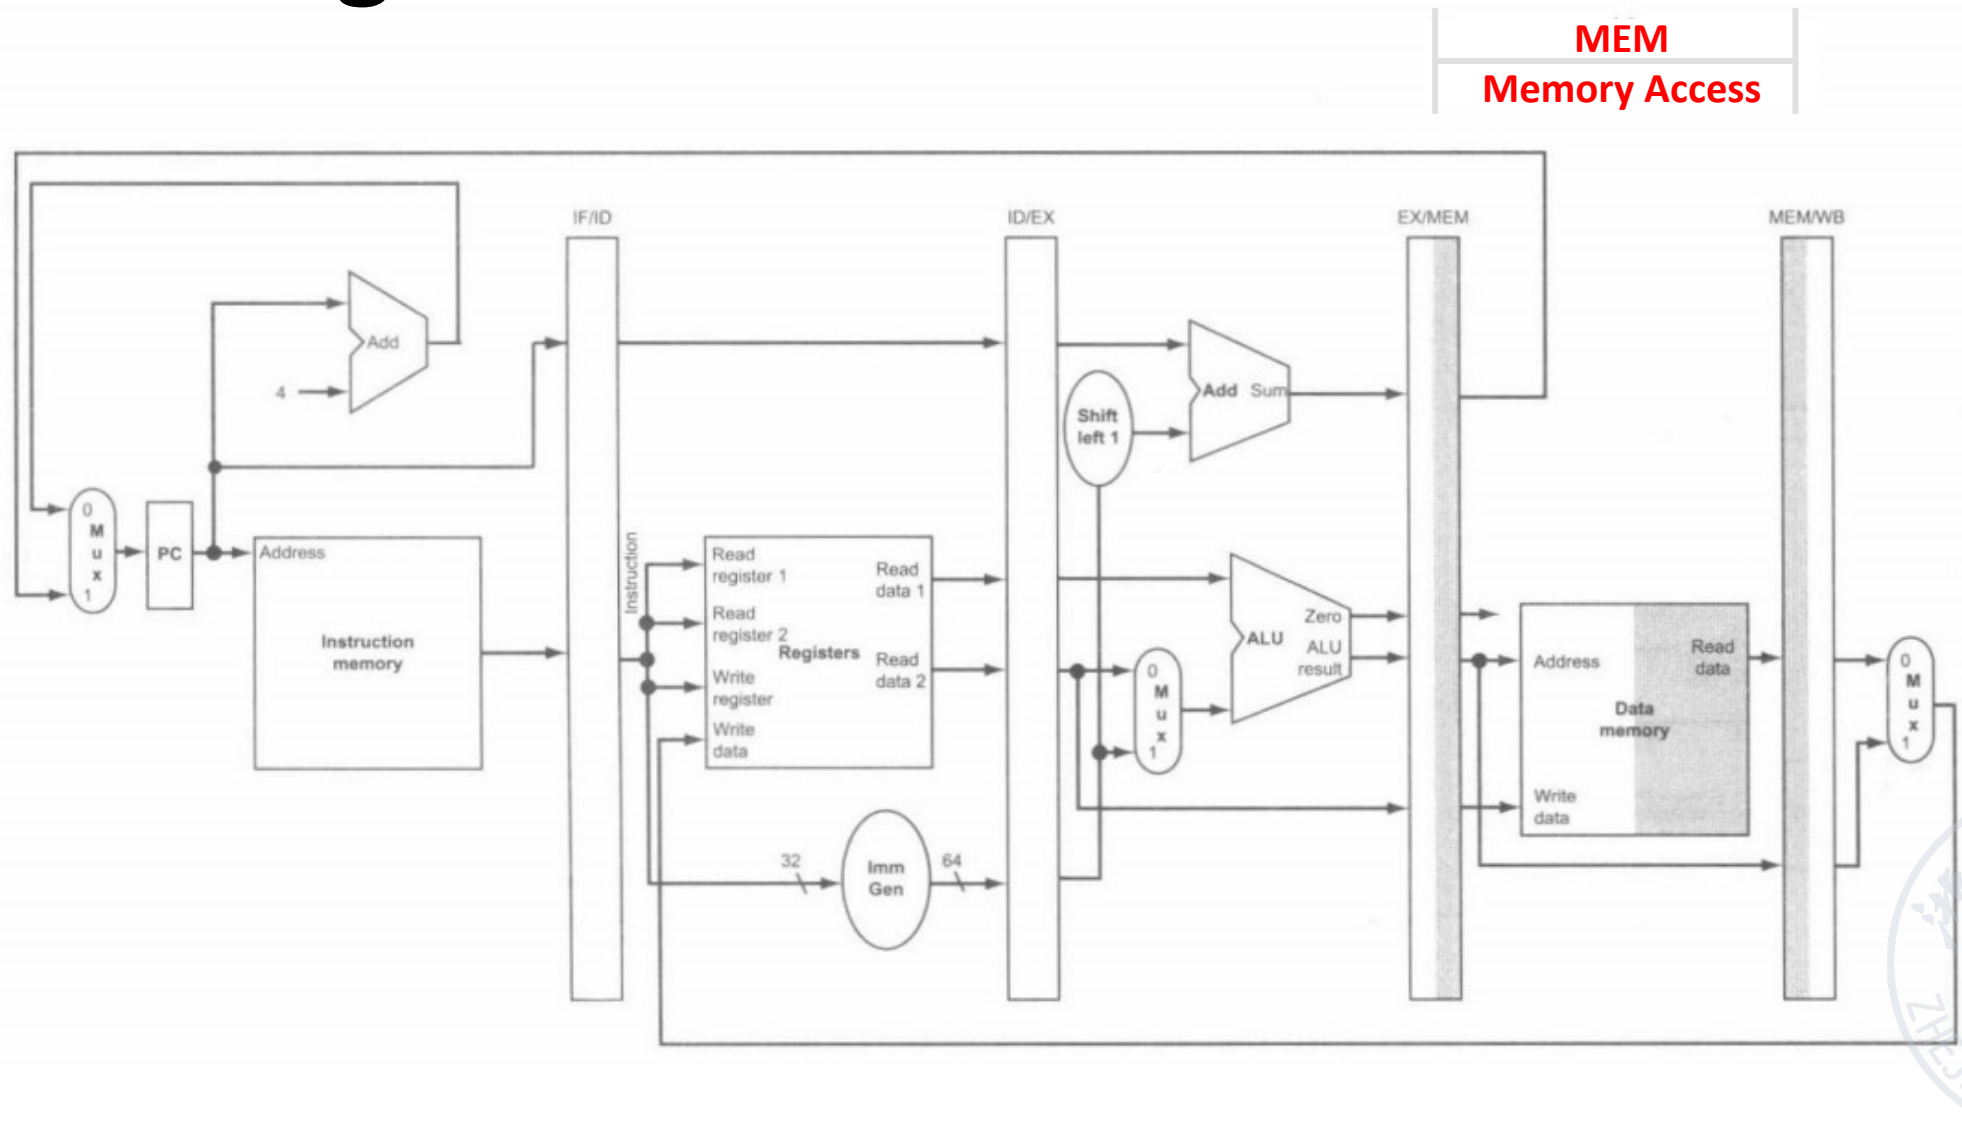


EX Stage of **Load** Instruction

EX
Execution

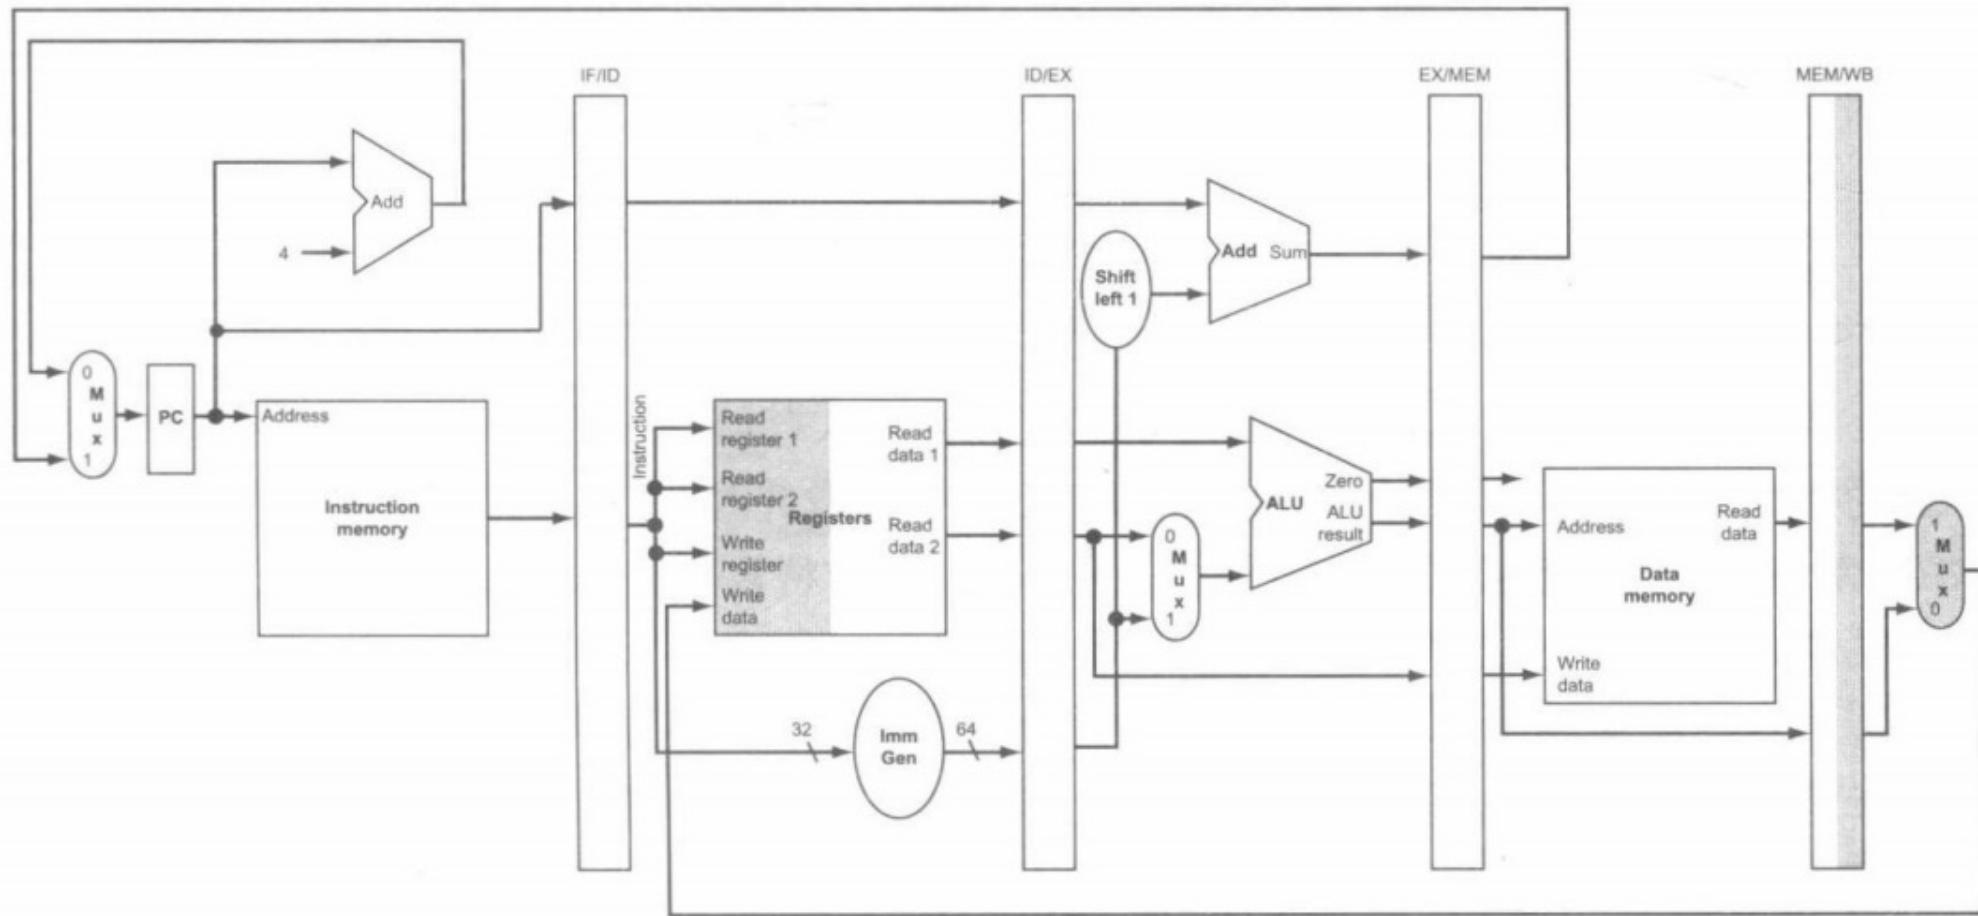


MEM Stage of Load Instruction



WB Stage of **Load** Instruction

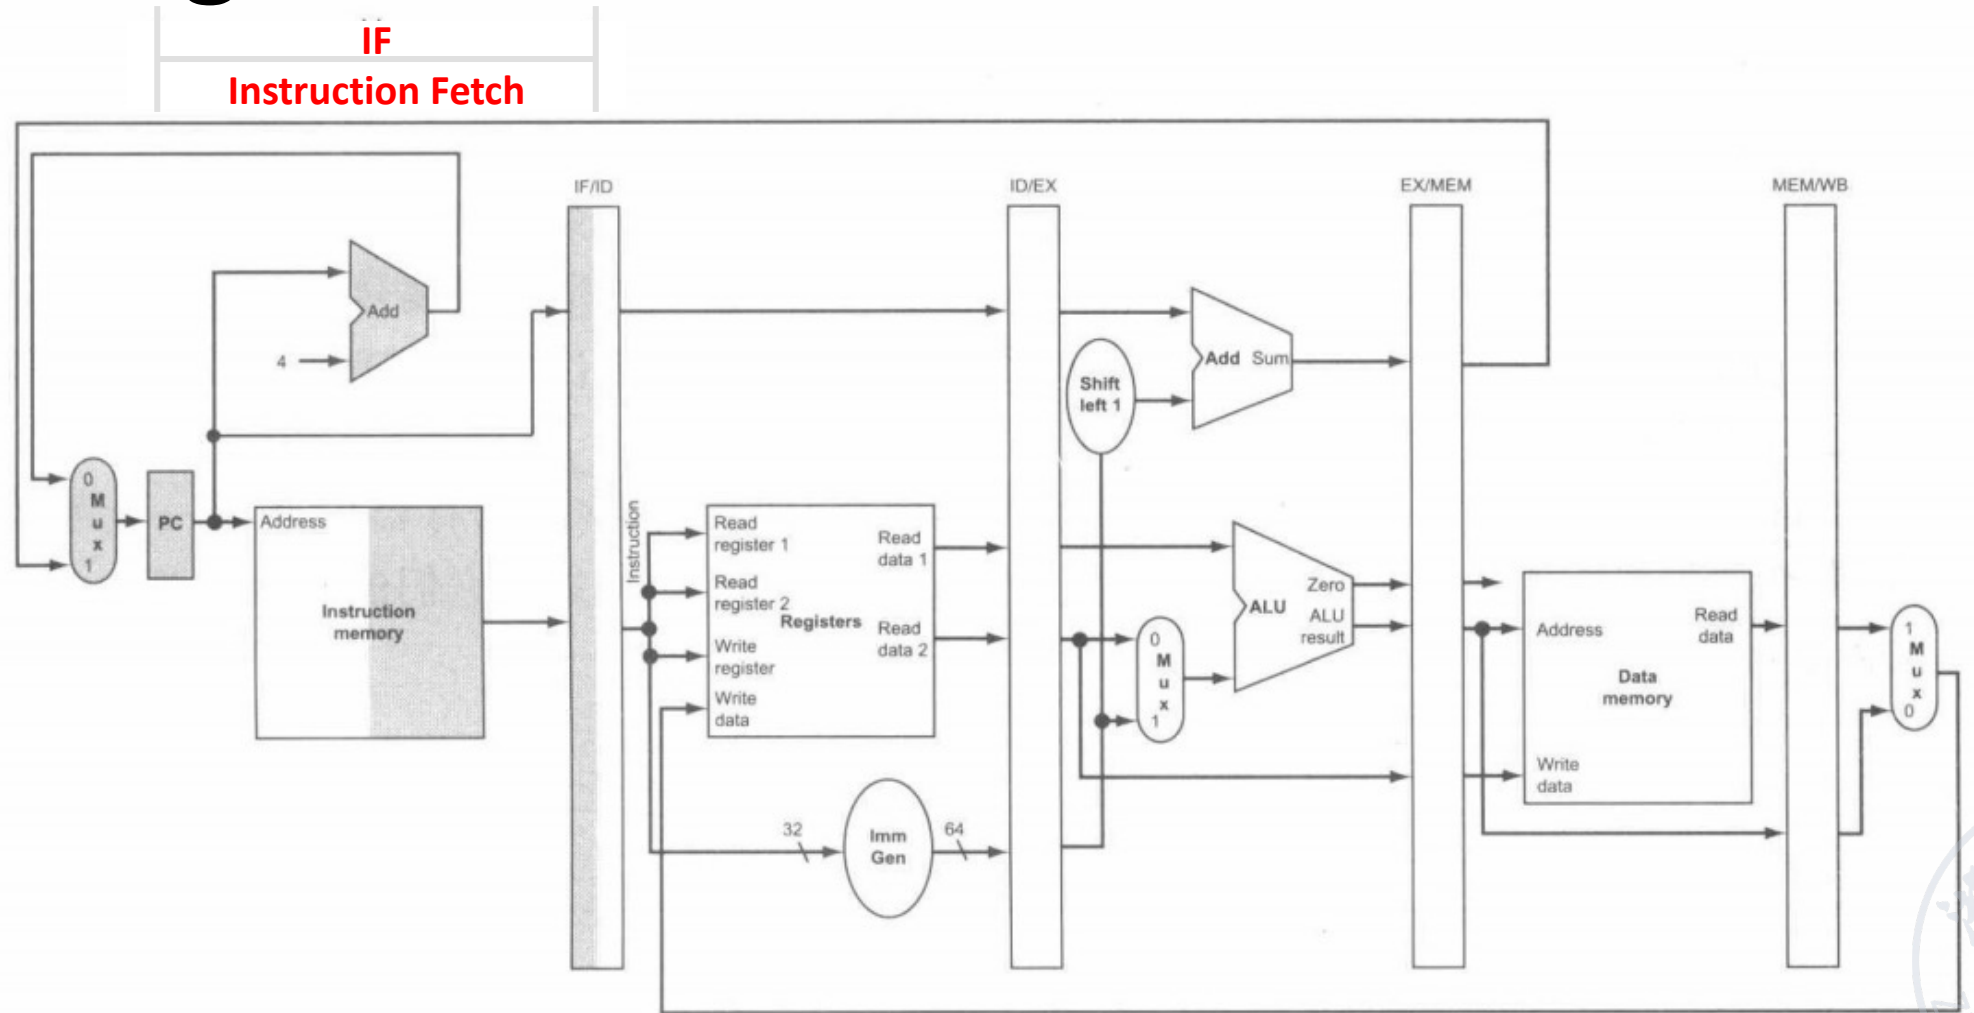
WB
Write Back



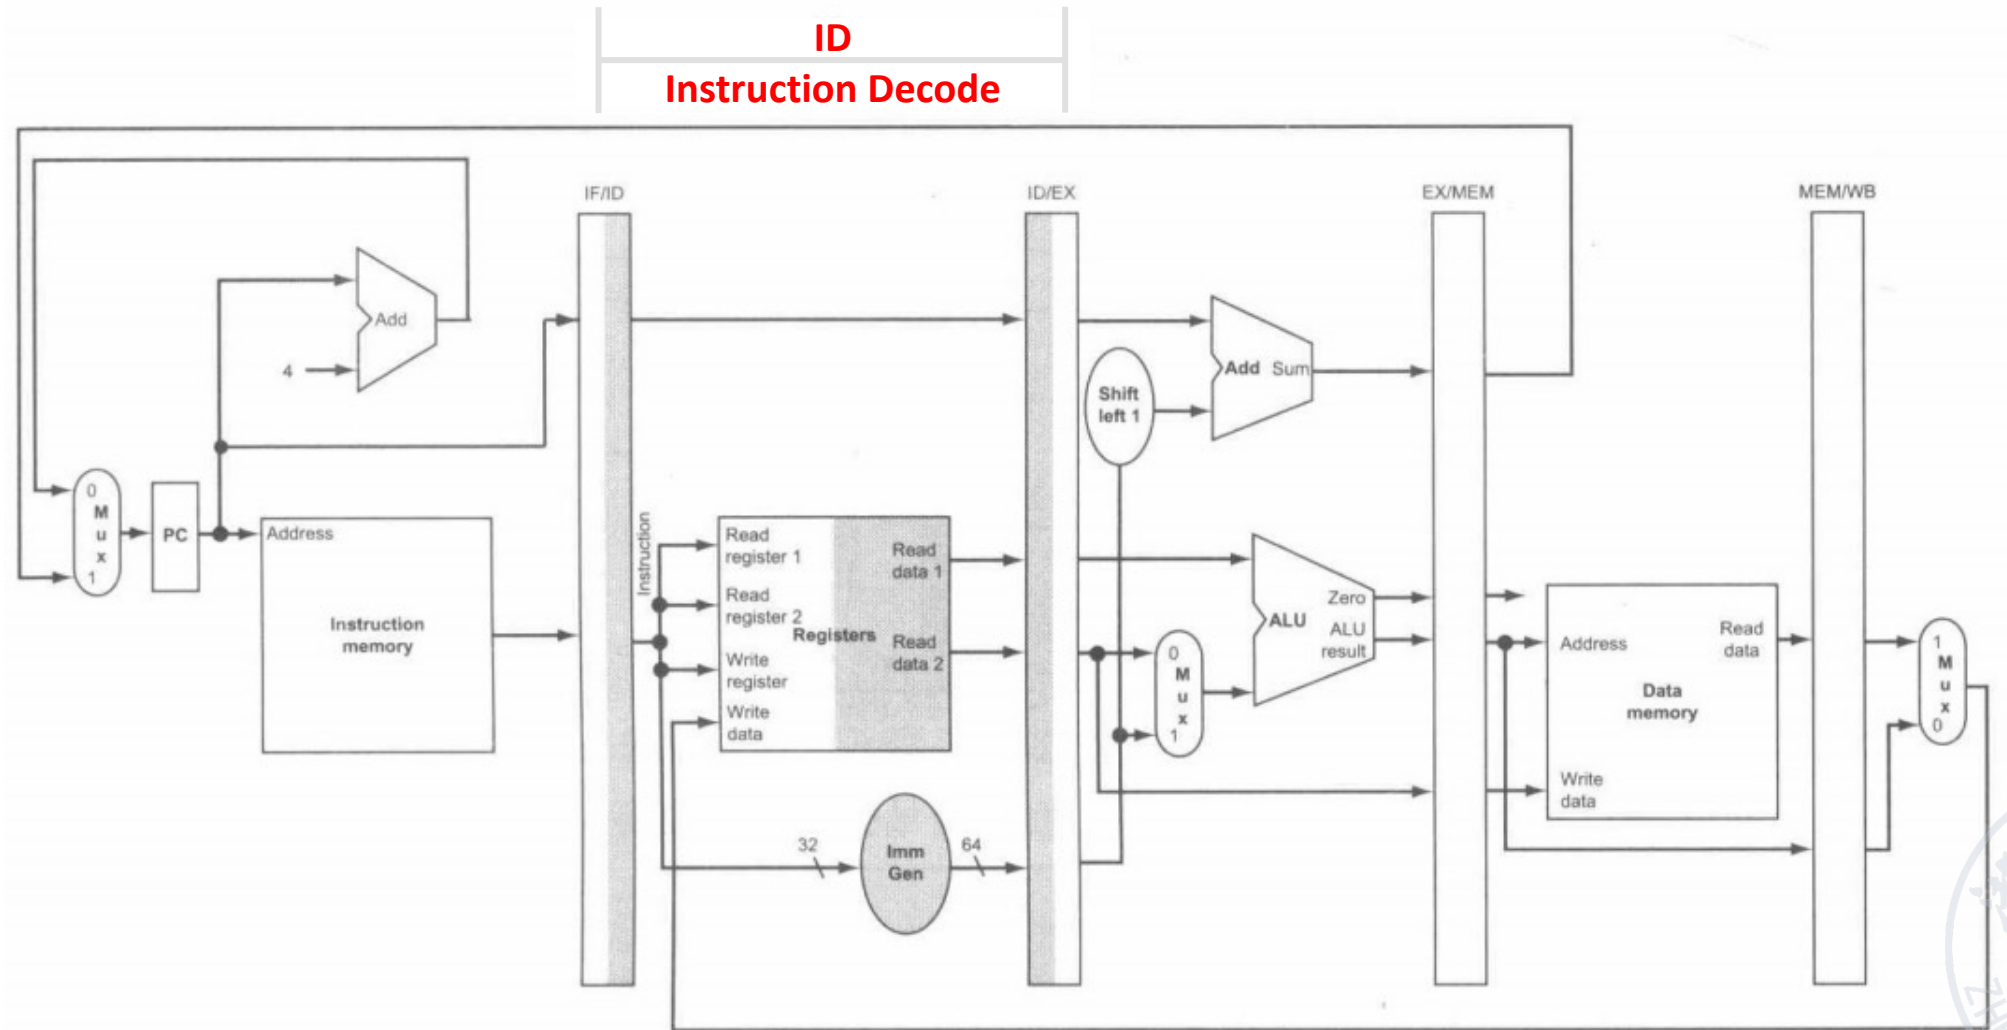
Store Instruction



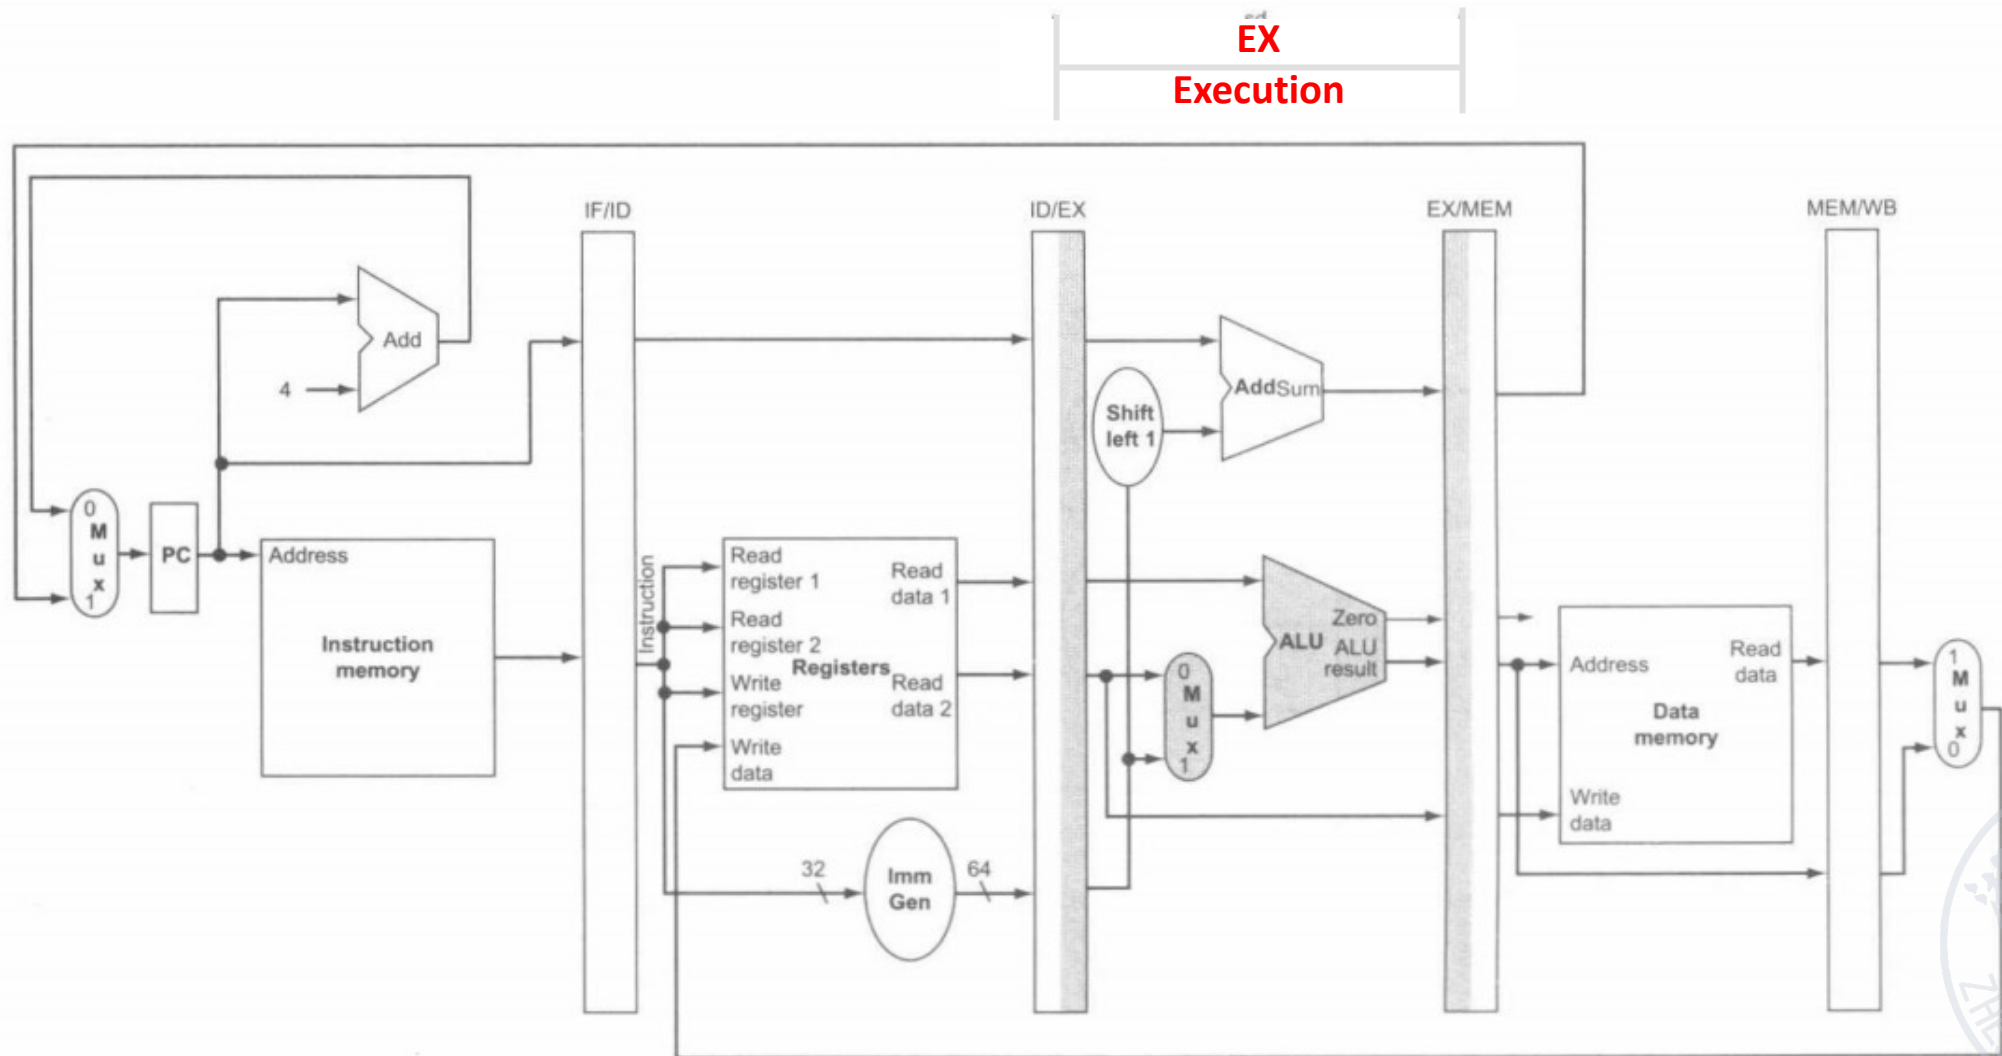
IF Stage of **Store** Instruction



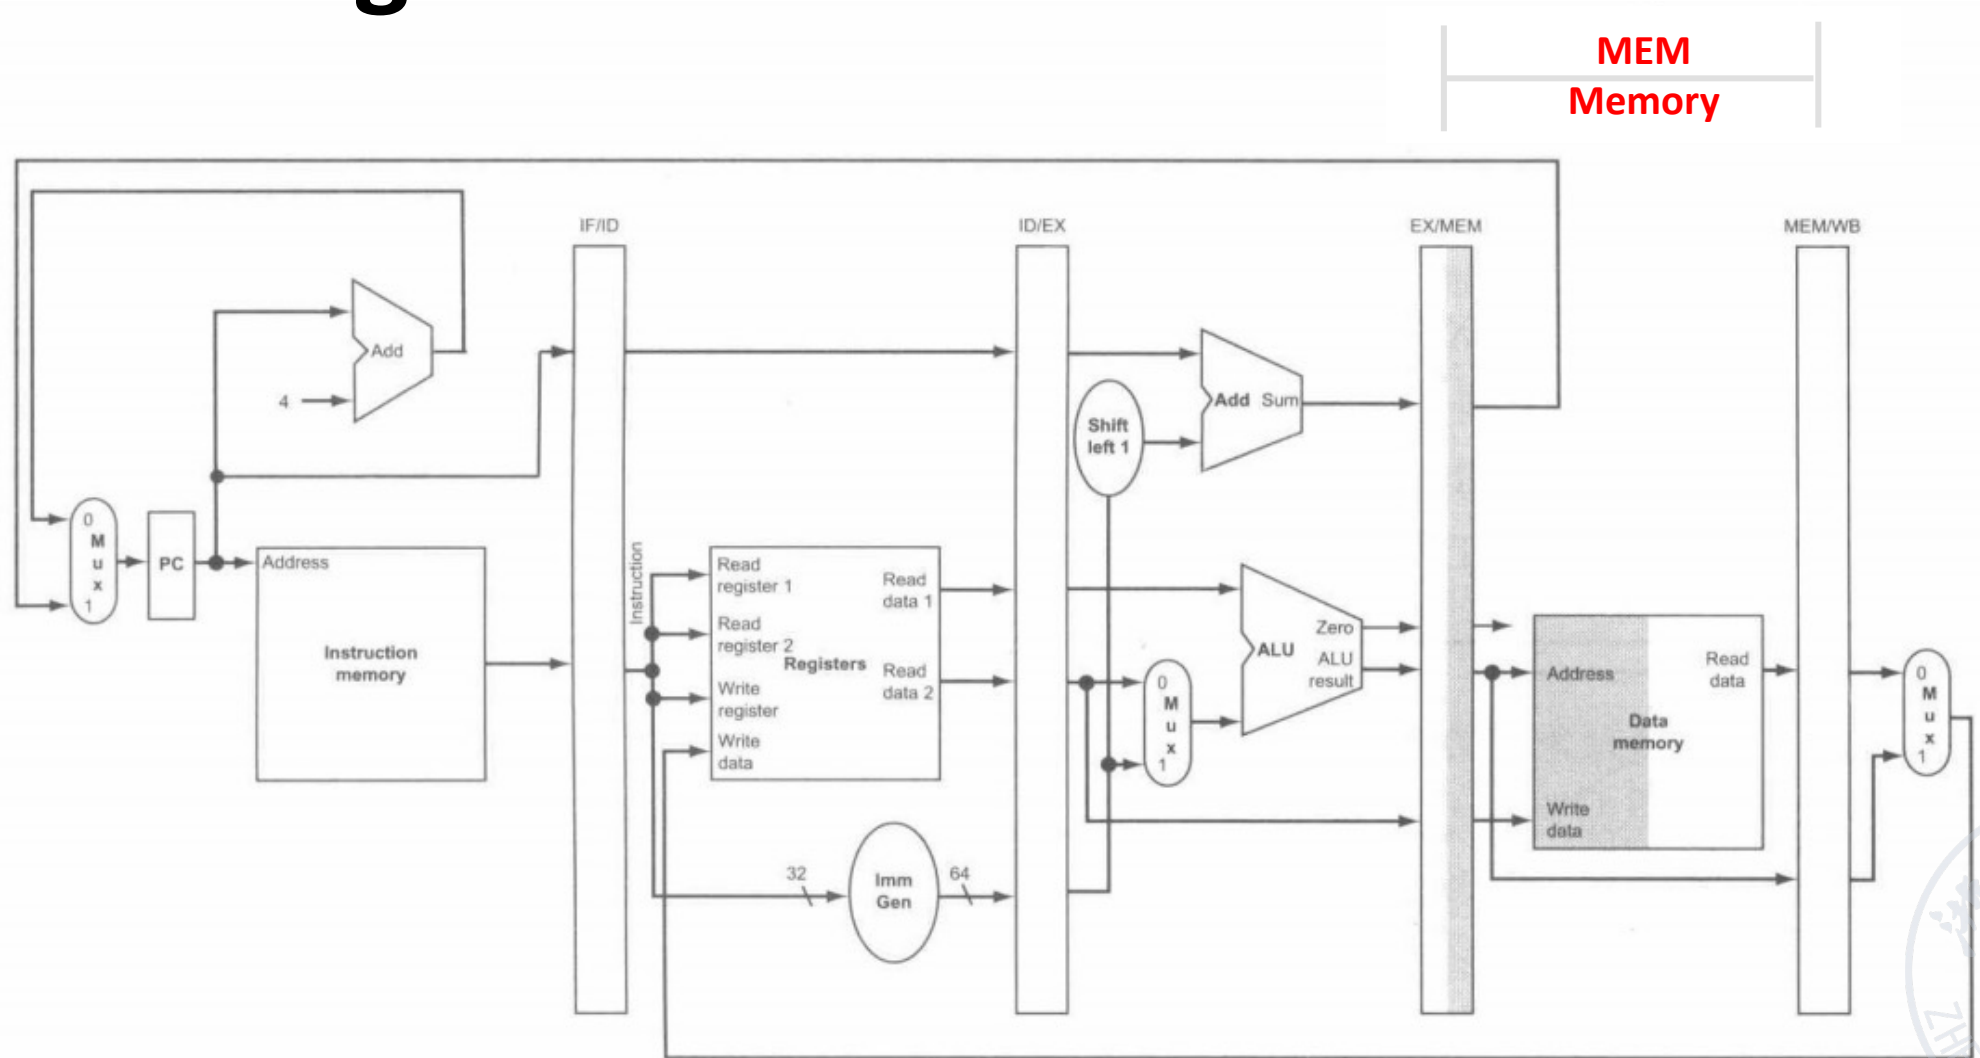
ID Stage of **Store** Instruction



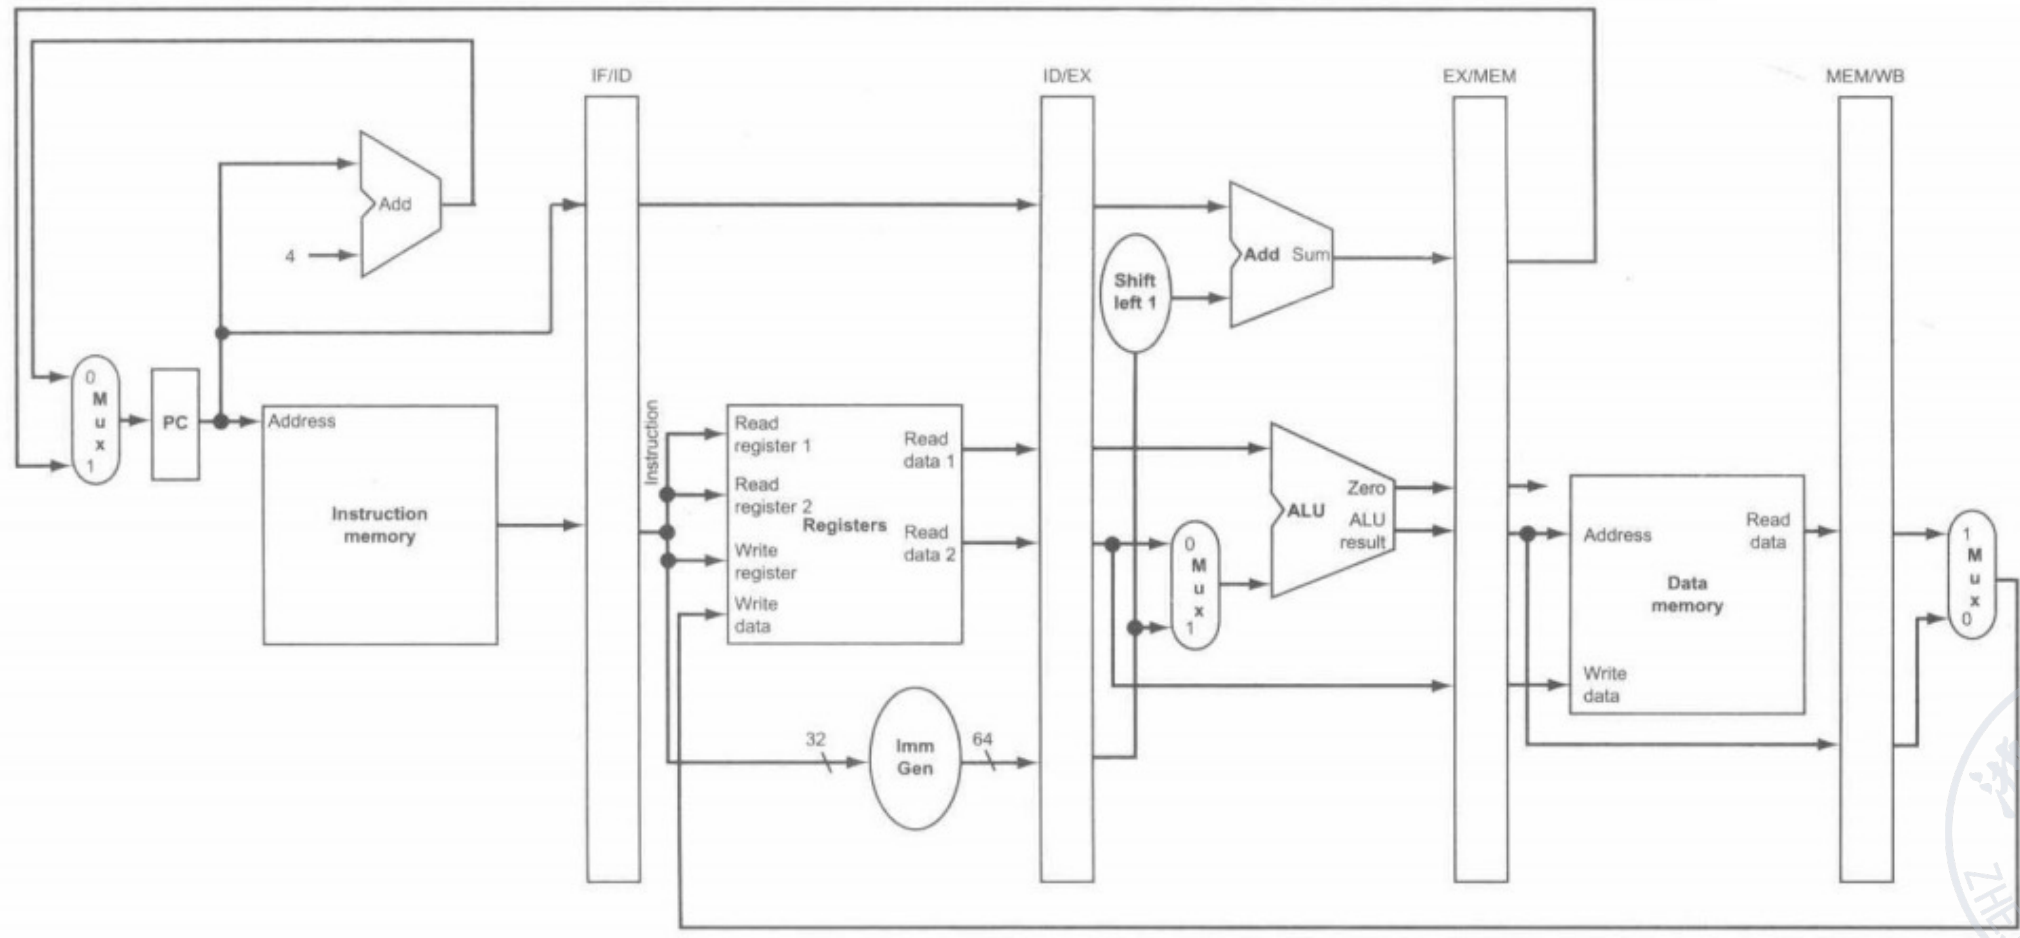
EX Stage of **Store** Instruction



MEM Stage of **Store** Instruction



WB Stage of **Store** Instruction



Review of **Store** Instruction

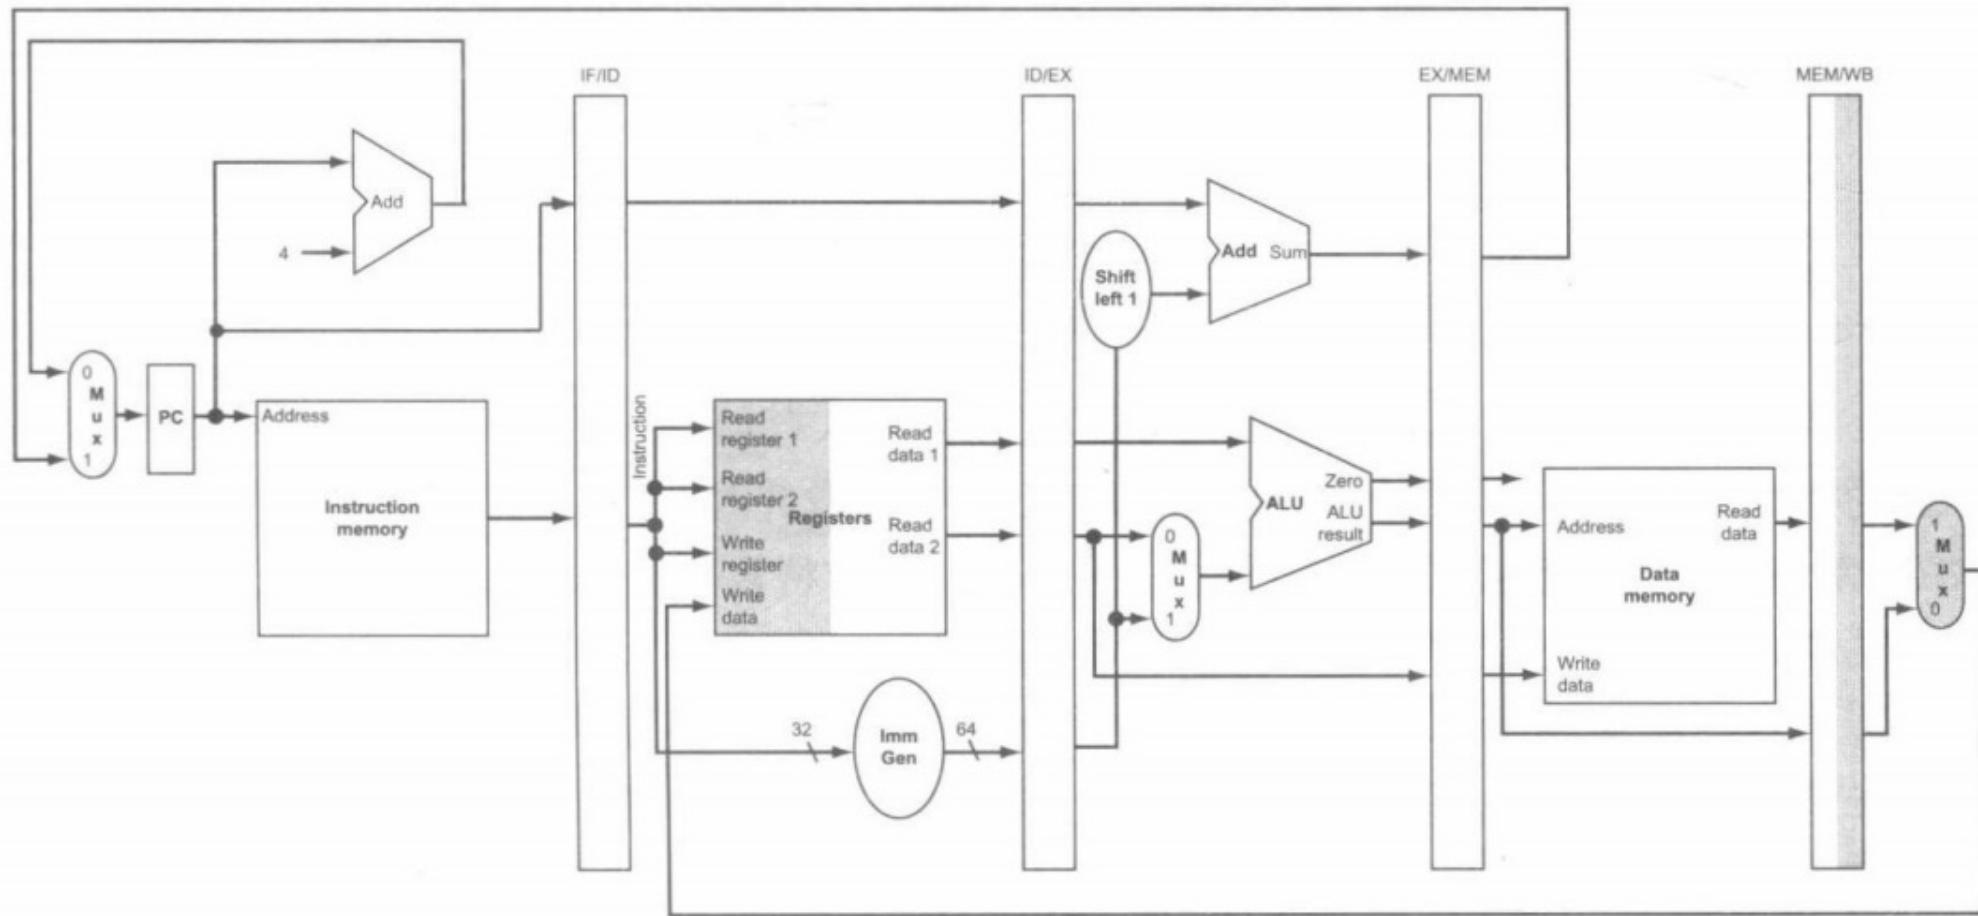


To pass something from an early pipe stage to a later pipe stage, the information must be placed in a pipeline register



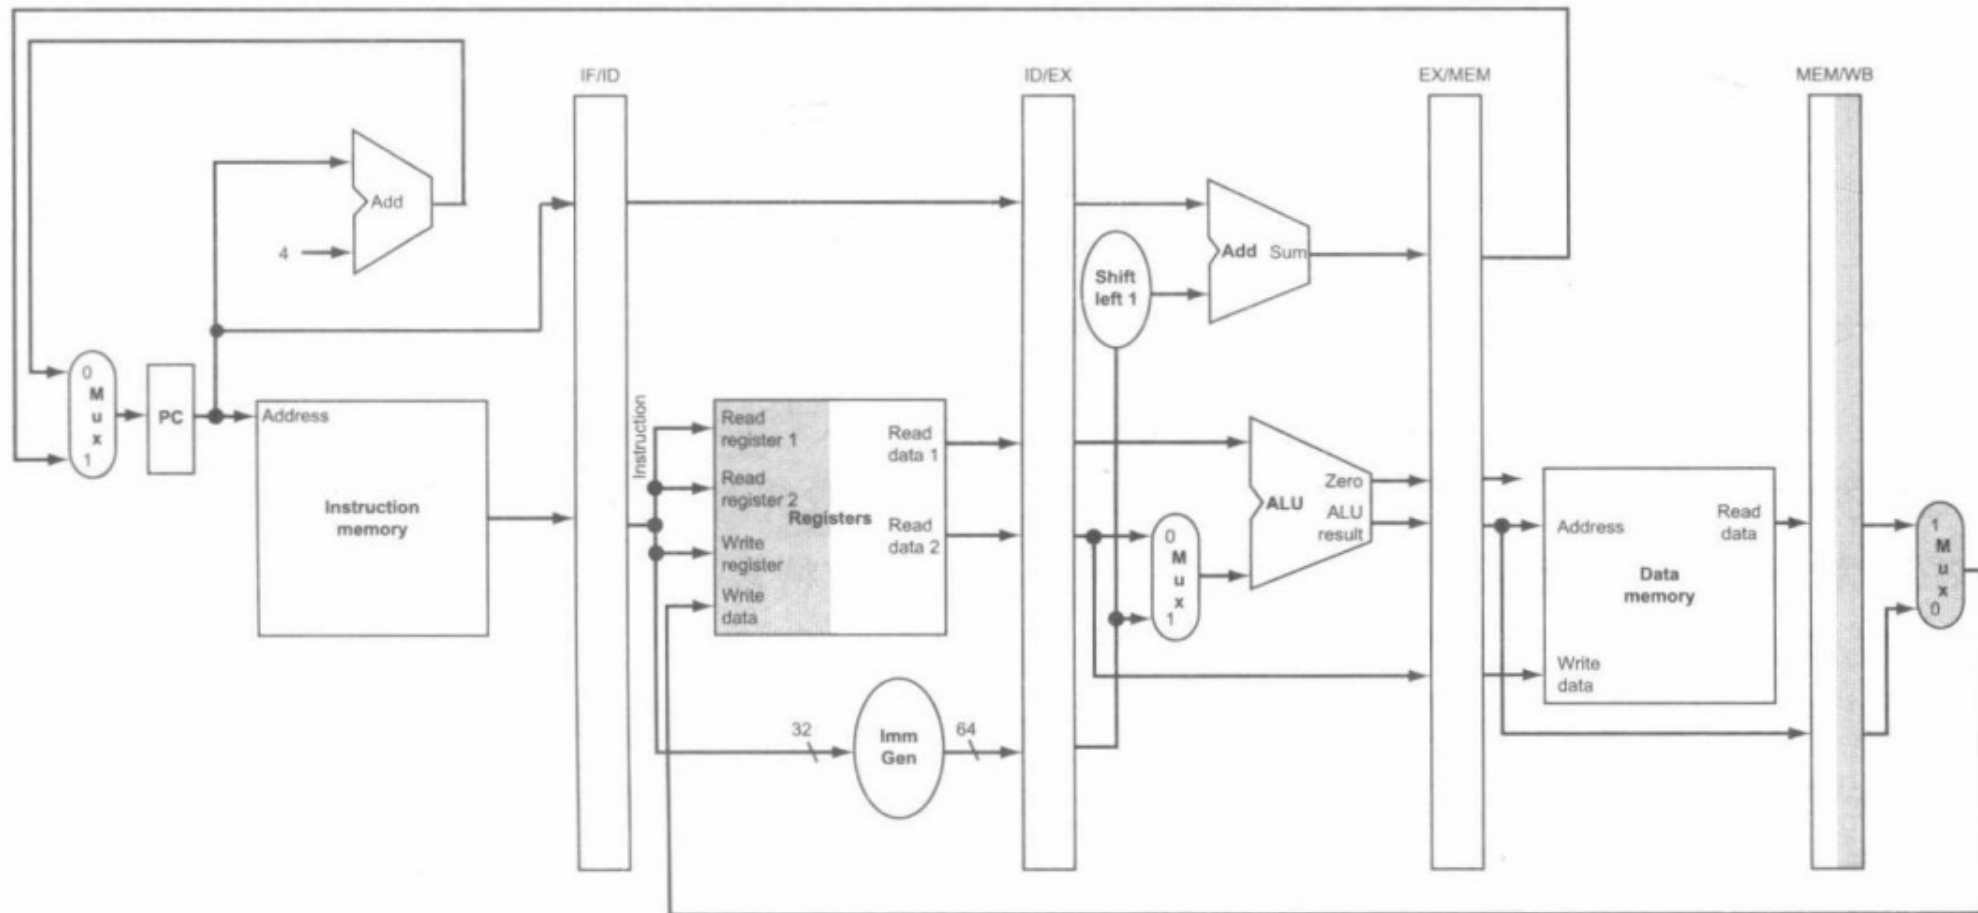
Review of **Load** Instruction in WB Stage

WB
Write Back



Review of **Load** Instruction in WB Stage

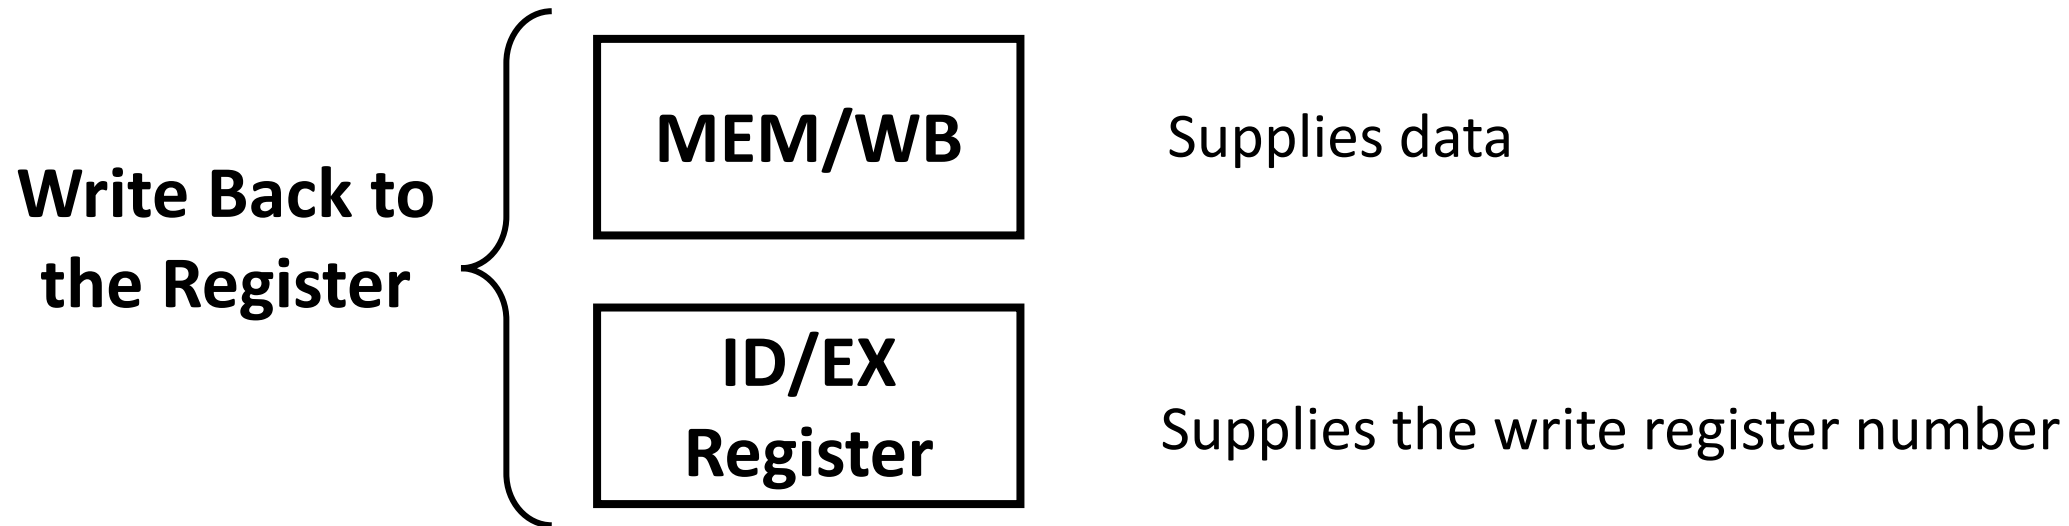
WB
Write Back



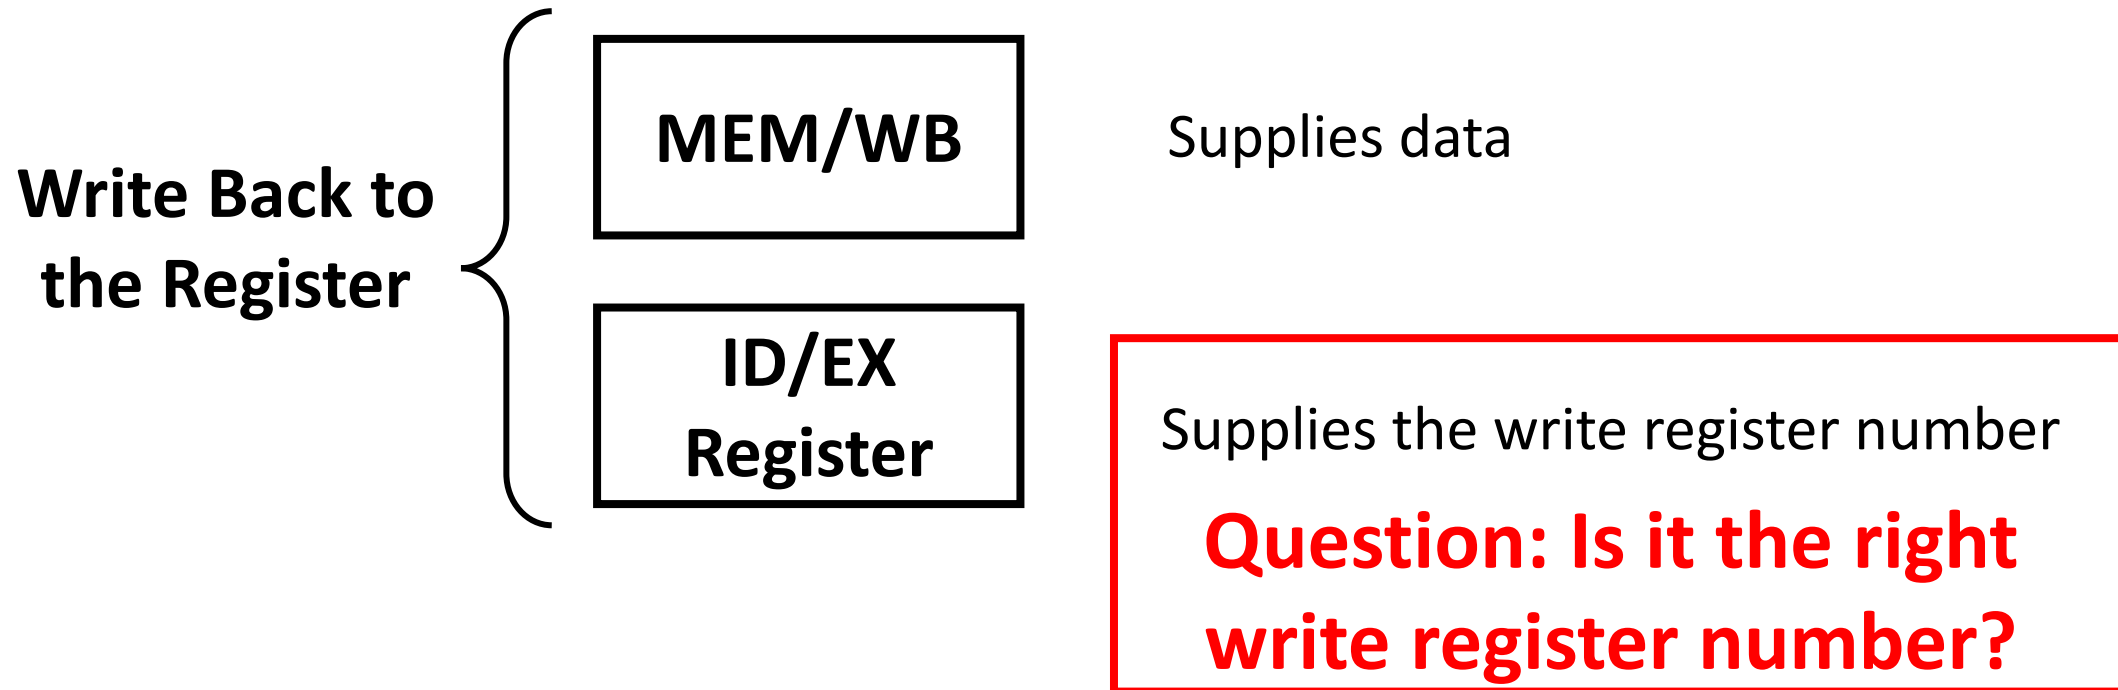
Let's uncover a bug in the design of the load instruction



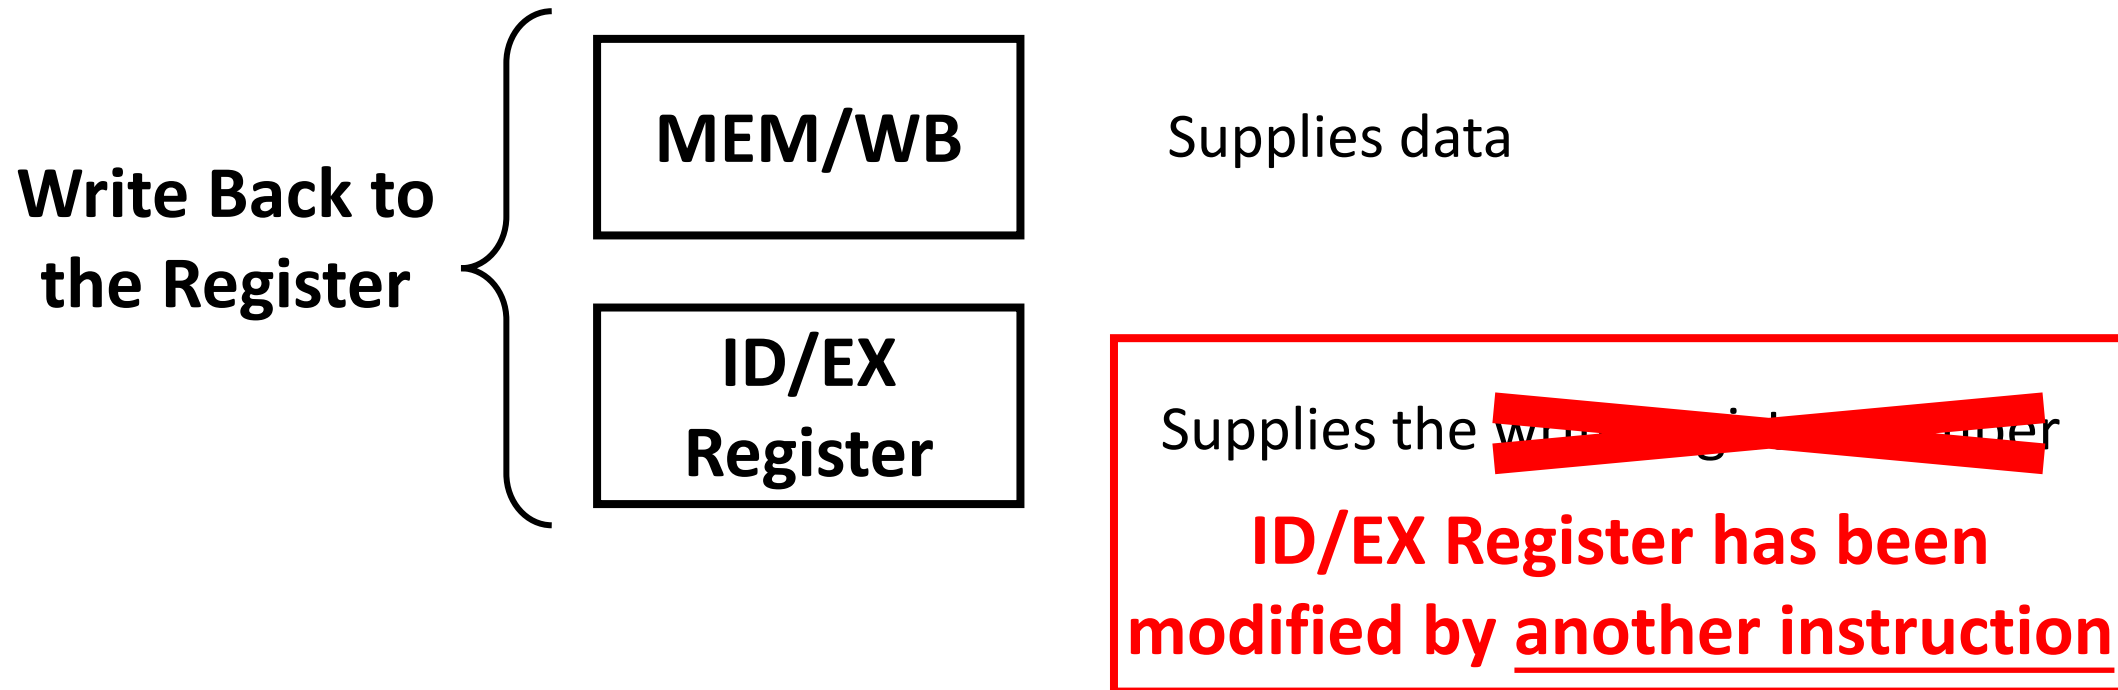
Review the **Load** Instruction in WB Stage



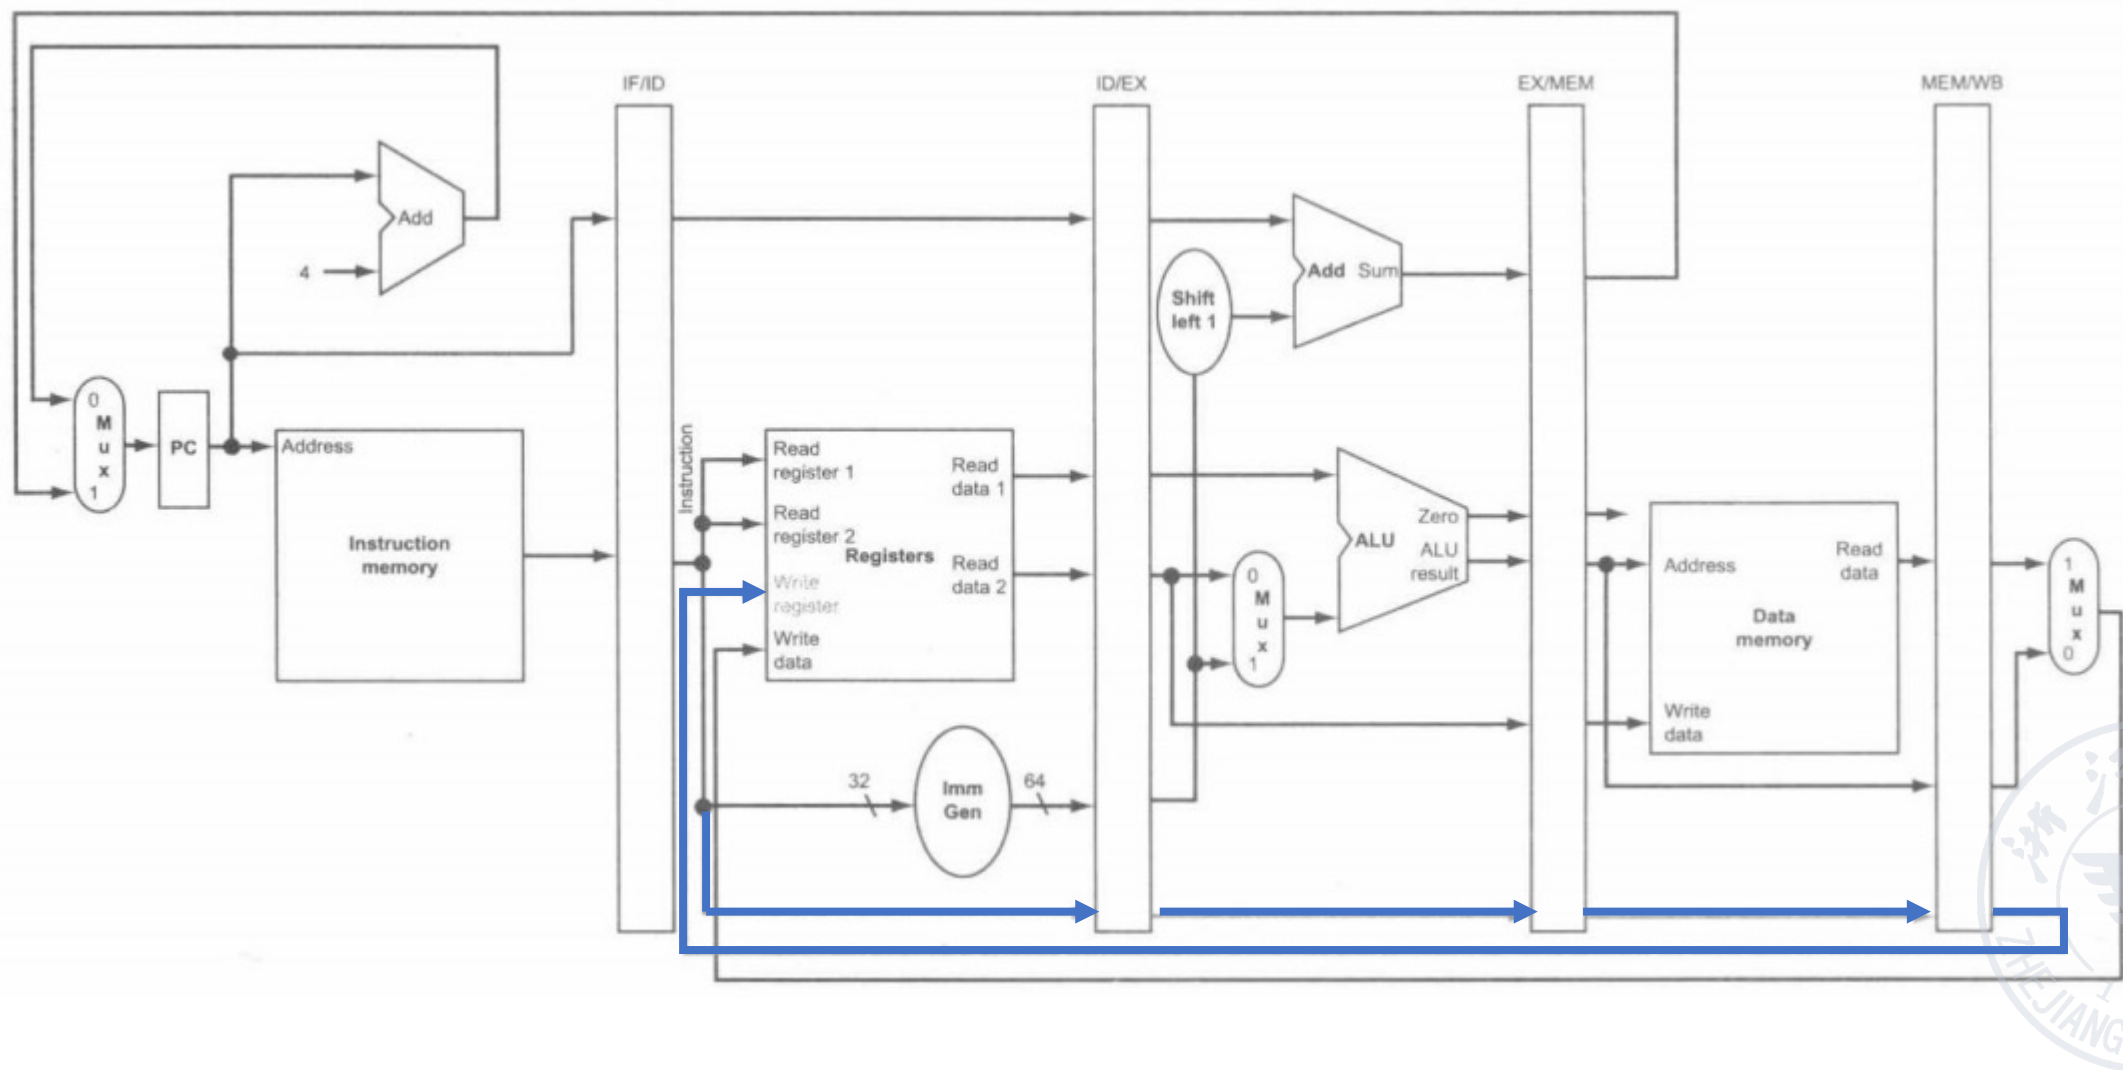
Review the **Load** Instruction in WB Stage



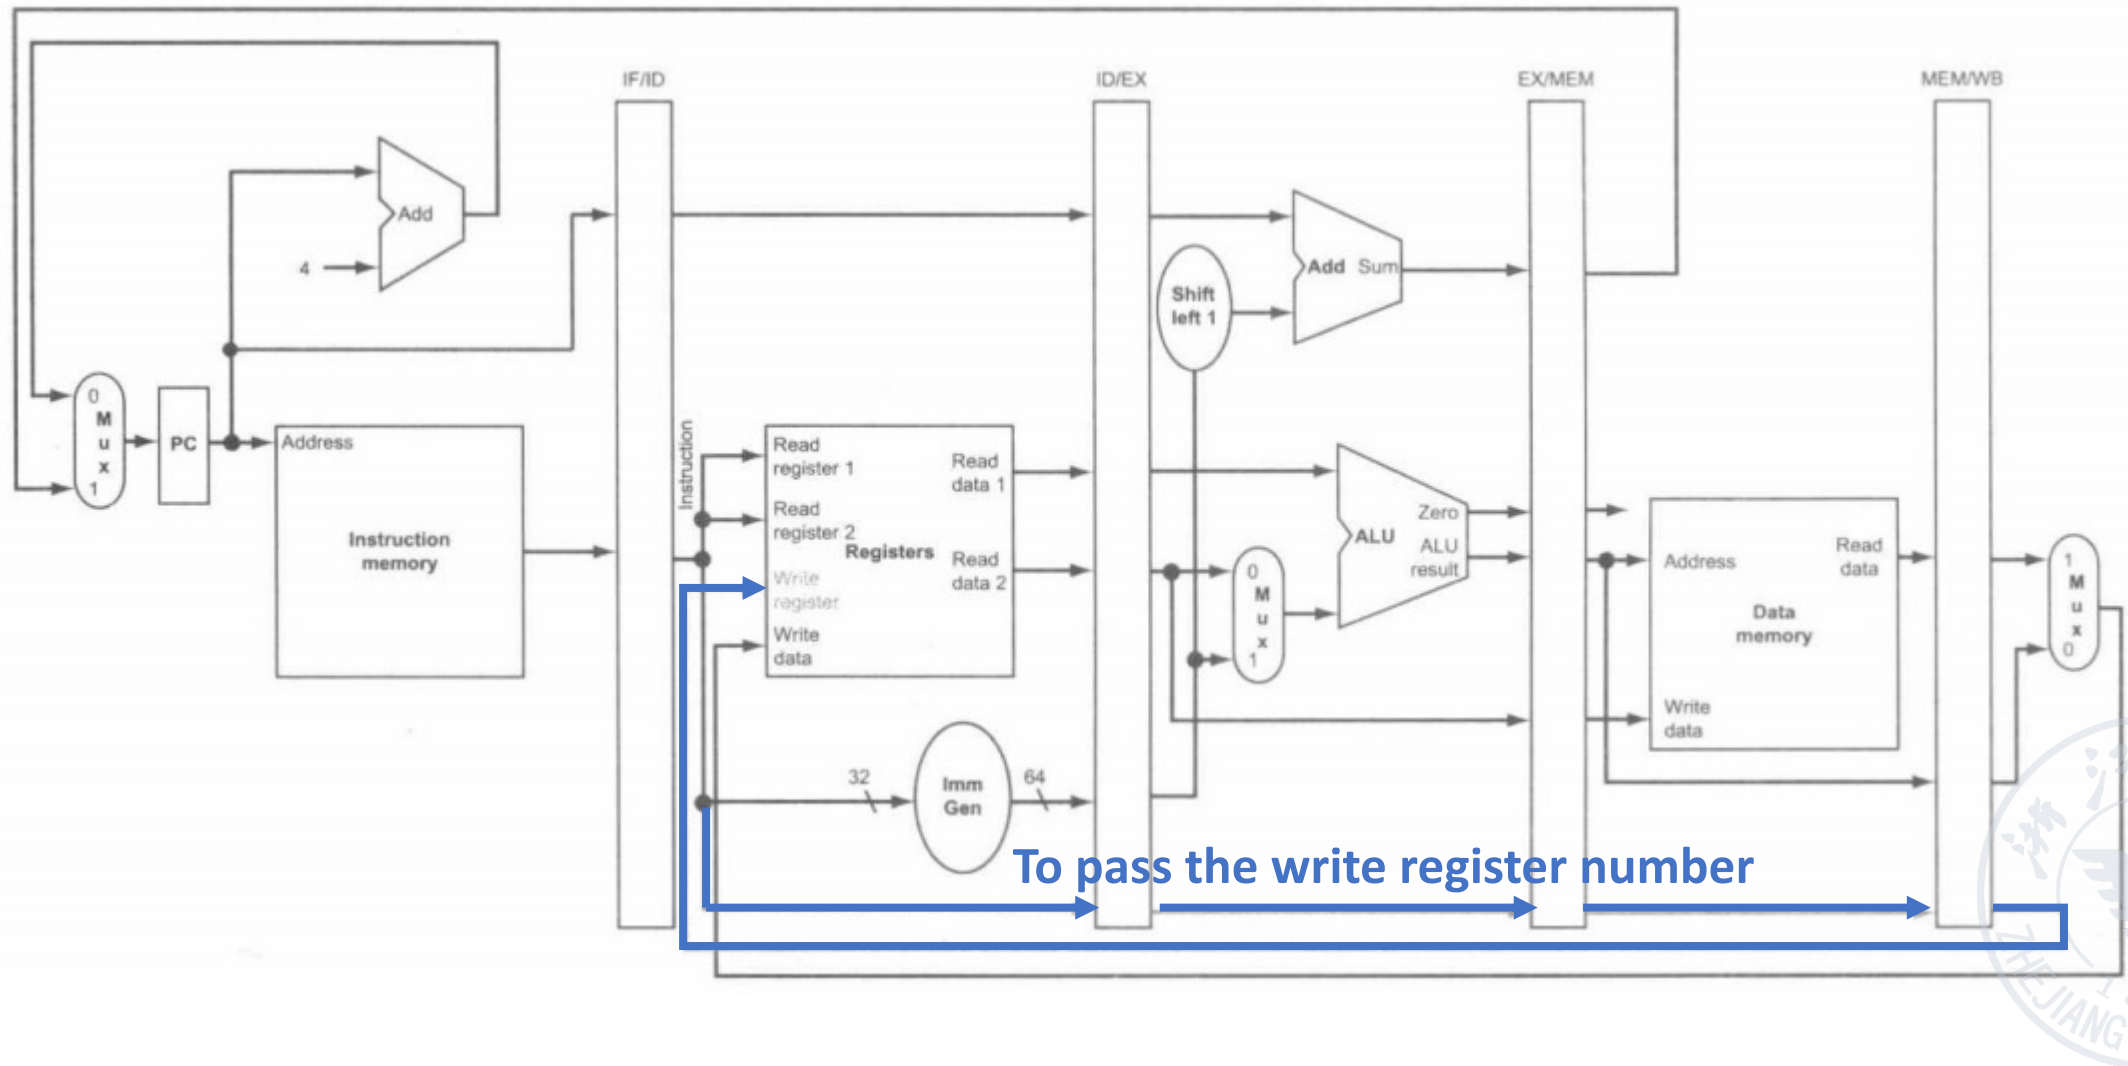
Review the **Load** Instruction in WB Stage



The Corrected Pipelined Datapath to Handle the **Load** Instruction Properly



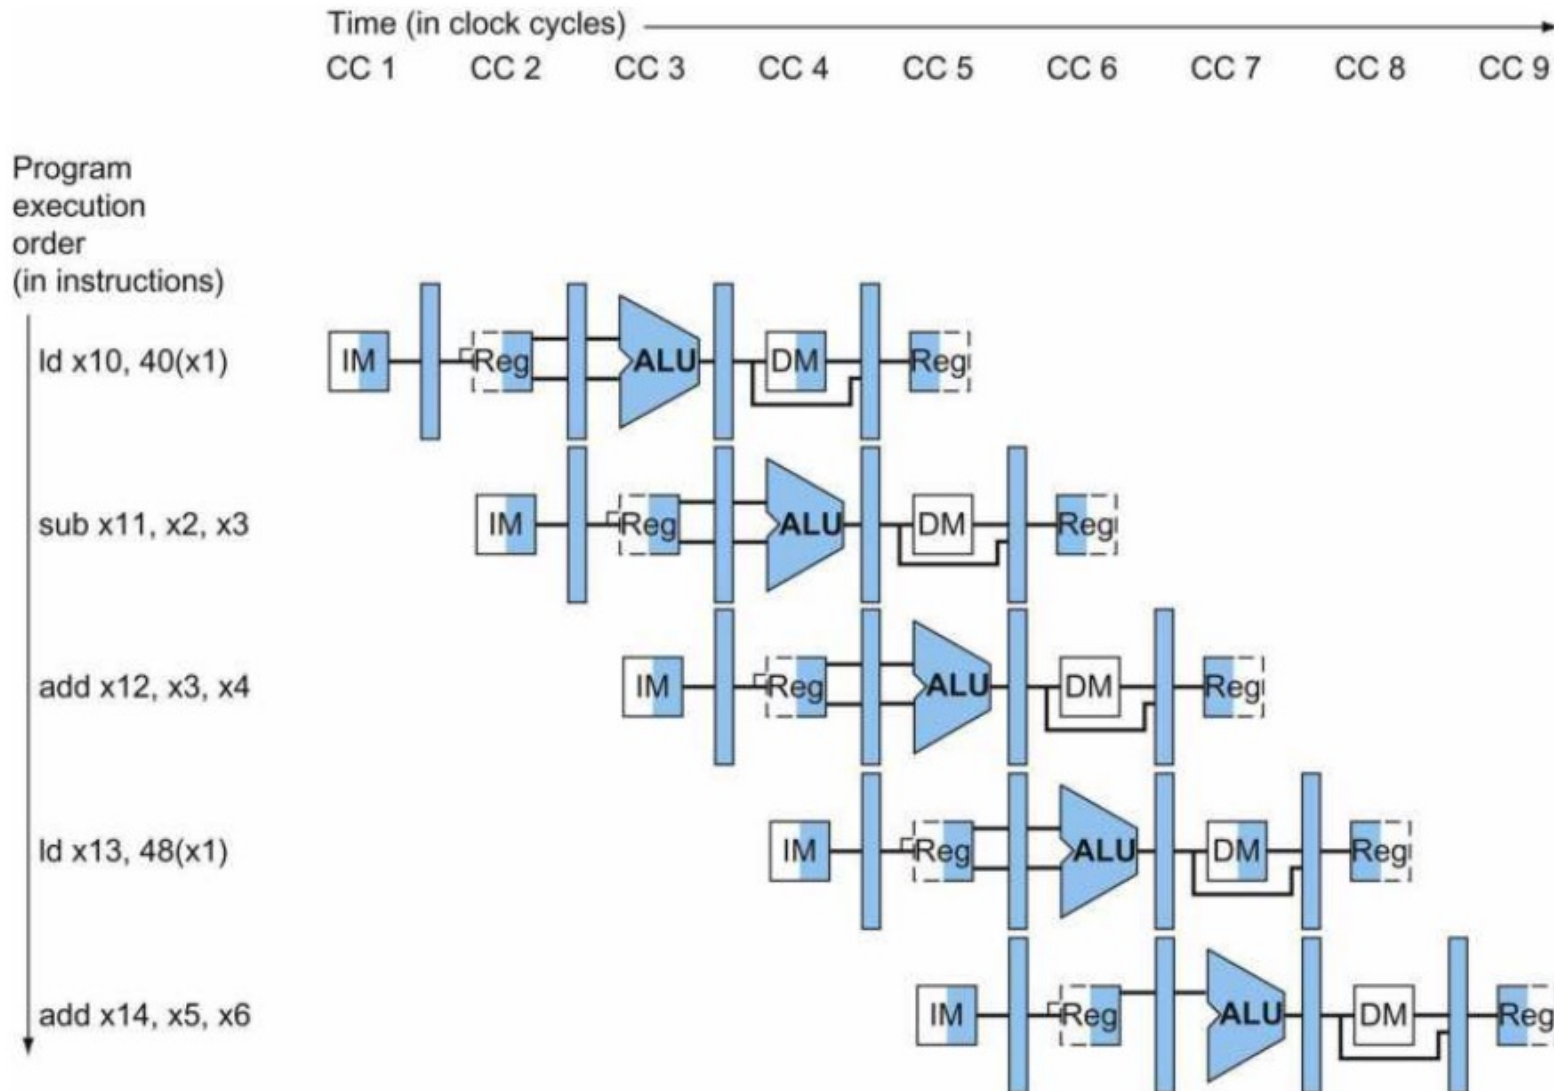
The Corrected Pipelined Datapath to Handle the **Load** Instruction Properly



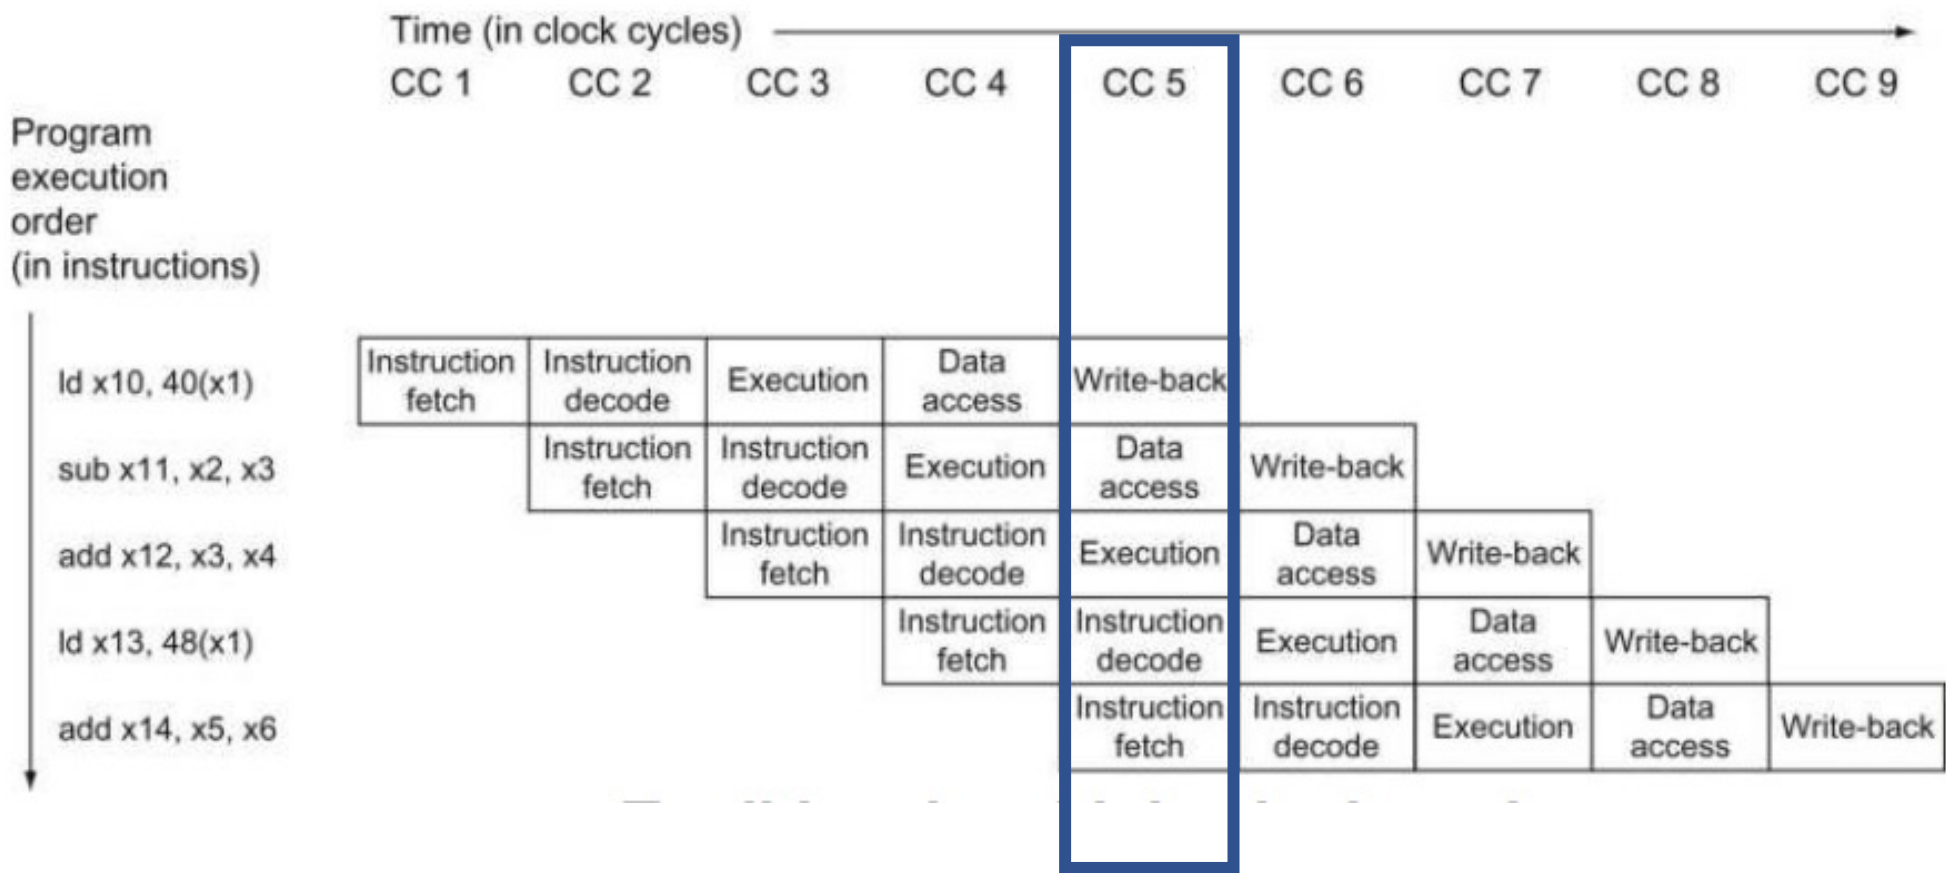
Graphically Representing Pipelines



Multiple-Clock-Cycle Pipeline Diagram of Five Instructions



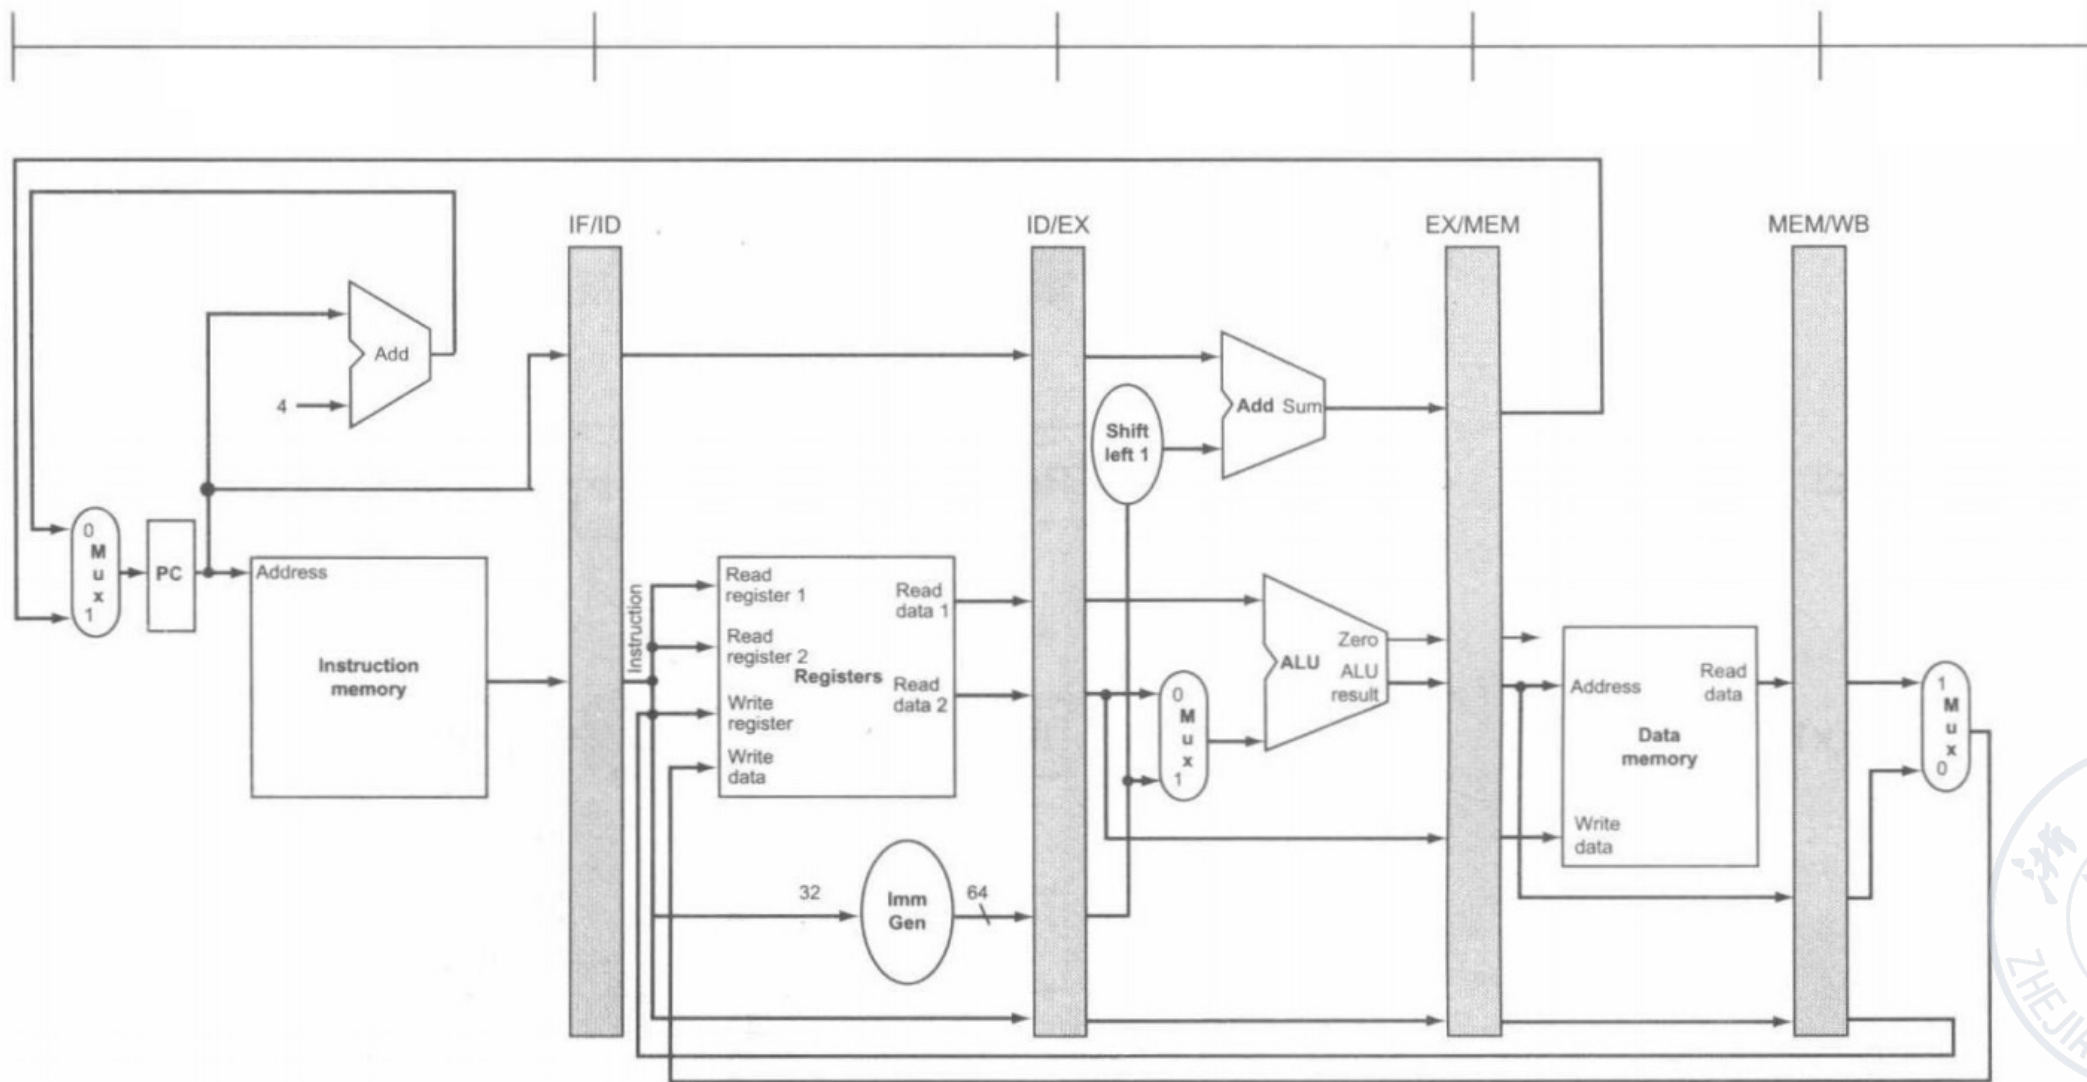
Traditional Multiple-Clock-Cycle Pipeline Diagram



Now draw the Single-clock-cycle pipeline diagram at CC5




Single-Clock-Cycle Pipeline Diagram at CC5

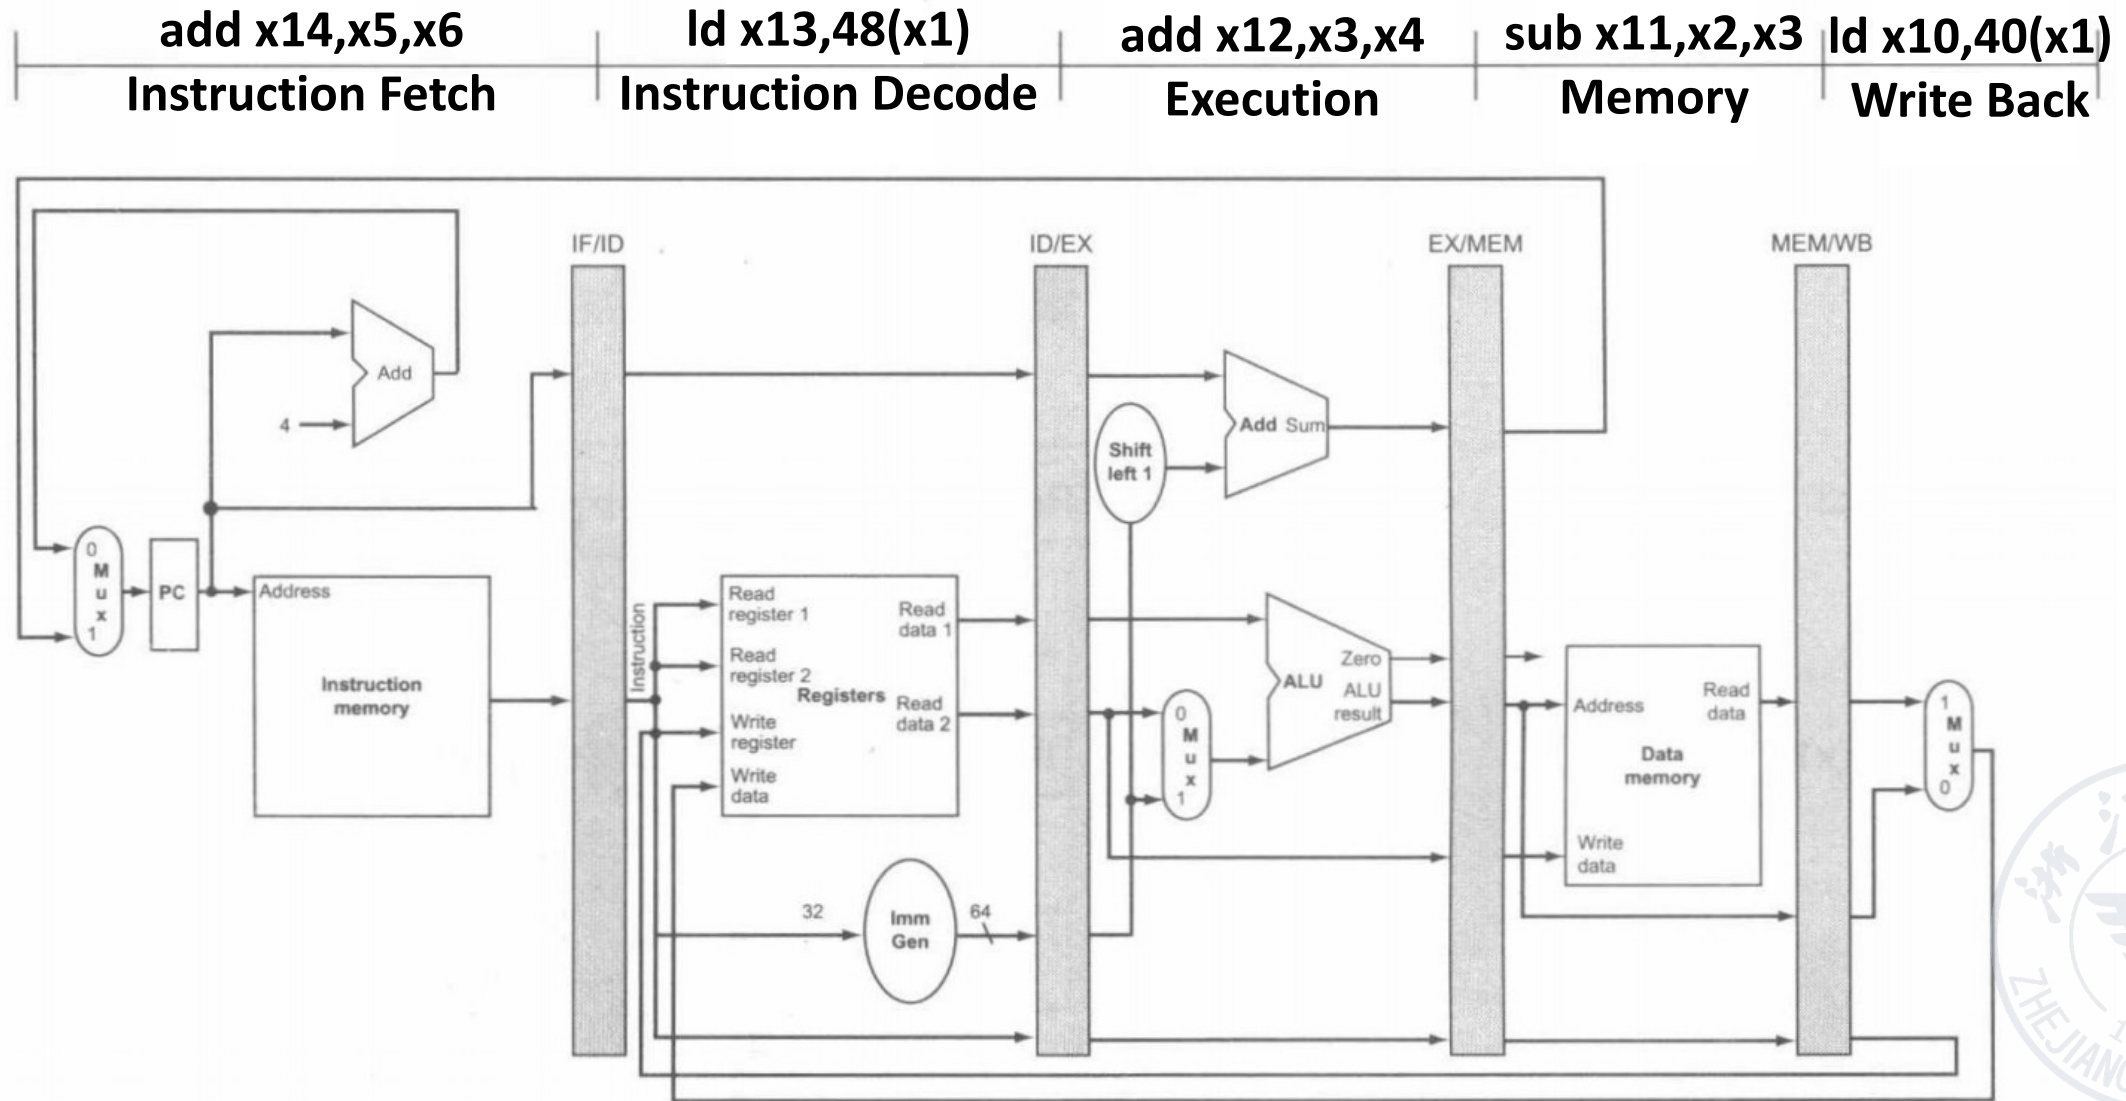


Single-Clock-Cycle Pipeline Diagram at CC5

Question: Write name and stage of instructions



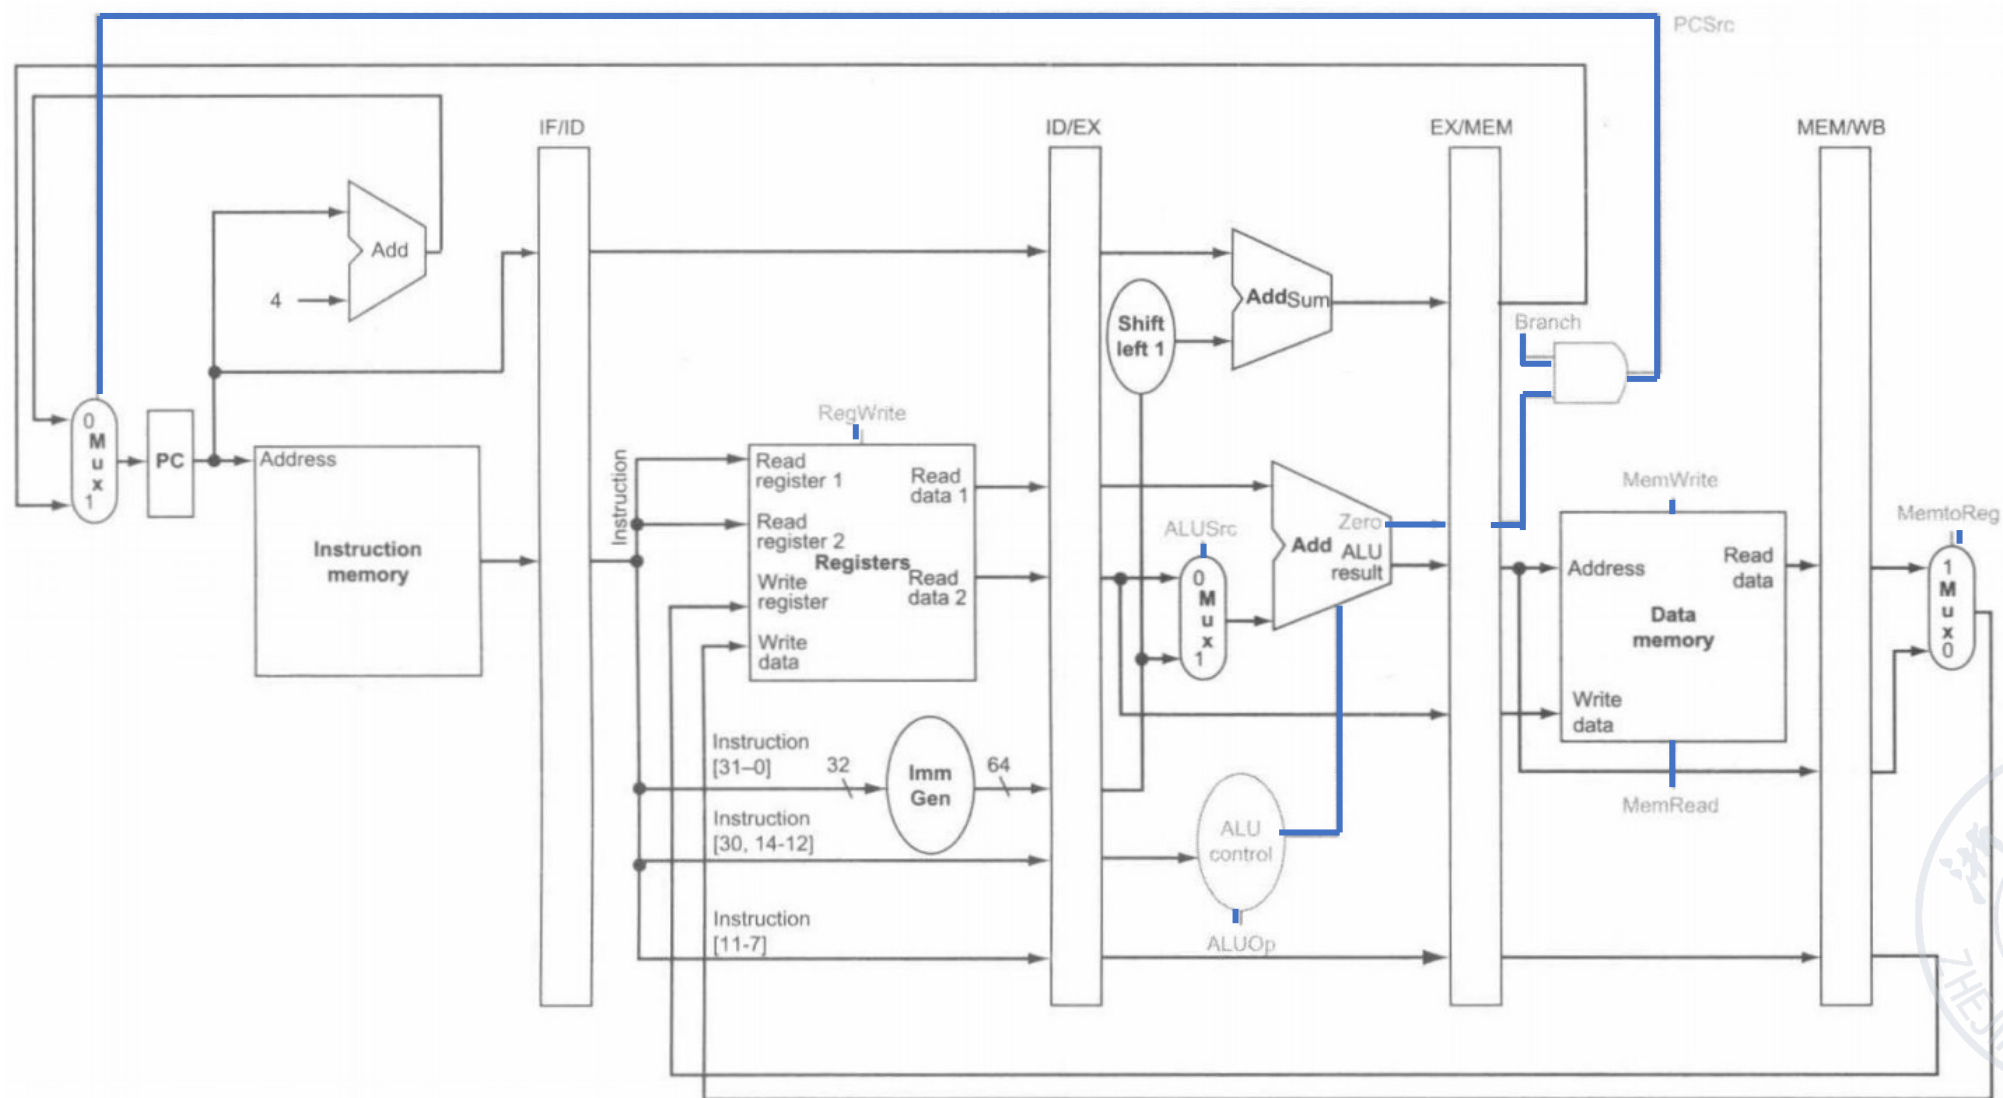
Single-Clock-Cycle Pipeline Diagram at CC5



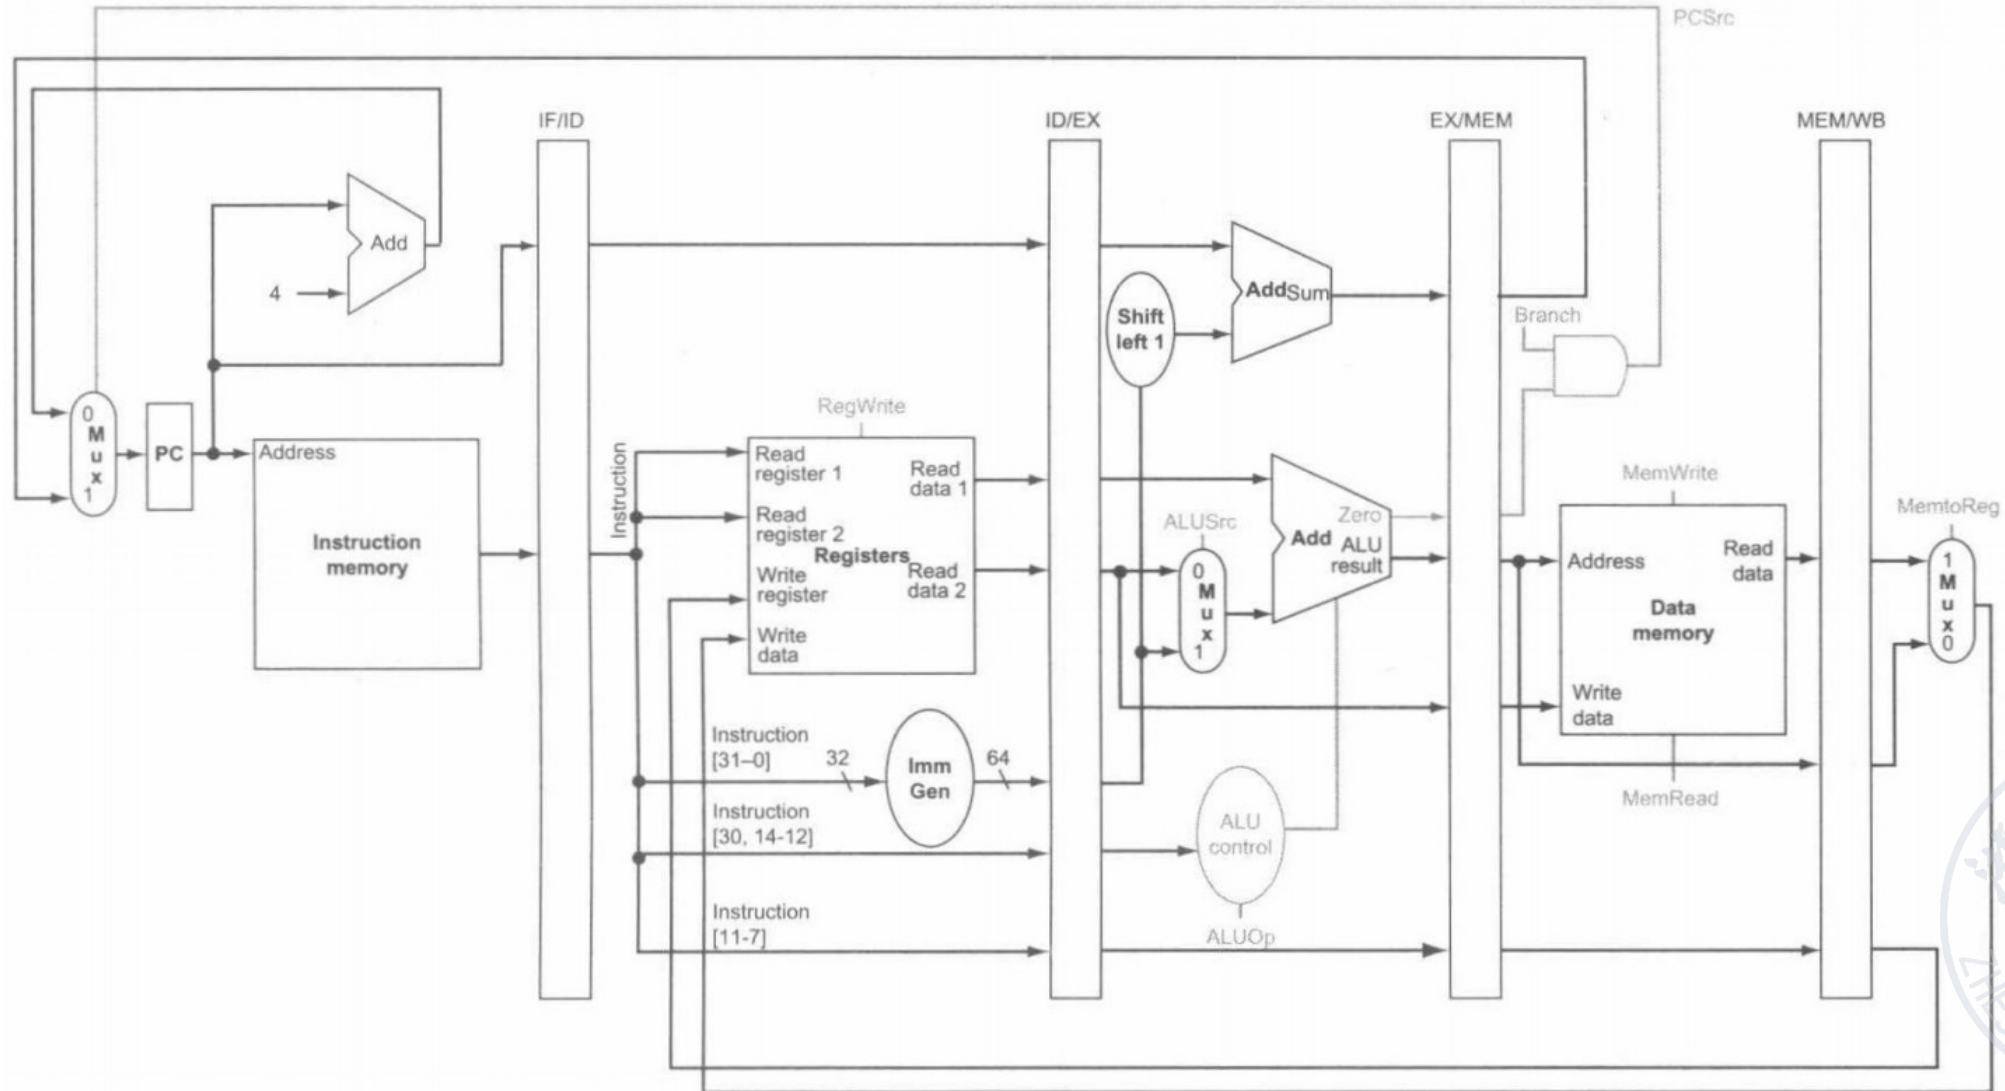
Pipelined Control



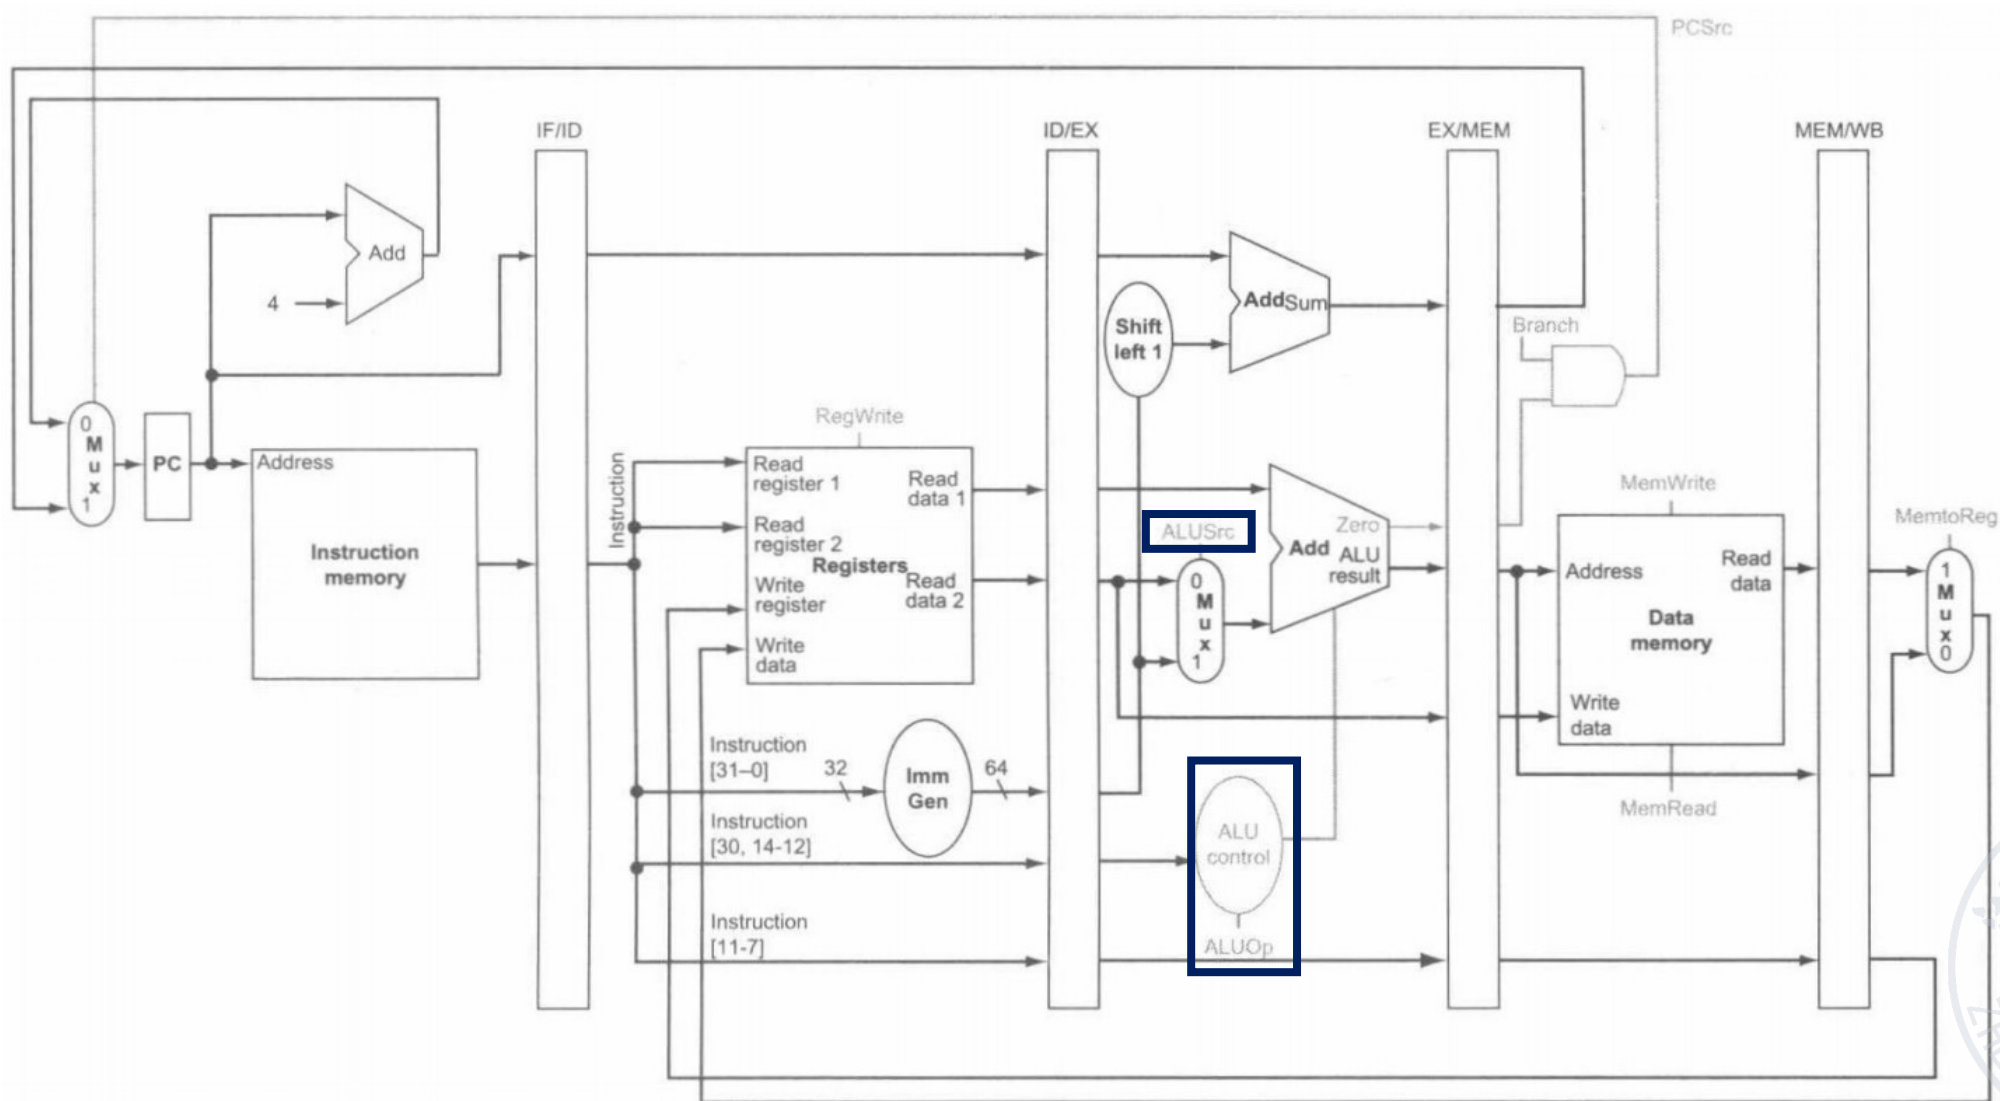
Pipelined Datapath with Control Signals



No Extra Design at Instruction Fetch and Instruction Decode



Control Signals at Execution Stage



Signals Related to ALU Operations

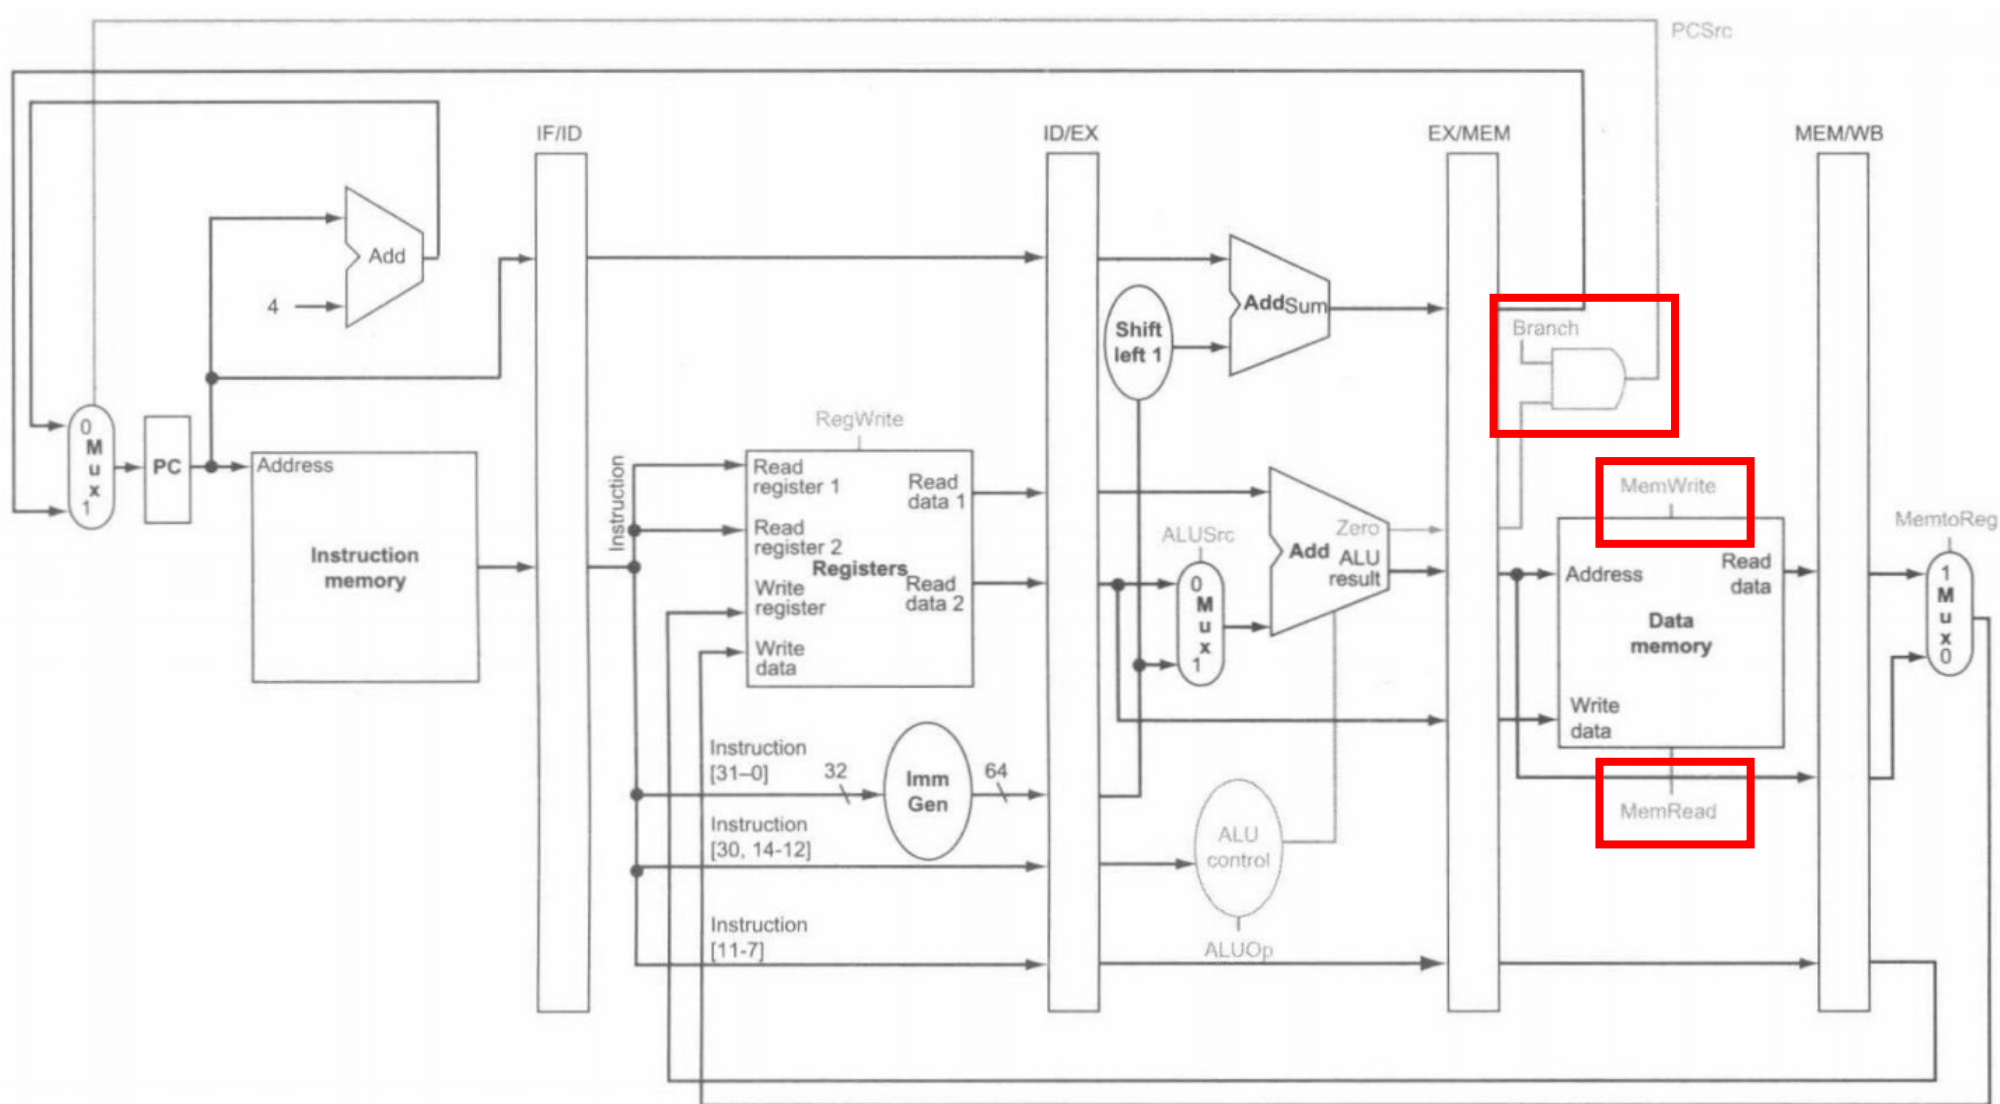
Instruction	ALUOp	operation	Funct7 field	Funct3 field	Desired ALU action	ALU control input
ld	00	load doubleword	XXXXXXX	XXX	add	0010
sd	00	store doubleword	XXXXXXX	XXX	add	0010
beq	01	branch if equal	XXXXXXX	XXX	subtract	0110
R-type	10	add	0000000	000	add	0010
R-type	10	sub	0100000	000	subtract	0110
R-type	10	and	0000000	111	AND	0000
R-type	10	or	0000000	110	OR	0001

These signals
select the
ALU operations

Signal name	Effect when deasserted	Effect when asserted
RegWrite	None.	The register on the Write register input is written with the value on the Write data input.
ALUSrc	The second ALU operand comes from the second register file output (Read data 2).	The second ALU operand is the sign-extended, 12 bits of the instruction.
PCSrc	The PC is replaced by the output of the adder that computes the value of PC + 4.	The PC is replaced by the output of the adder that computes the branch target.
MemRead	None.	Data memory contents designated by the address input are put on the Read data output.
MemWrite	None.	Data memory contents designated by the address input are replaced by the value on the Write data input.
MemtoReg	The value fed to the register Write data input comes from the ALU.	The value fed to the register Write data input comes from the data memory.



Control Signals at Memory Access Stage

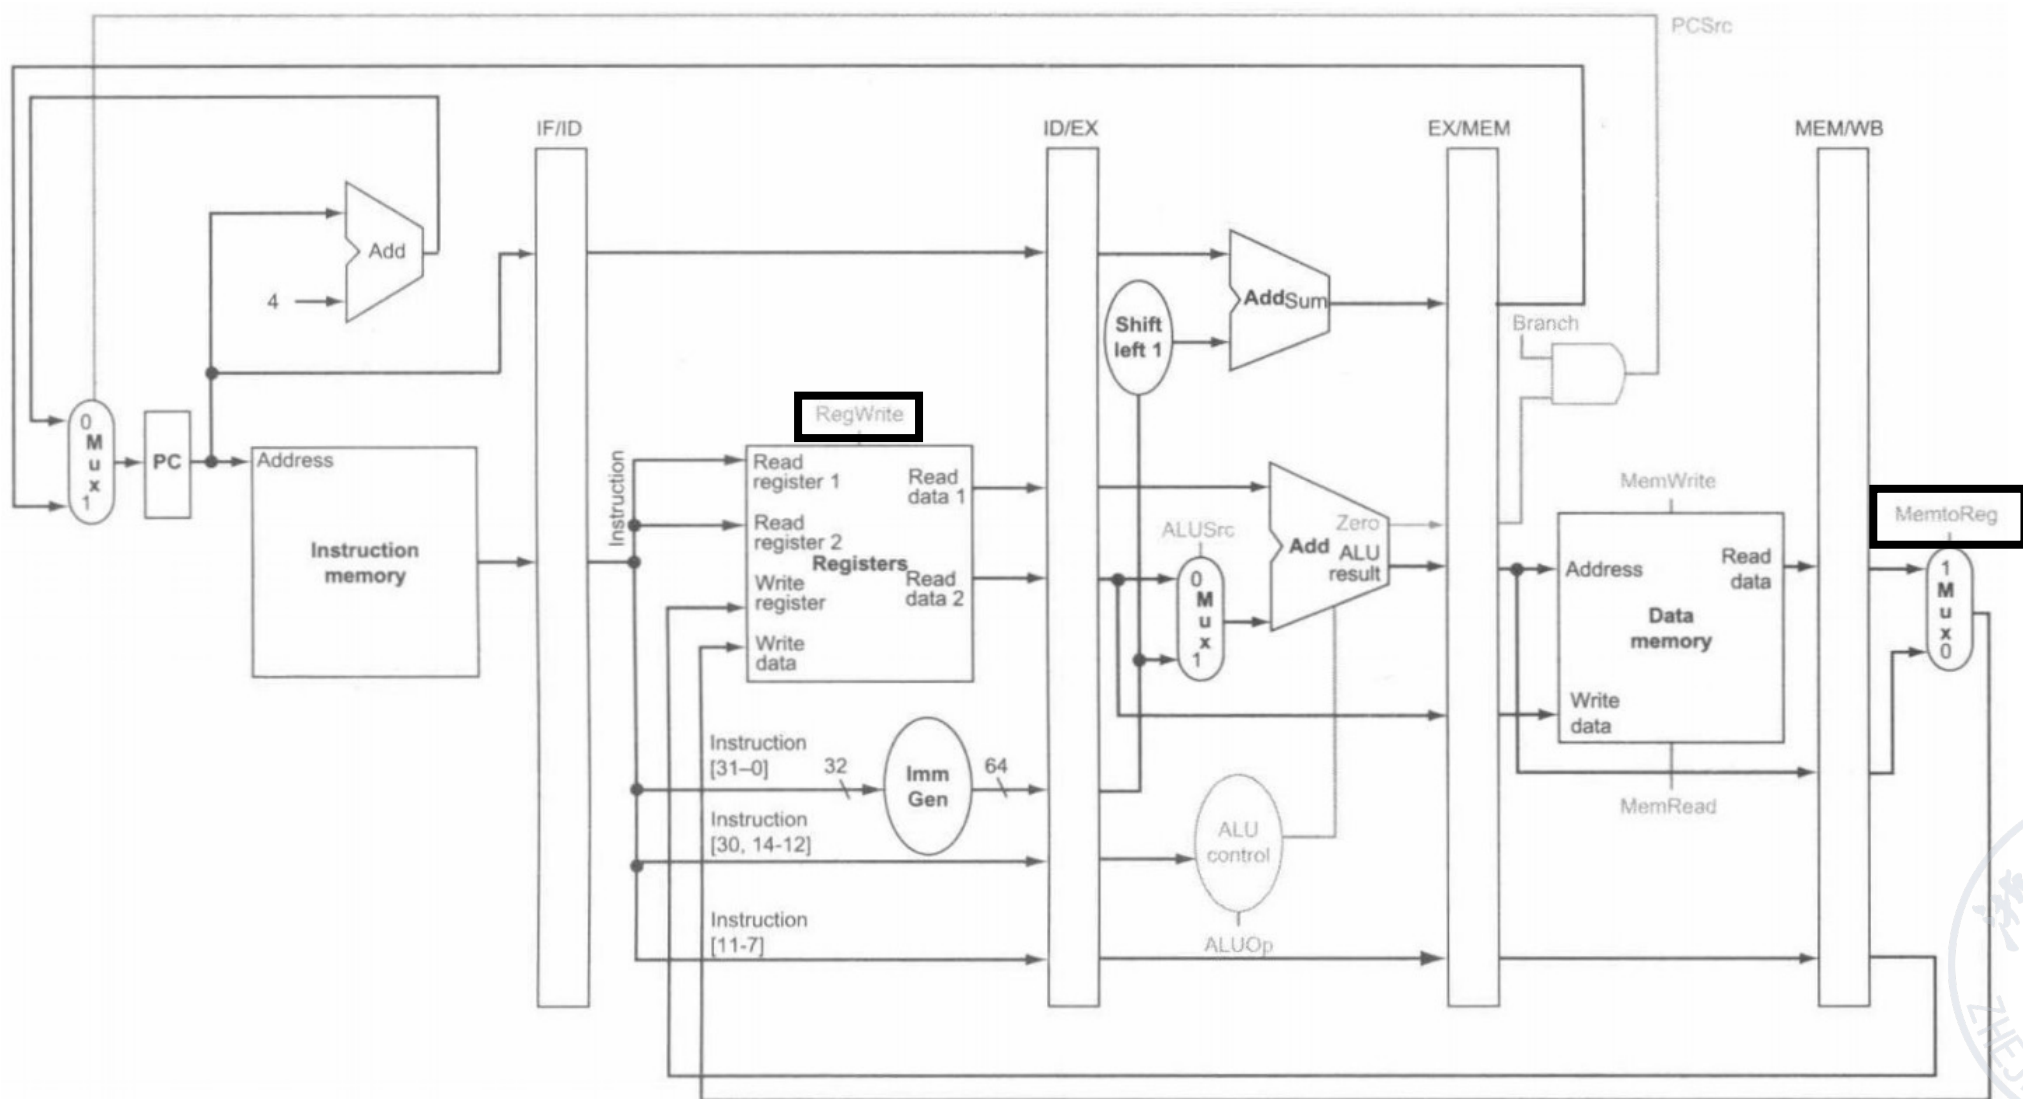


Signals Related to Branch If Equal, Load, and Store Instructions

Signal name	Effect when deasserted	Effect when asserted
RegWrite	None.	The register on the Write register input is written with the value on the Write data input.
ALUSrc	The second ALU operand comes from the second register file output (Read data 2).	The second ALU operand is the sign-extended, 12 bits of the instruction.
PCSrc	The PC is replaced by the output of the adder that computes the value of $PC + 4$.	The PC is replaced by the output of the adder that computes the branch target.
MemRead	None.	Data memory contents designated by the address input are put on the Read data output.
MemWrite	None.	Data memory contents designated by the address input are replaced by the value on the Write data input.
MemtoReg	The value fed to the register Write data input comes from the ALU.	The value fed to the register Write data input comes from the data memory.



Control Signals at Write Back Stage

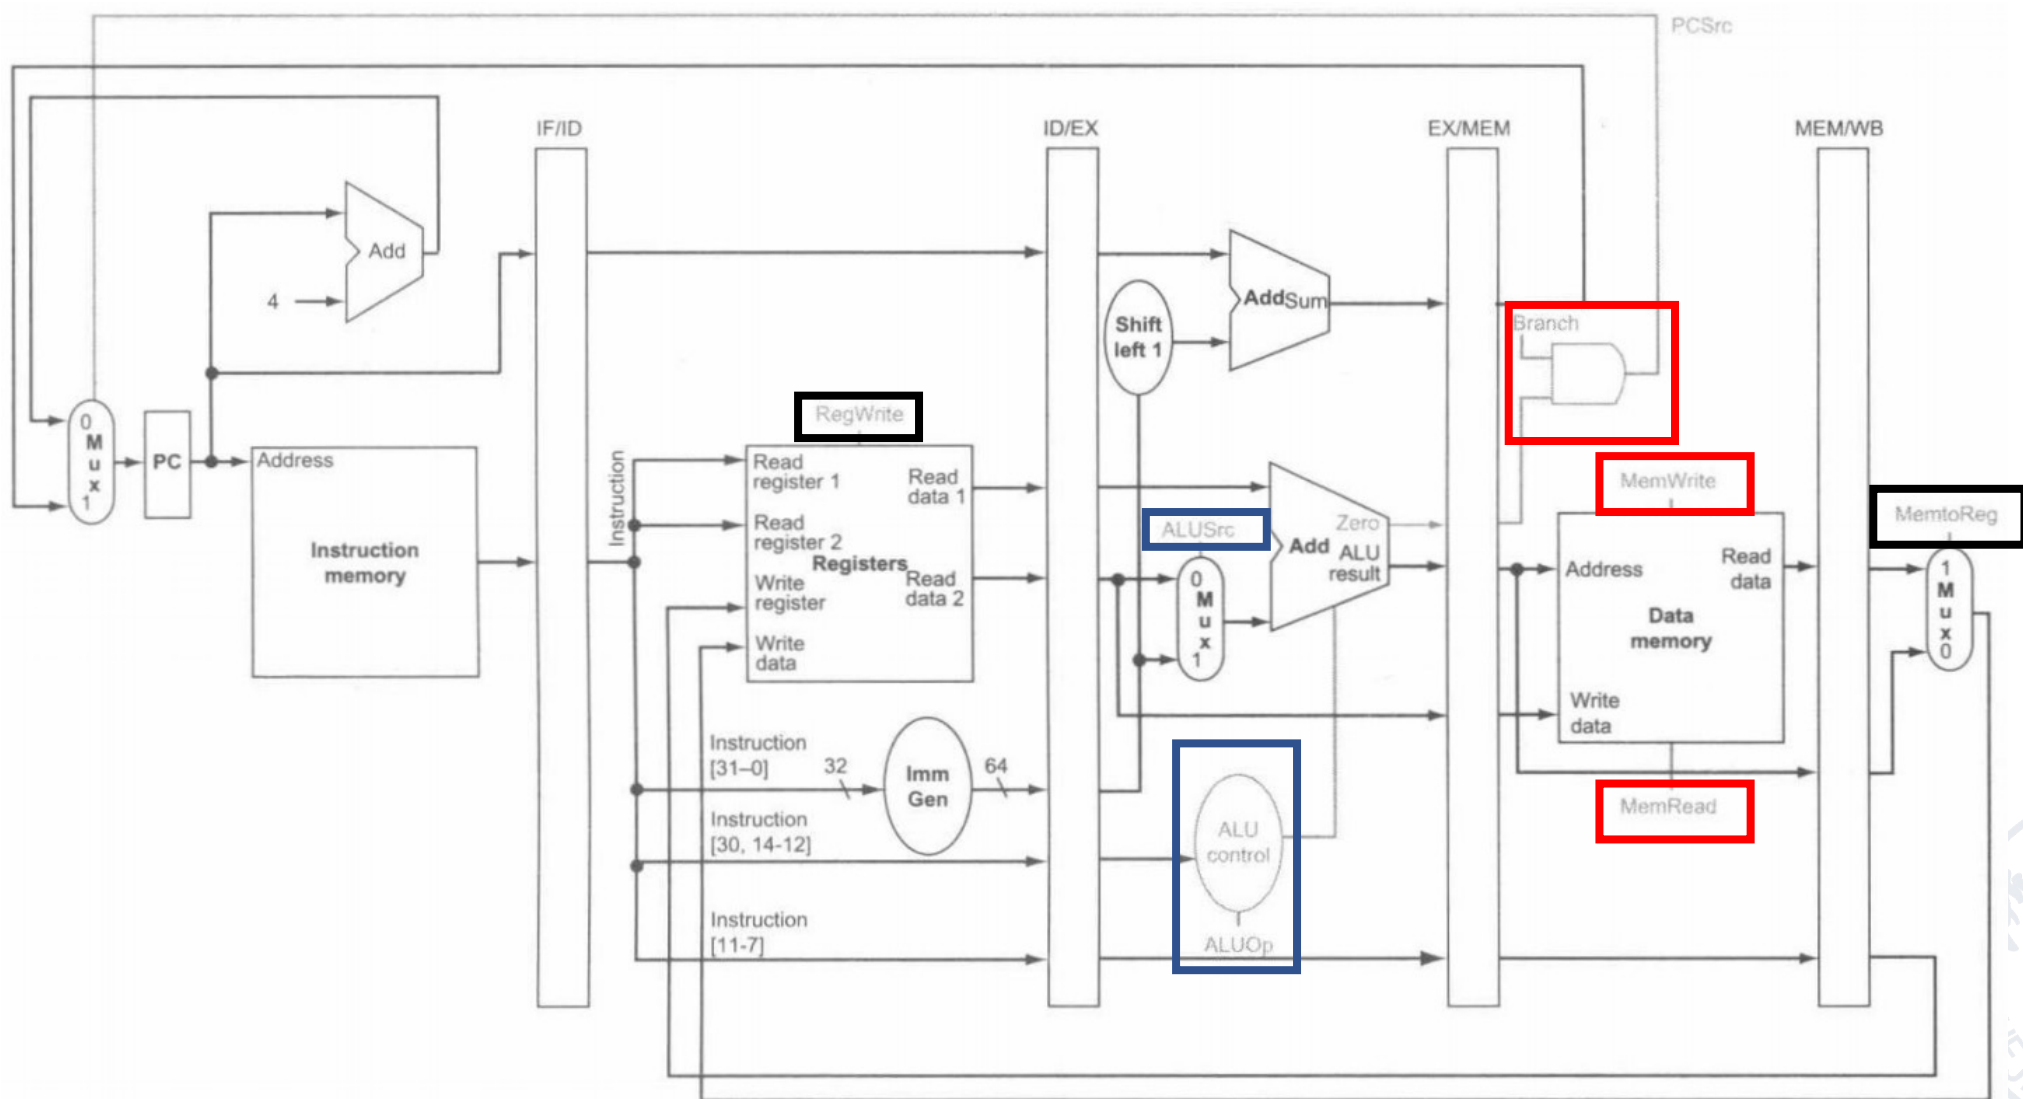


Signals Related to the Write Operation

Signal name	Effect when deasserted	Effect when asserted
RegWrite	None.	The register on the Write register input is written with the value on the Write data input.
ALUSrc	The second ALU operand comes from the second register file output (Read data 2).	The second ALU operand is the sign-extended, 12 bits of the instruction.
PCSrc	The PC is replaced by the output of the adder that computes the value of $PC + 4$.	The PC is replaced by the output of the adder that computes the branch target.
MemRead	None.	Data memory contents designated by the address input are put on the Read data output.
MemWrite	None.	Data memory contents designated by the address input are replaced by the value on the Write data input.
MemtoReg	The value fed to the register Write data input comes from the ALU.	The value fed to the register Write data input comes from the data memory.



Review of 7 Control Lines



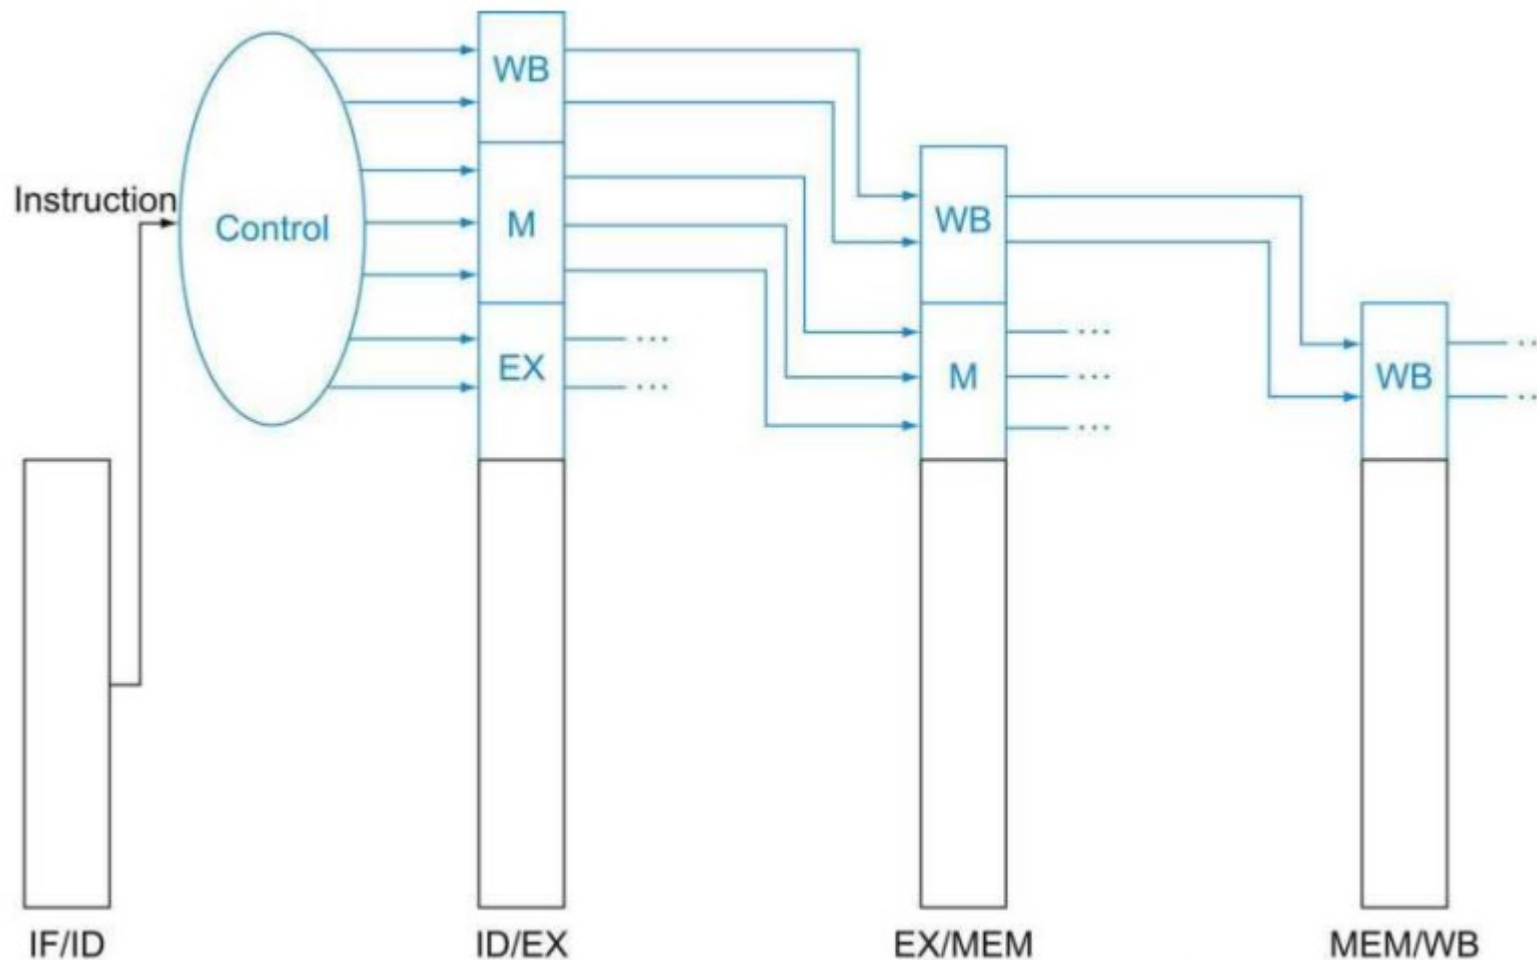
Seven Control Lines in Last Three Stages

Instruction	Execution/address calculation stage control lines		Memory access stage control lines			Write-back stage control lines	
	ALUOp	ALUSrc	Branch	Mem-Read	Mem-Write	Reg-Write	Memto-Reg
R-format	10	0	0	0	0	1	0
ld	00	1	0	1	0	1	1
sd	00	1	0	0	1	0	X
beq	01	0	1	0	0	0	X

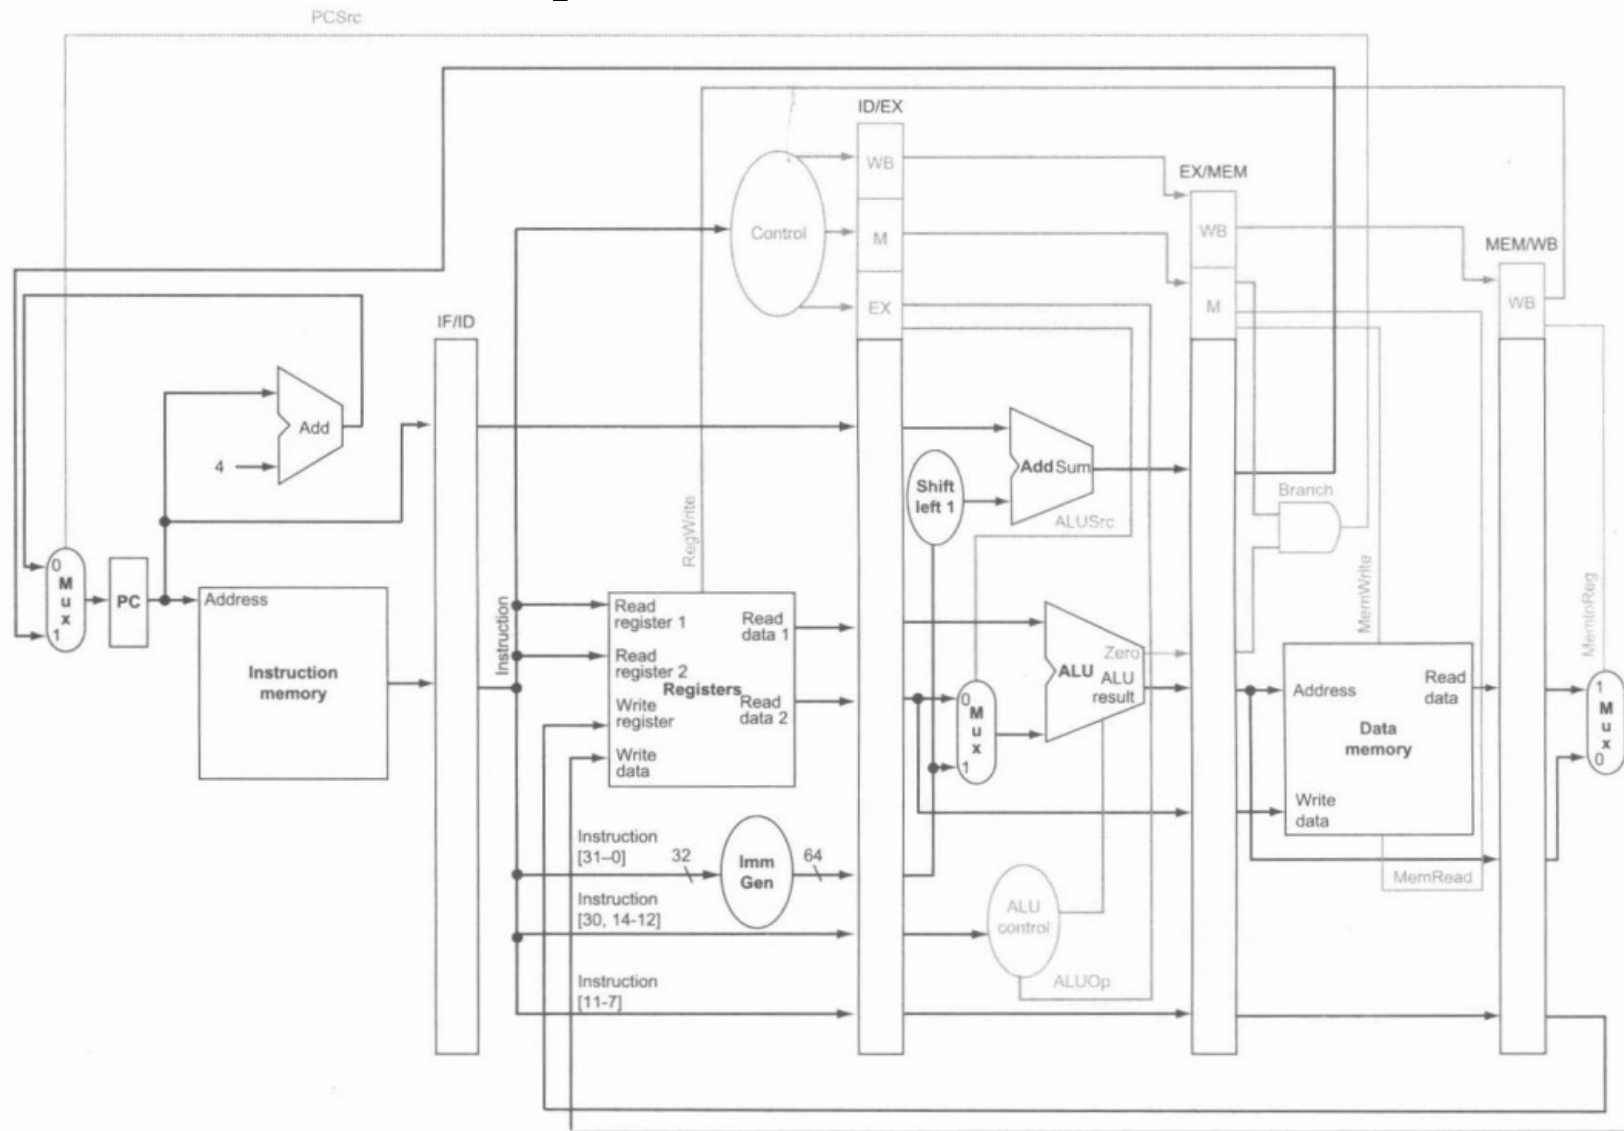




Extend Pipeline Registers to Include Control Information



Pipelined Datapath with the Control Signals

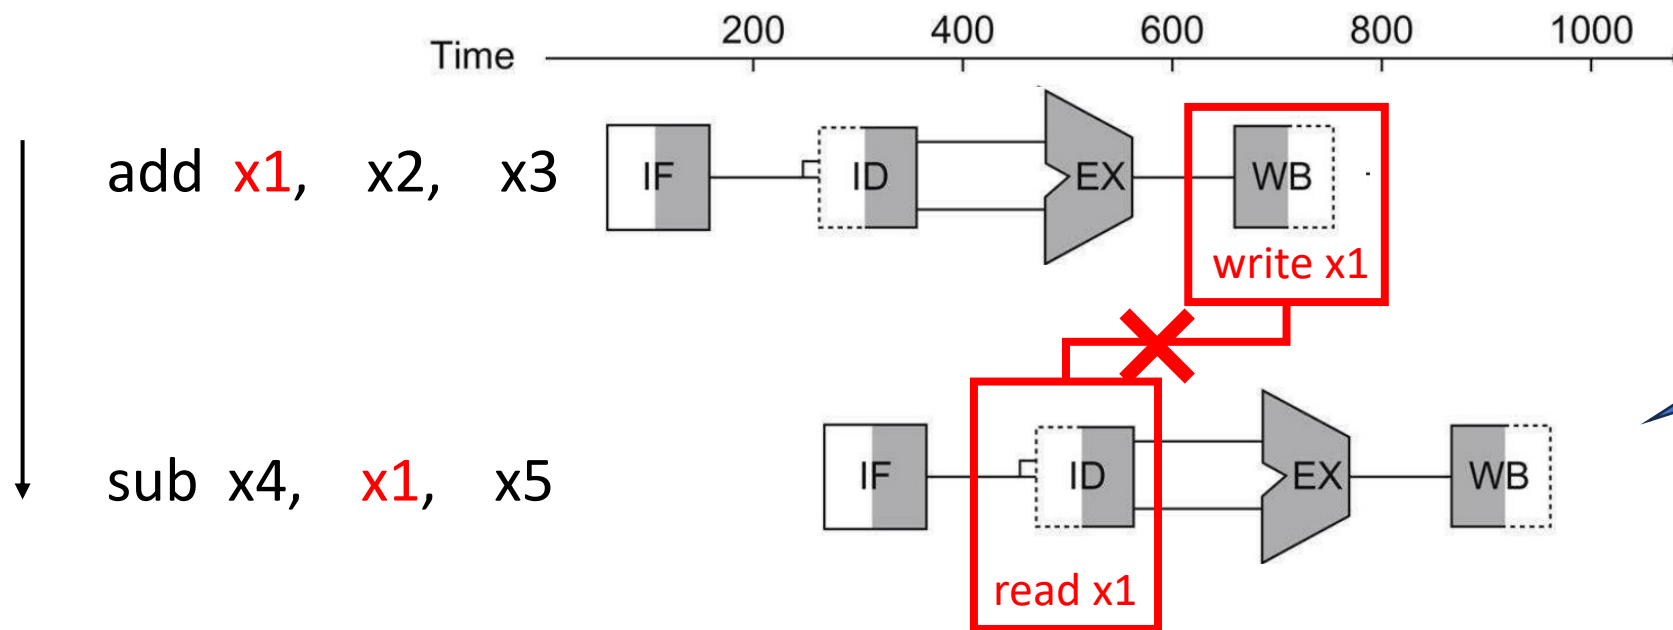


Pipelining with Stall



Stall Condition Detection

- Data dependency between instructions



Hazard! But how to formulate it?



Stall Condition Detection

- Pass register numbers along pipeline
 - e.g., ID/EX.Rs1 = register number for Rs1 sitting in ID/EX pipeline register
- ALU operand register numbers in EX stage are given by
 - ID/EX.Rs1, ID/EX.Rs2
- Data hazards when
 - 1a. EX/MEM. Rd = ID/EX. Rs1
 - 1b. EX/MEM. Rd = ID/EX. Rs2

Fwd from EX/MEM
pipeline reg



Stall Condition Detection

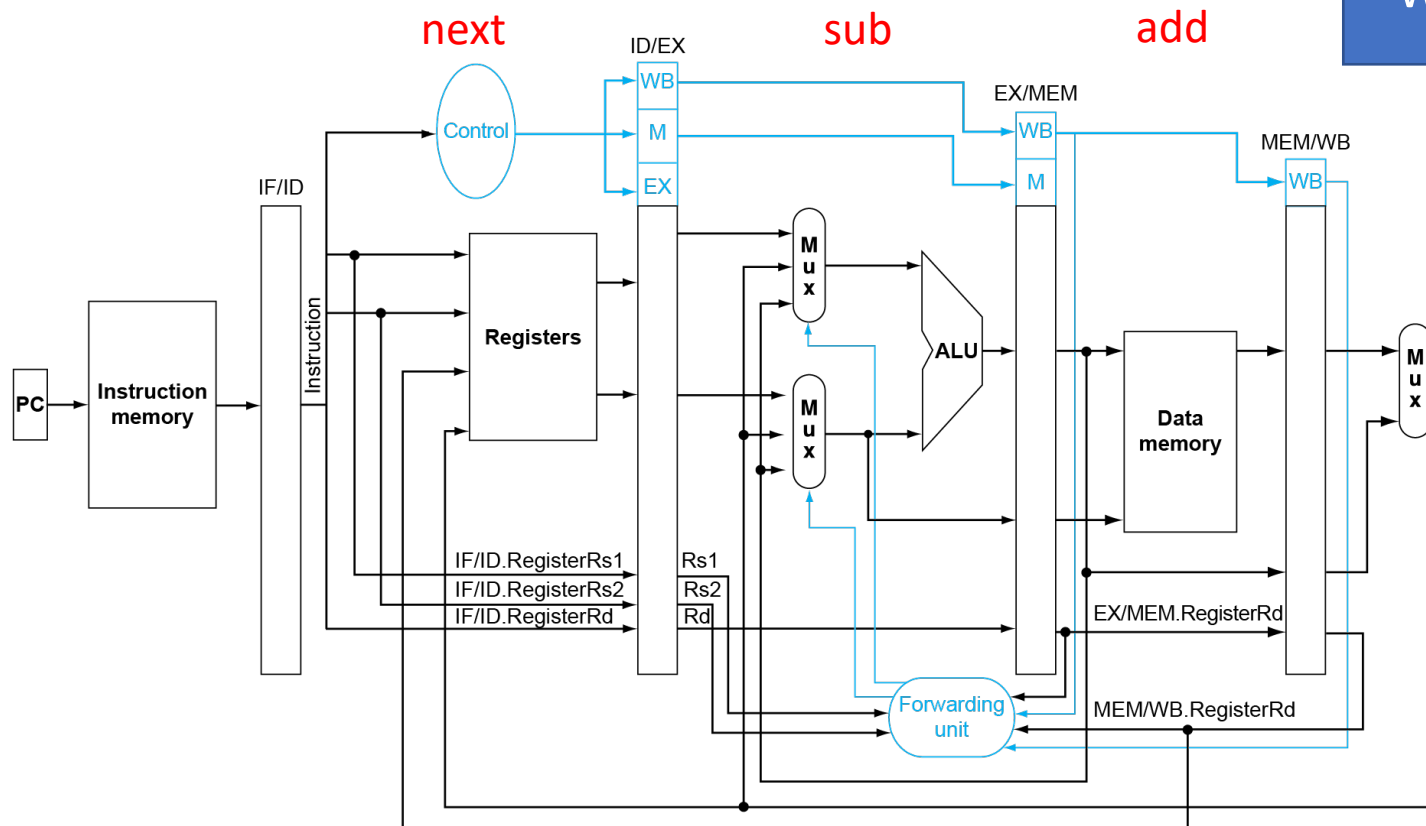
- But only if forwarding instruction will write to a register!
 - EX/MEM.RegWrite
- And only if Rd for that instruction is not x0
 - EX/MEM. Rd \neq 0



How to Stall the Pipeline

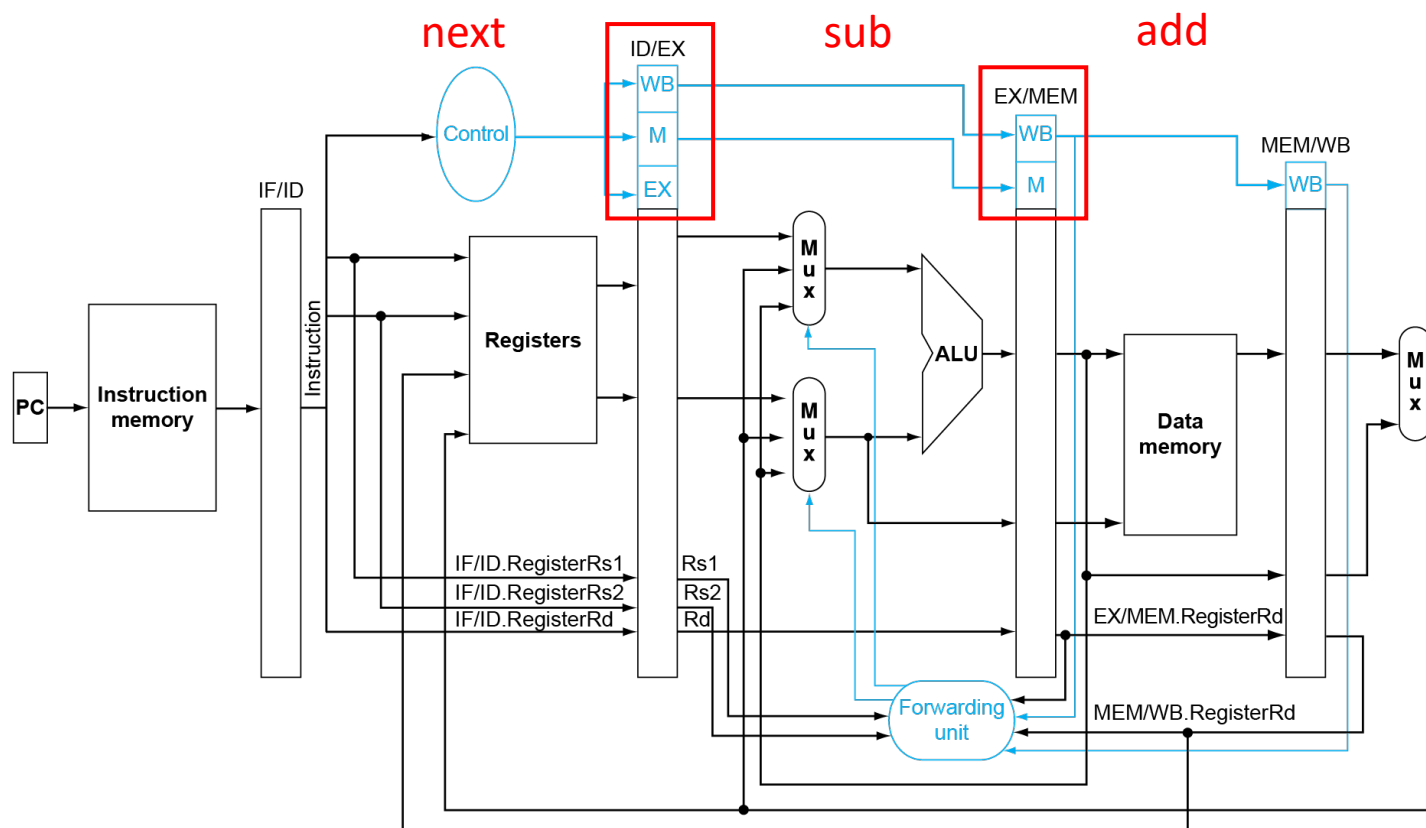
- Force control values in EX/MEM register to 0
 - EX, MEM and WB do nop (no-operation)

Which one?



How to Stall the Pipeline

- Force control values in EX/MEM register to 0
 - EX, MEM and WB do nop (no-operation)



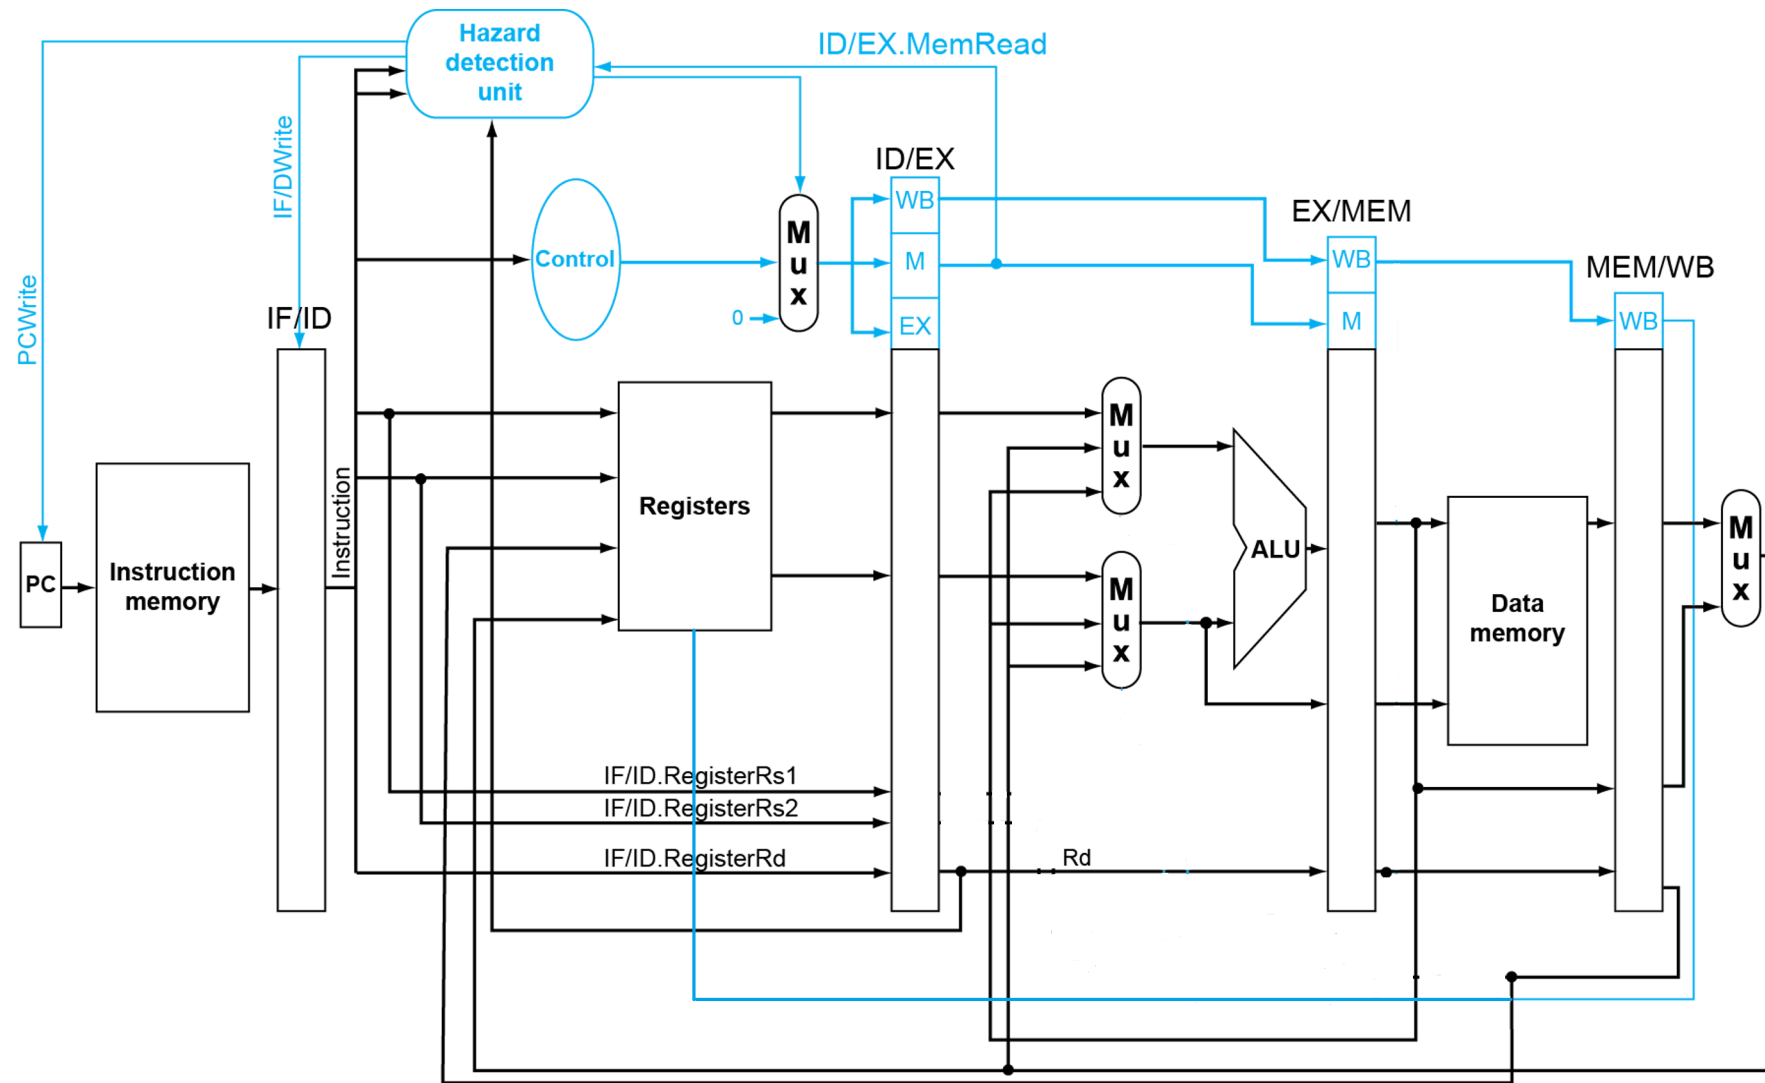
What's more?

How to Stall the Pipeline

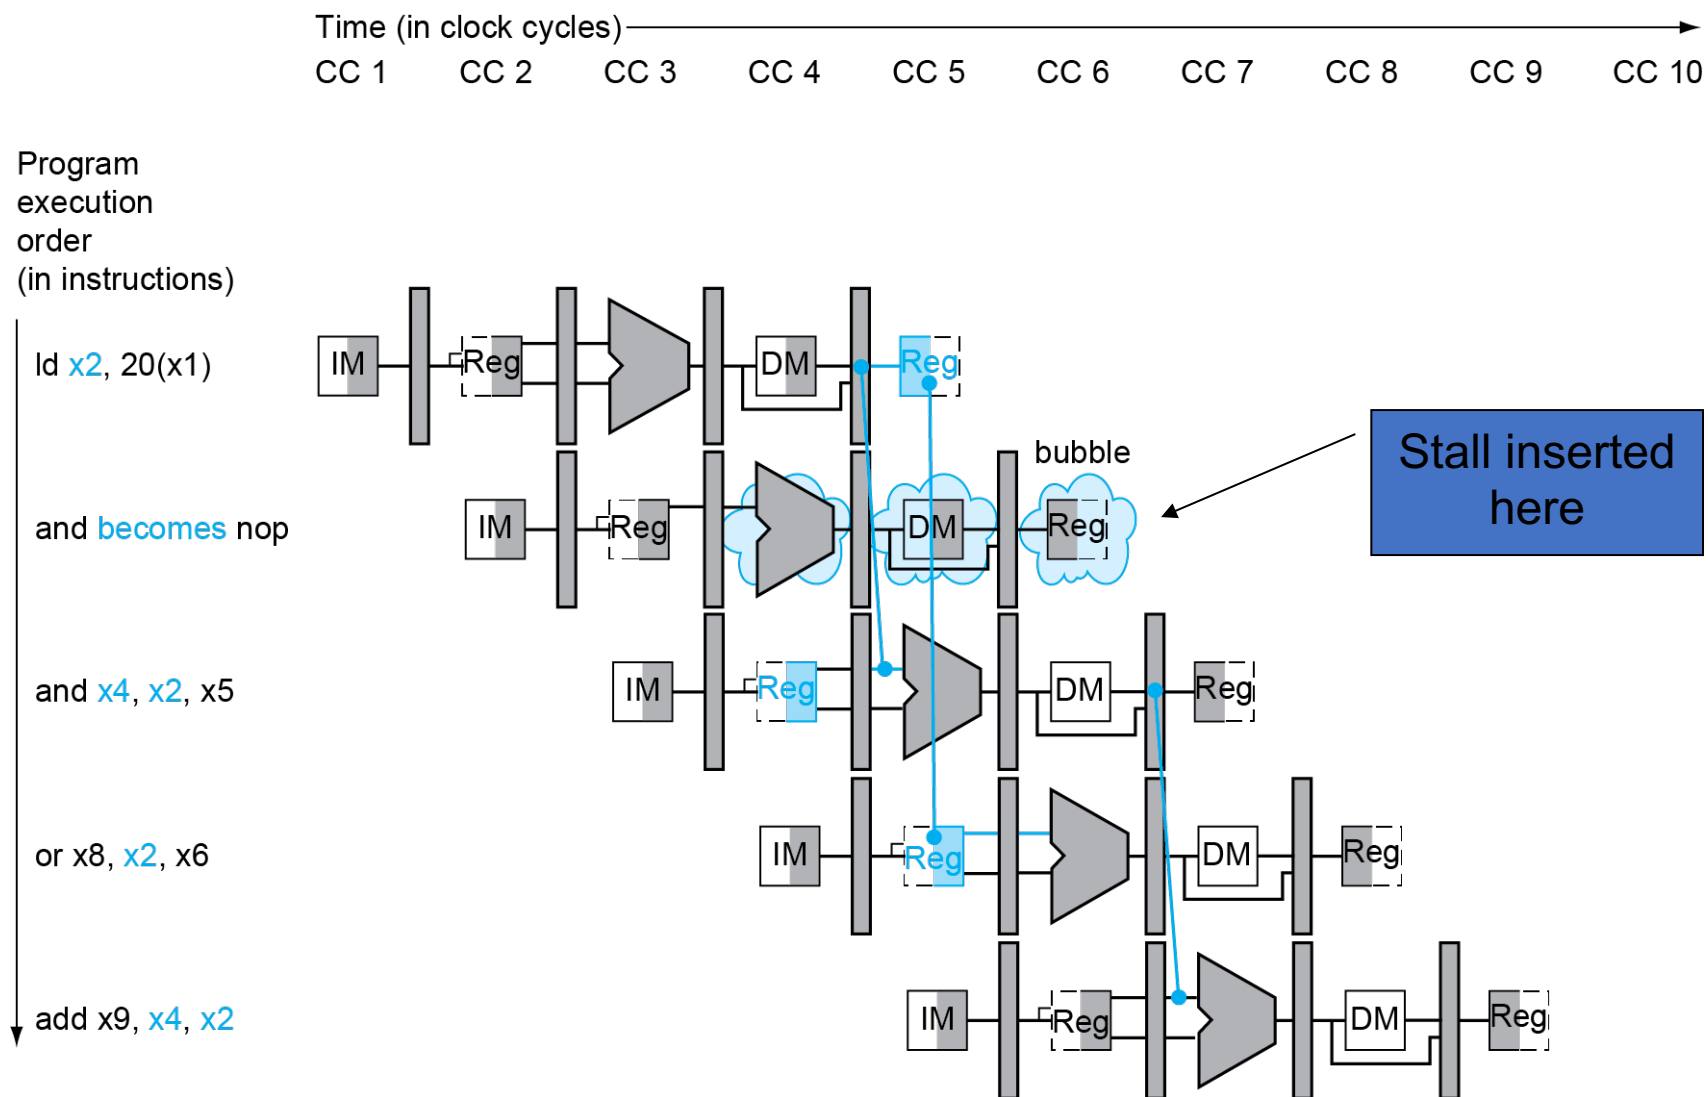
- Force control values in EX/MEM register to 0
 - EX, MEM and WB do nop (no-operation)
- Prevent update of PC and IF/ID register
 - Using instruction is decoded again
 - Following instruction is fetched again
 - 1-cycle stall allows MEM to read data for 1d
 - Can subsequently forward to EX stage



Datapath with Stall



Stall/Bubble in the Pipeline



Stall/Bubble in the Pipeline

