

Computer Systems II

Li Lu

Room 319, Yifu Business and Management Building

Yuquan Campus

li.lu@zju.edu.cn

<https://person.zju.edu.cn/lynnluli>



Pipeline Hazards

- Structural Hazard A required resource is busy
- Data Hazards
 - Data dependency between instructions
 - Need to wait for previous instruction to complete its data read/write
- Control Hazards Flow of execution depends on previous instruction



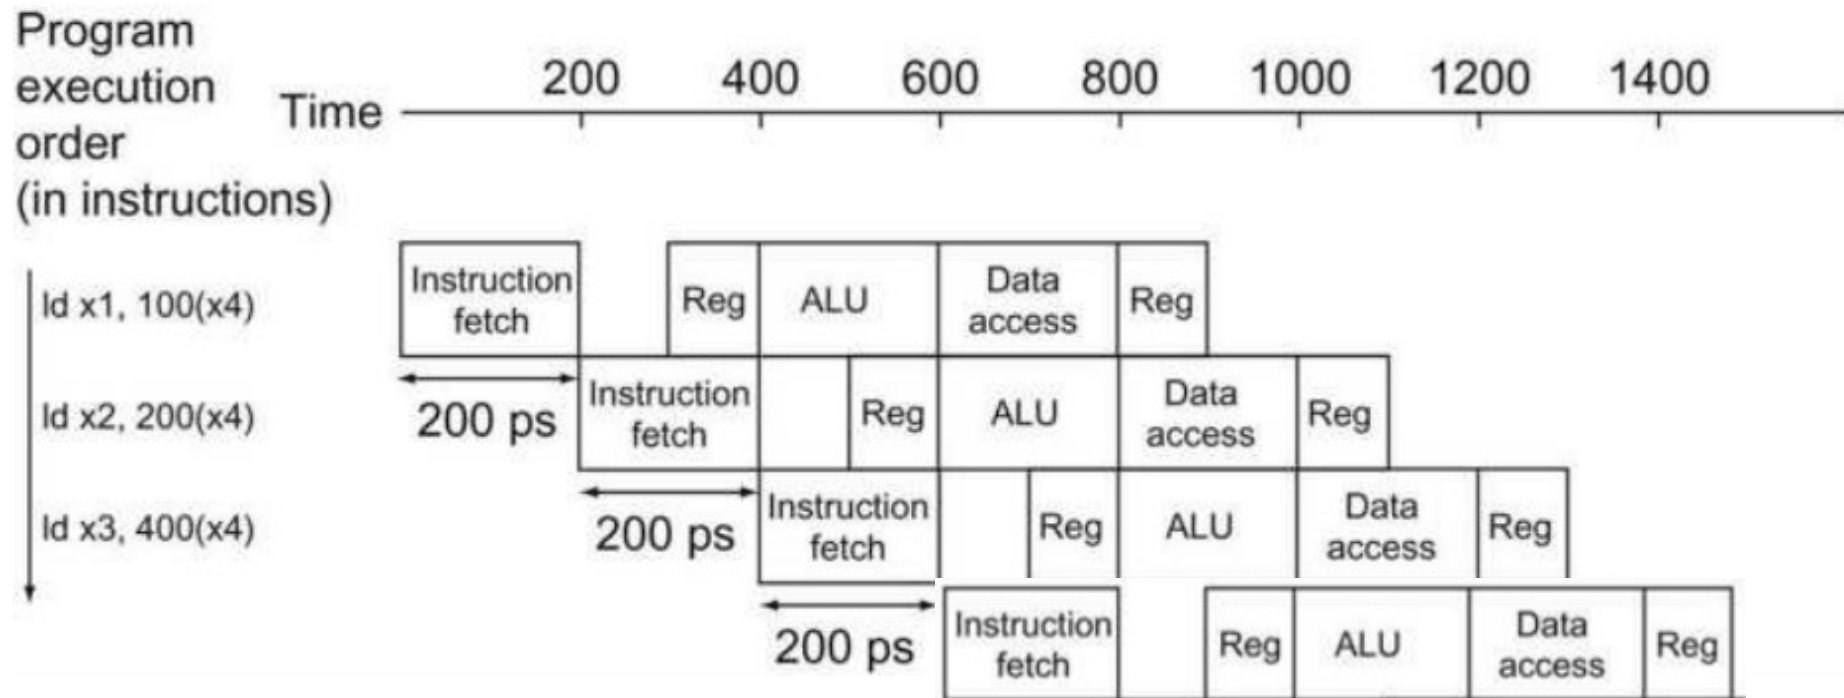
Structural Hazard



Structural Hazard

A required resource is busy

Example: Consider the situation while the pipeline only has a single memory



Question: Can the four instructions execute correctly?



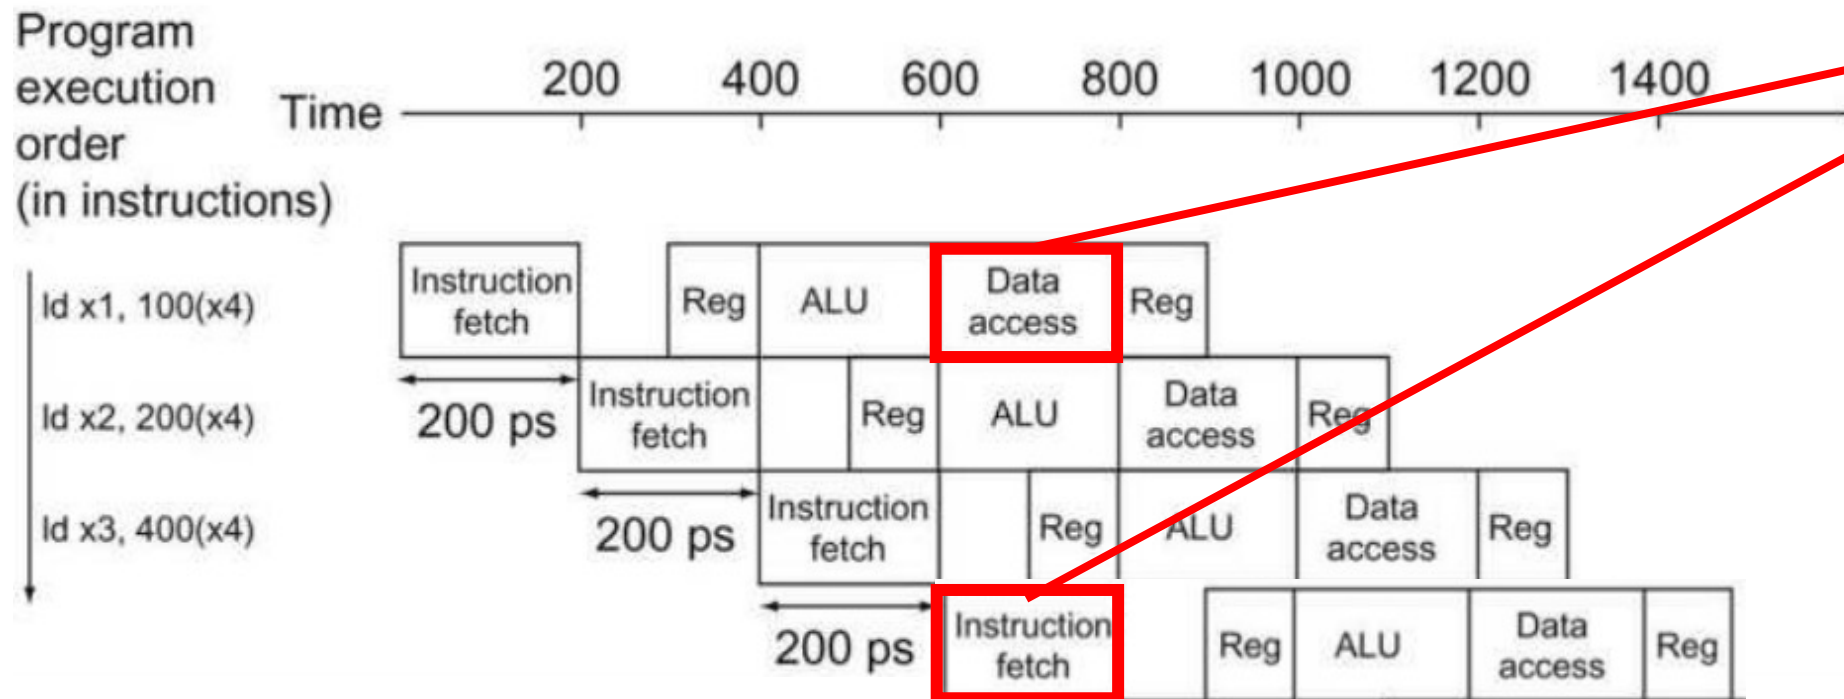
Structural Hazard

A required resource is busy

Solution:

Use Instruction and data memory simultaneously.

Example: Consider the situation while the pipeline only has **a single memory**



How to Deal with Structural Hazard?

Problem: Two or more instructions in the pipeline compete for access to a single physical resource

- Solution 1: Instructions take it in turns to use resource, some instructions have to stall.
- Solution 2: Add more hardware to machine.

Can always solve a structural hazard by adding more hardware



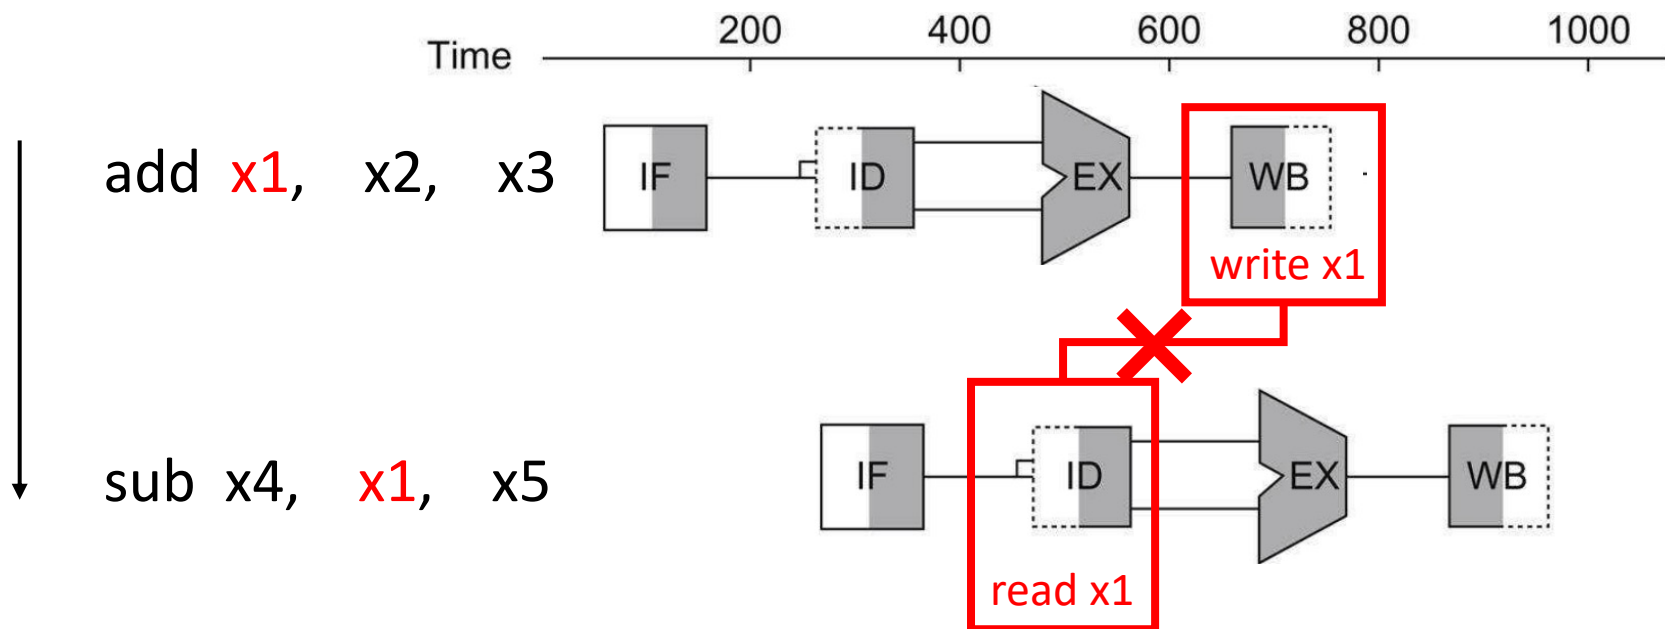
Data Hazards



Data Hazard

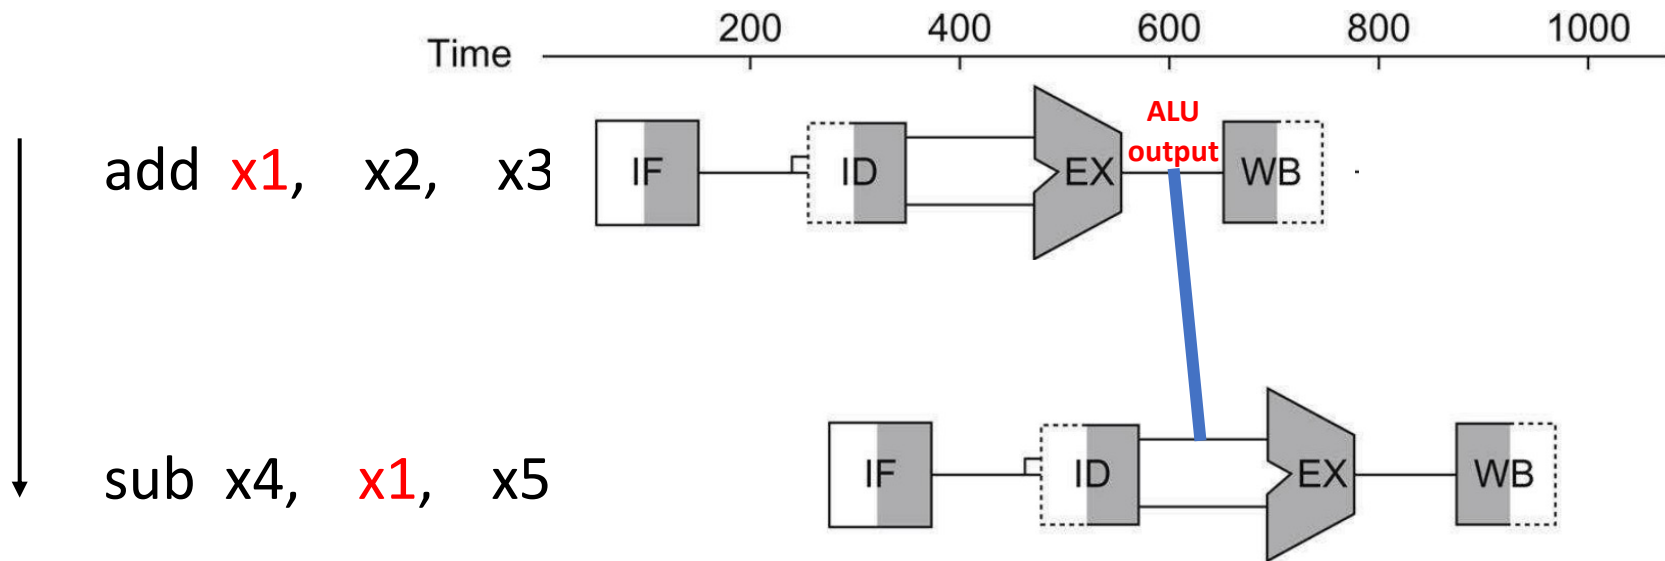
- Data dependency between instructions
- Need to wait for previous instruction to complete its data read/write

Problem: Instruction depends on result from previous



Data Hazard

Solution “forwarding”: Adding extra hardware to retrieve the missing item early from the internal resources



Data Hazard in ALU Instructions

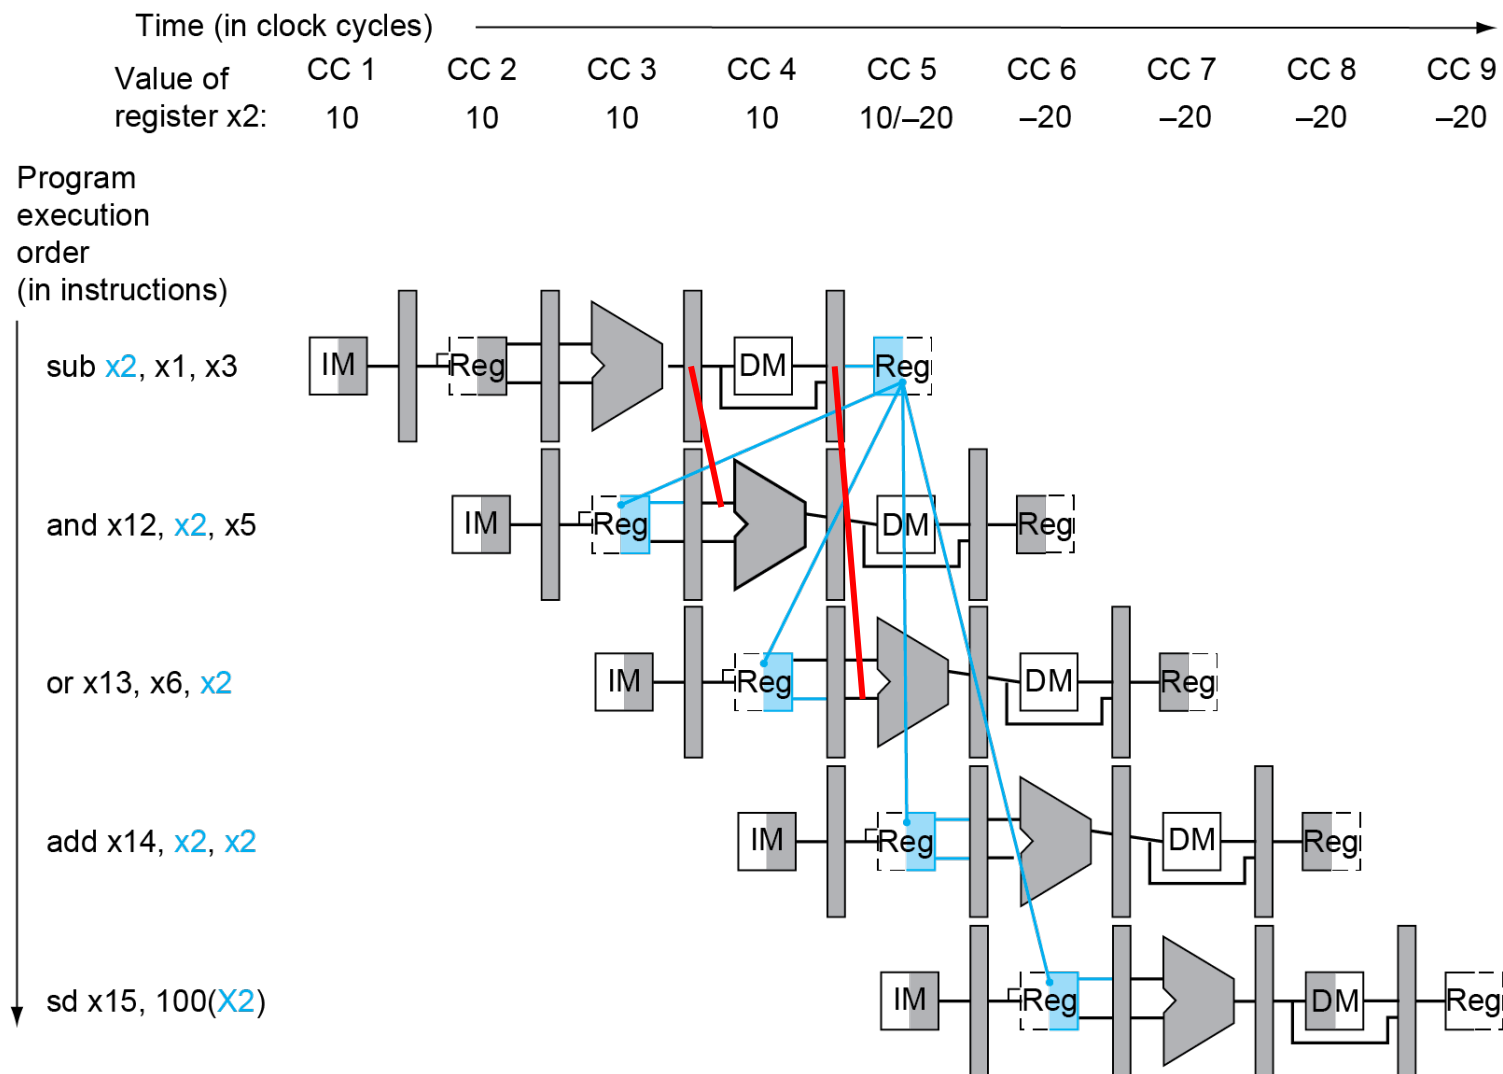
- Consider this sequence:

```
sub  x2, x1, x3
and  x12, x2, x5
or   x13, x6, x2
add  x14, x2, x2
sd   x15, 100(x2)
```

- We can resolve hazards with forwarding
 - How do we detect when to forward?




Dependencies & Forwarding



Detecting the Need to Forward

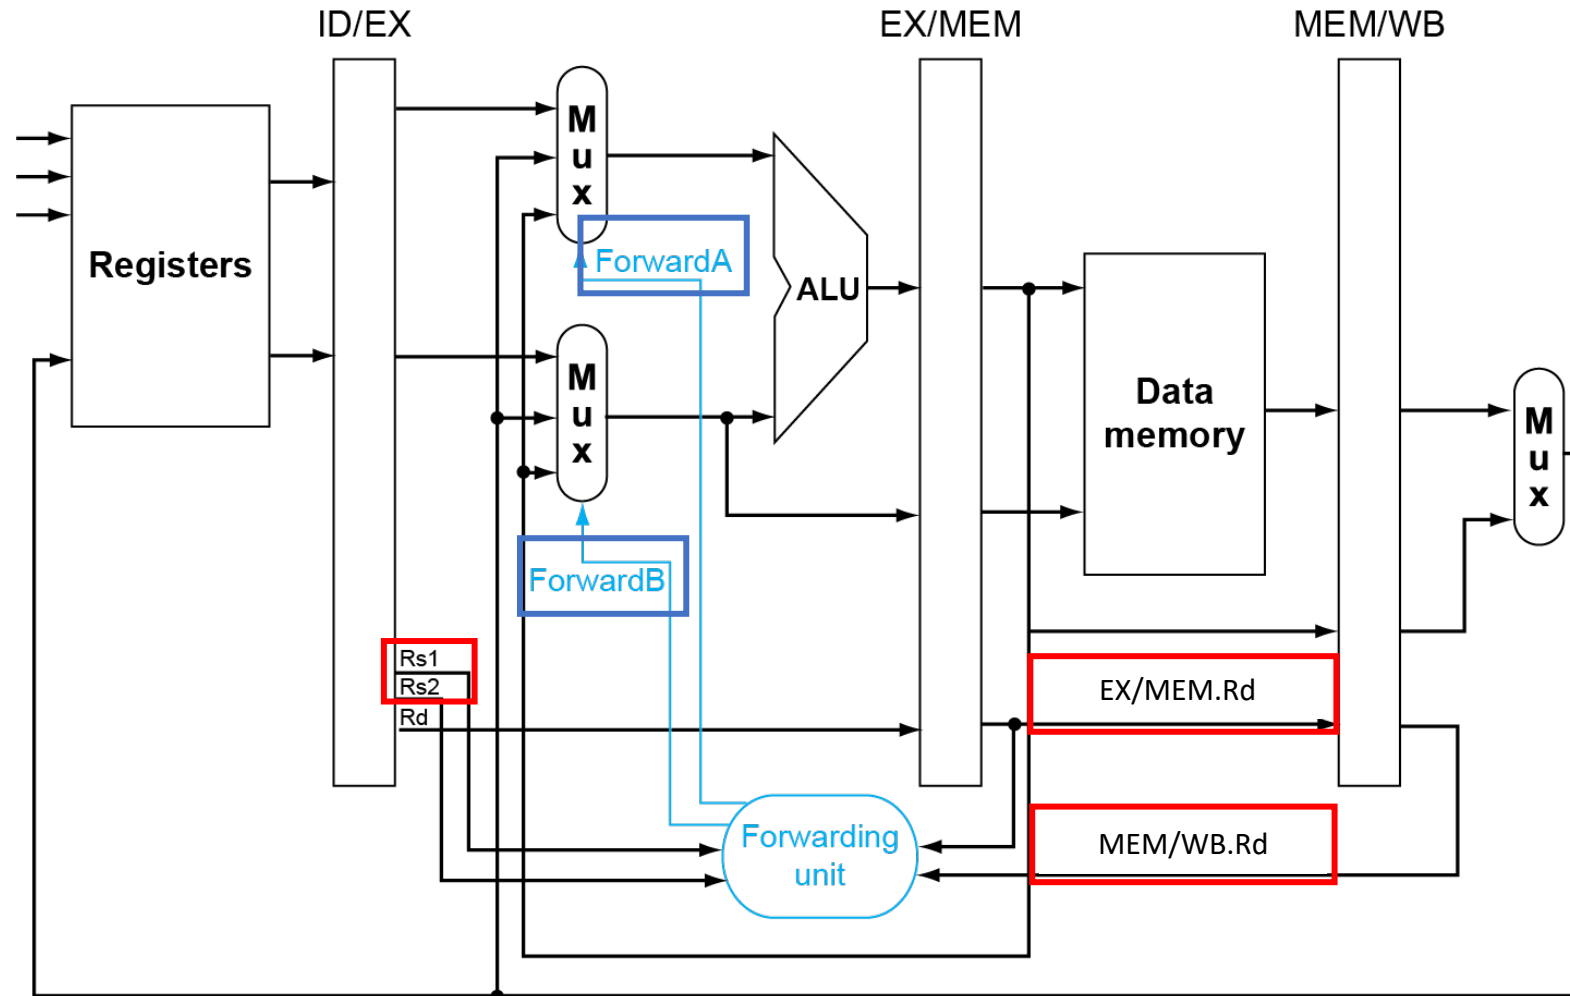
- Pass register numbers along pipeline
 - e.g., ID/EX.Rs1 = register number for Rs1 sitting in ID/EX pipeline register
- ALU operand register numbers in EX stage are given by
 - ID/EX.Rs1, ID/EX.Rs2
- Data hazards when
 - 1a. EX/MEM. Rd = ID/EX. Rs1
 - 1b. EX/MEM. Rd = ID/EX. Rs2
 - 2a. MEM/WB. Rd = ID/EX. Rs1
 - 2b. MEM/WB. Rd = ID/EX. Rs2



Detecting the Need to Forward

- But only if forwarding instruction will write to a register!
 - EX/MEM.RegWrite, MEM/WB.RegWrite
- And only if Rd for that instruction is not x0
 - EX/MEM. Rd \neq 0,
MEM/WB. Rd \neq 0

Forwarding Paths

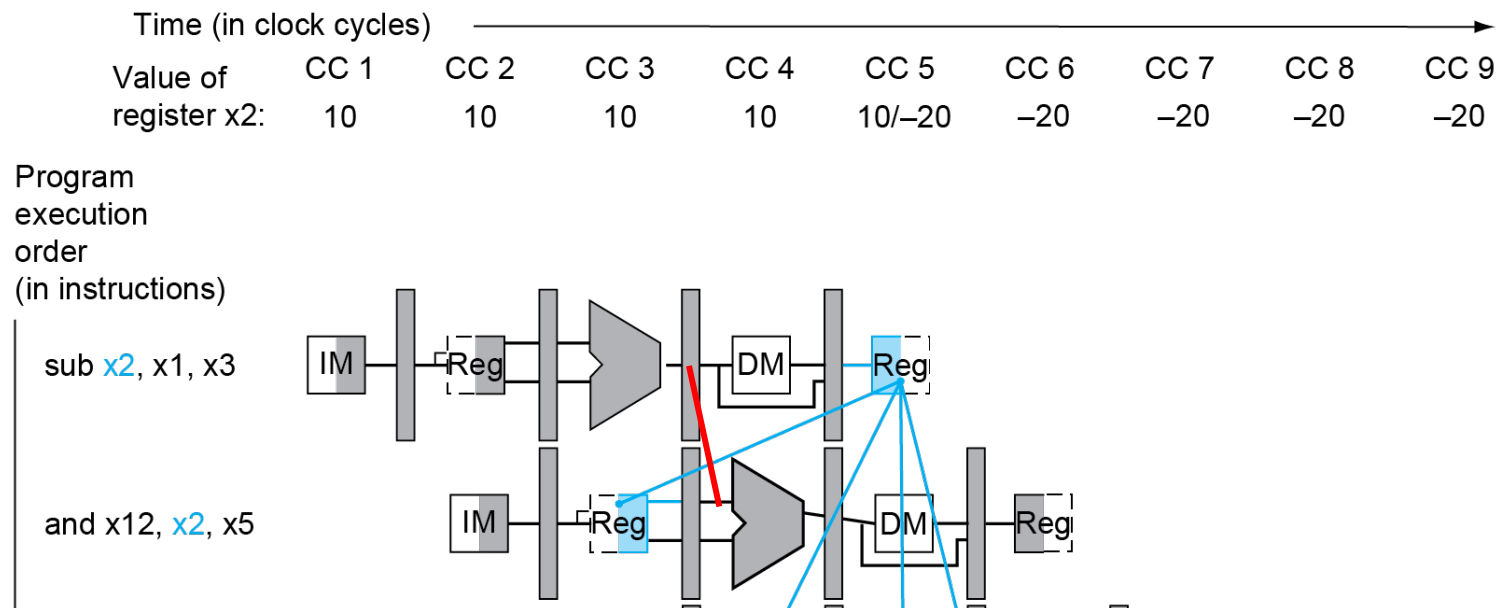


Forwarding Conditions

Mux control	Source	Explanation
ForwardA = 00	ID/EX	The first ALU operand comes from the register file.
ForwardA = 10	EX/MEM	The first ALU operand is forwarded from the prior ALU result.
ForwardA = 01	MEM/WB	The first ALU operand is forwarded from data memory or an earlier ALU result.
ForwardB = 00	ID/EX	The second ALU operand comes from the register file.
ForwardB = 10	EX/MEM	The second ALU operand is forwarded from the prior ALU result.
ForwardB = 01	MEM/WB	The second ALU operand is forwarded from data memory or an earlier ALU result.

Example 1

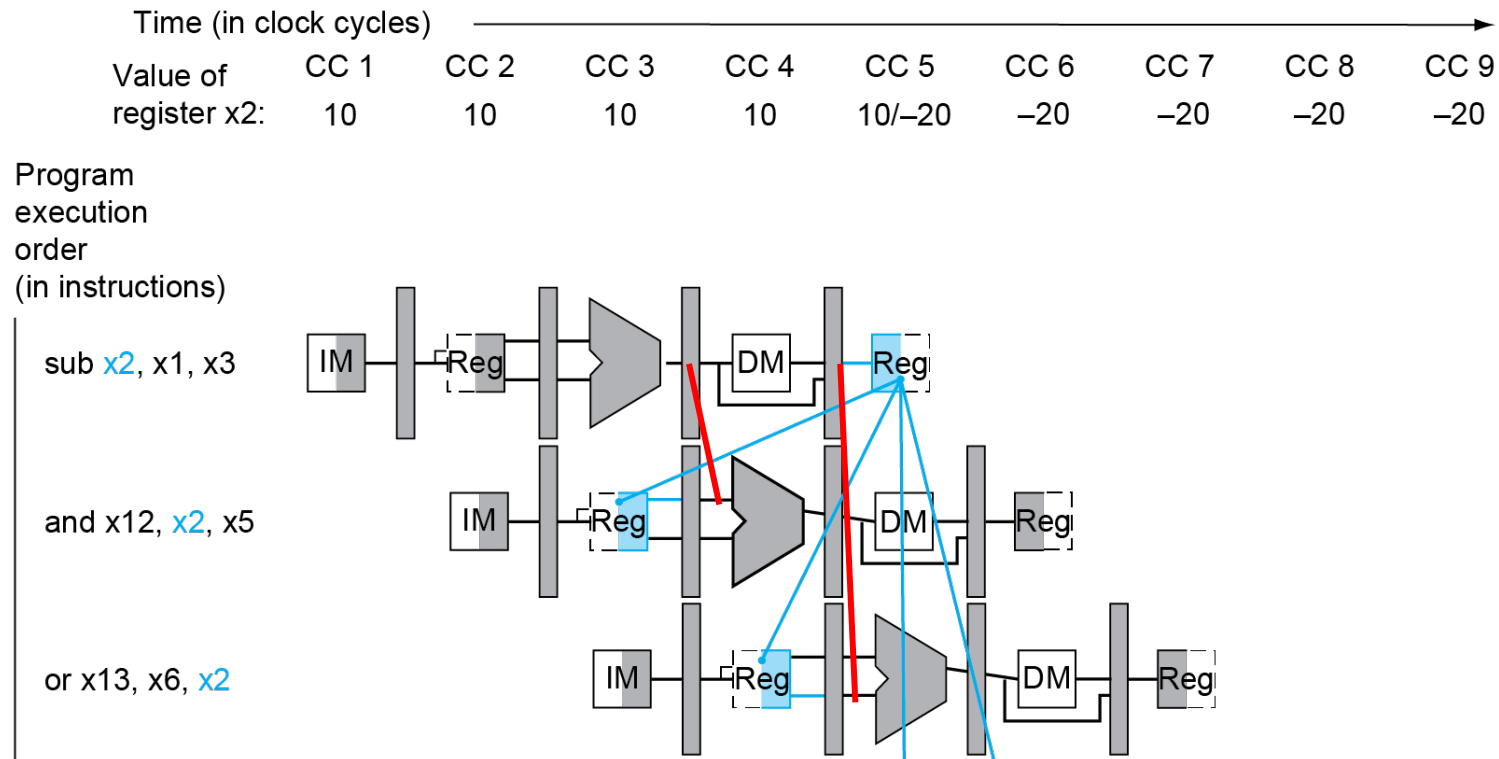
What is the Data hazards condition in the following case?



EX/MEM. Rd = ID/EX. Rs1

Example 2

What is the Data hazards condition in the following case?



MEM/WB. Rd = ID/EX. Rs2

Forwarding Conditions

Mux control	Source	Explanation
ForwardA/B = 00	ID/EX	No forwarding
ForwardA/B = 10	EX/MEM	Forwarding with data hazard in EX/MEM
ForwardA/B = 01	MEM/WB	Forwarding with data hazard in MEM/WB

Forwarding Conditions

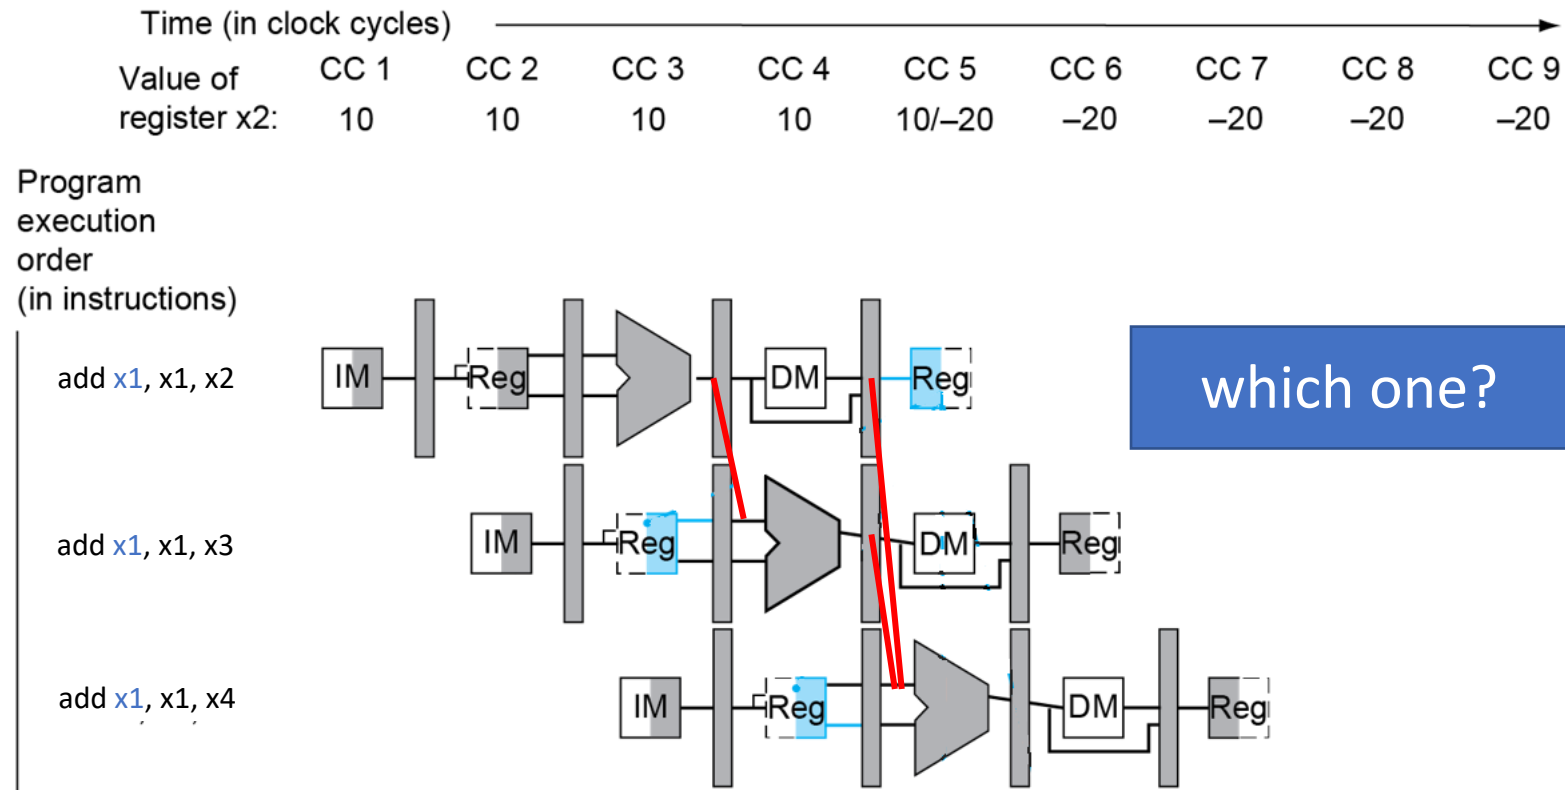
- EX hazard
 - if (EX/MEM.RegWrite and (EX/MEM. Rd \neq 0)
and (EX/MEM. Rd = ID/EX. Rs1))
ForwardA = 10
 - if (EX/MEM.RegWrite and (EX/MEM. Rd \neq 0)
and (EX/MEM. Rd = ID/EX. Rs2))
ForwardB = 10
- MEM hazard
 - if (MEM/WB.RegWrite and (MEM/WB. Rd \neq 0)
and (MEM/WB. Rd = ID/EX. Rs1))
ForwardA = 01
 - if (MEM/WB.RegWrite and (MEM/WB. Rd \neq 0)
and (MEM/WB. Rd = ID/EX. Rs2))
ForwardB = 01



Double Data Hazard

- Consider the sequence:
 - add **x1**, x1, x2
 - add **x1**, **x1**, x3
 - add x1, **x1**, x4
- Both hazards occur
 - Want to use the most recent

Double Data Hazard



Such an exception should be added into MEM hazards!

Double Data Hazard

- Consider the sequence:
 - add x1, x1, x2
 - add x1, x1, x3
 - add x1, x1, x4
- Both hazards occur
 - Want to use the most recent
- Revise MEM hazard condition
 - Only fwd if EX hazard condition isn't true

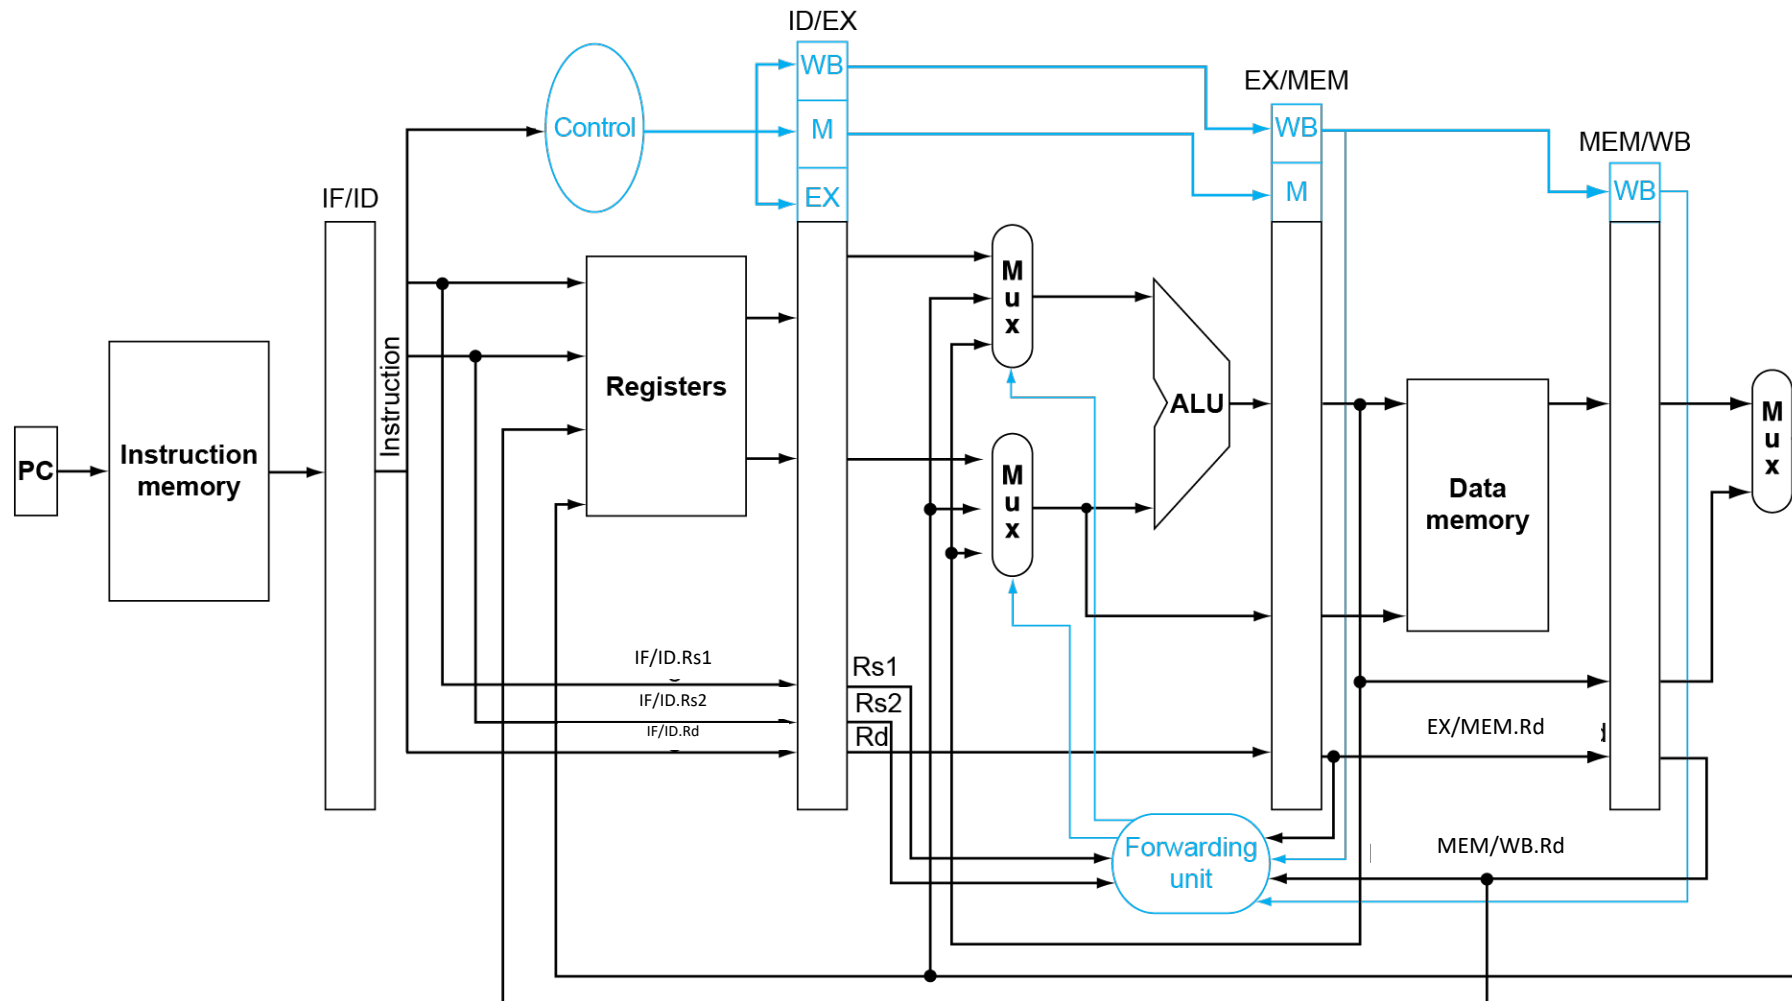
Revised Forwarding Condition

- MEM hazard
 - if (MEM/WB.RegWrite and (MEM/WB. Rd \neq 0)
and not(EX/MEM.RegWrite and (EX/MEM. Rd \neq 0)
and (EX/MEM. Rd = ID/EX. Rs1))
and (MEM/WB. Rd = ID/EX. Rs1))
ForwardA = 01
 - if (MEM/WB.RegWrite and (MEM/WB. Rd \neq 0)
and not(EX/MEM.RegWrite and (EX/MEM. Rd \neq 0)
and (EX/MEM. Rd = ID/EX. Rs2))
and (MEM/WB. Rd = ID/EX. Rs2))
ForwardB = 01

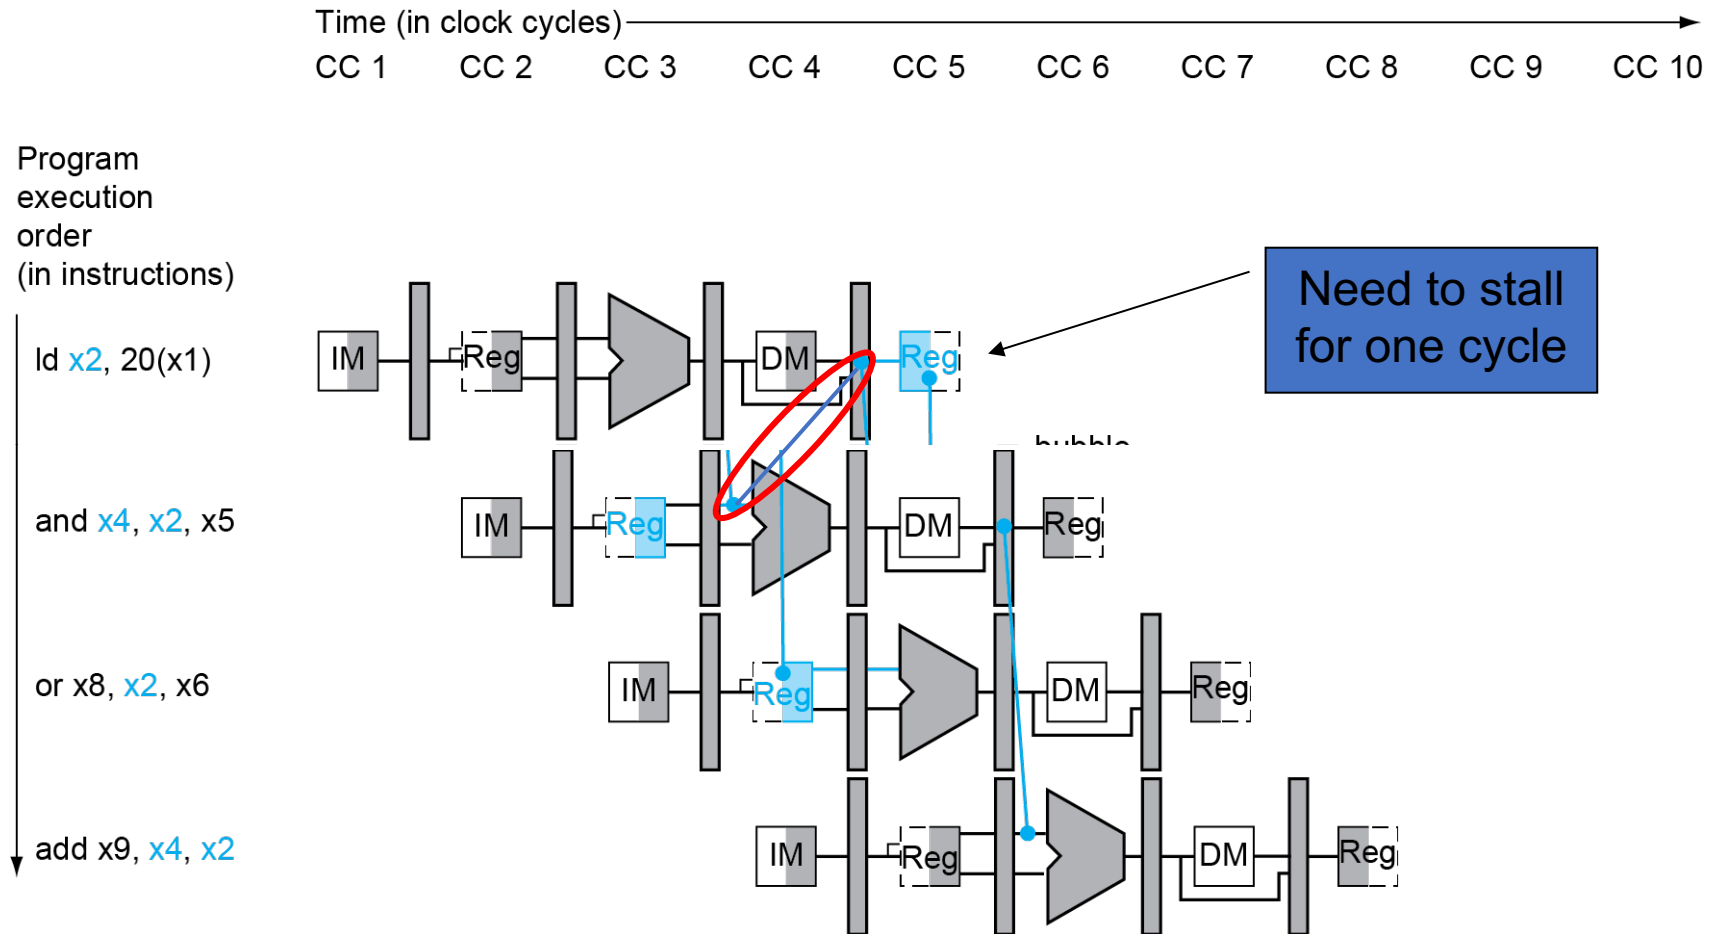
Revised Forwarding Condition

- MEM hazard
 - if (MEM/WB.RegWrite and (MEM/WB. Rd \neq 0)
and not(EX hazard)
and (MEM/WB. Rd = ID/EX. Rs1))
ForwardA = 01
 - if (MEM/WB.RegWrite and (MEM/WB. Rd \neq 0)
and not(EX hazard)
and (MEM/WB. Rd = ID/EX. Rs2))
ForwardB = 01

Datapath with Forwarding



Load-Use Data Hazard



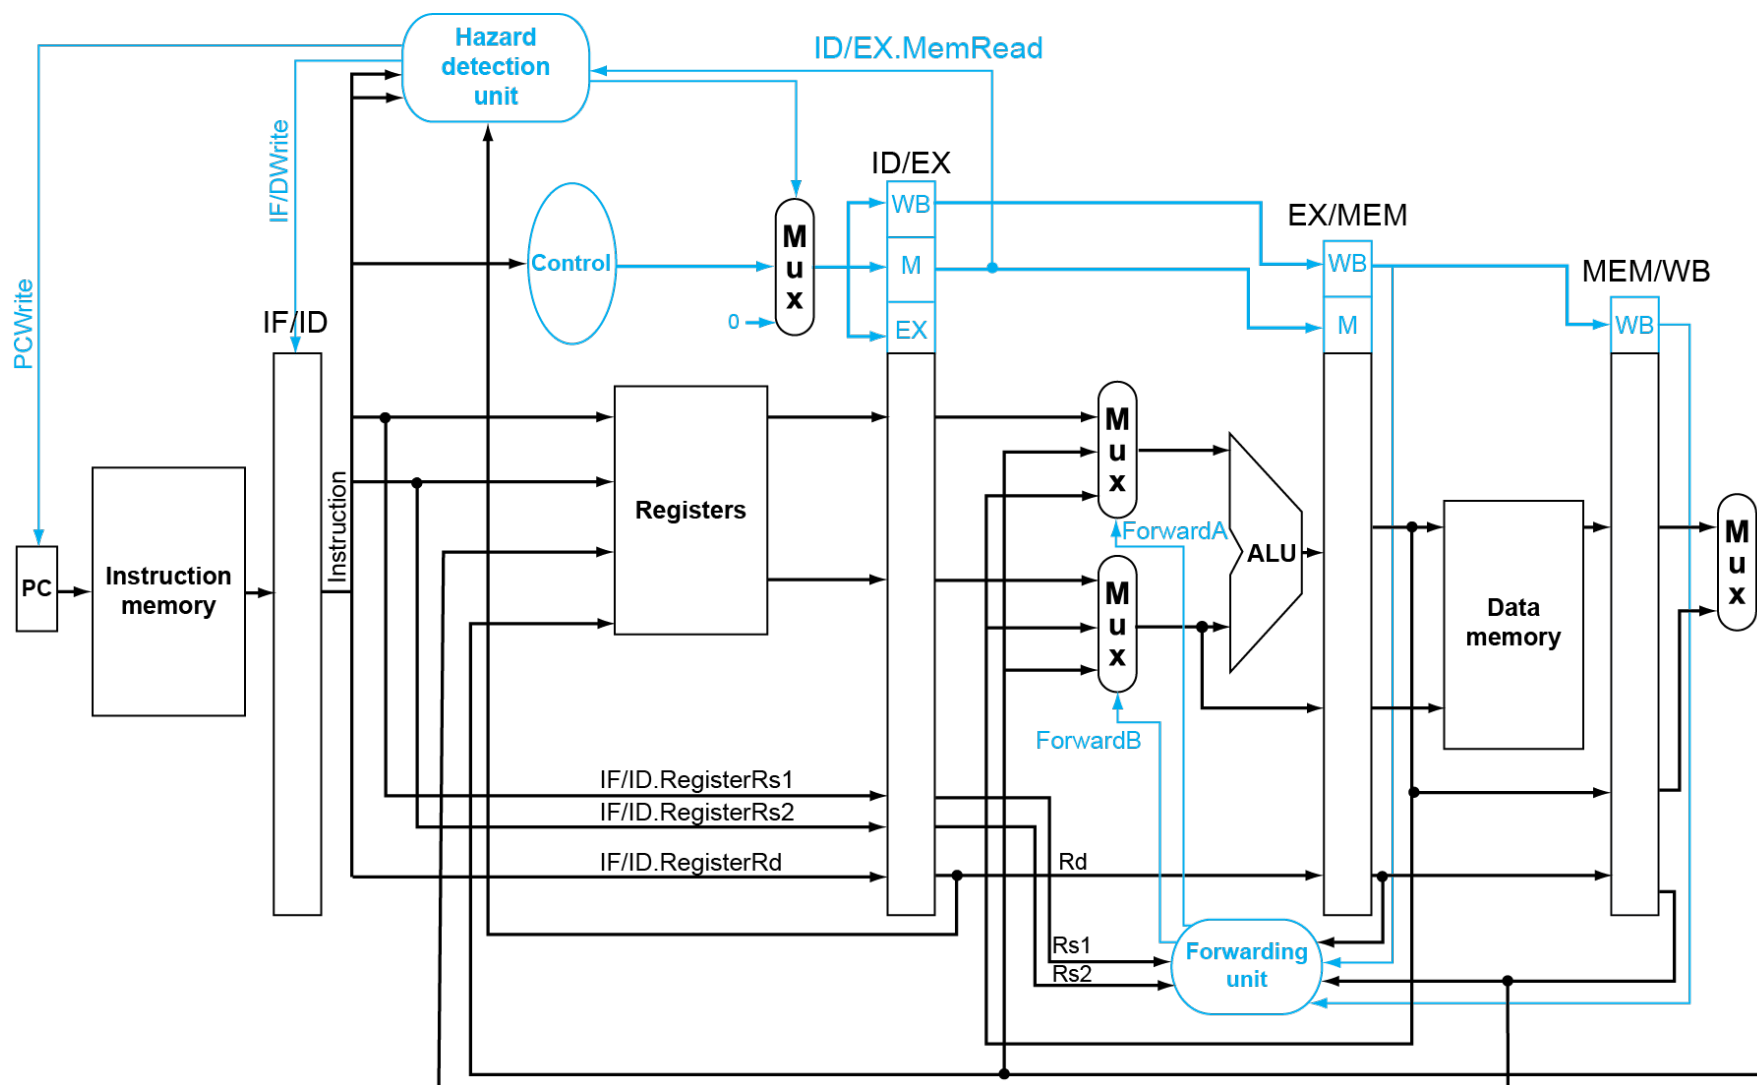
Load-Use Hazard Detection

- Check when using instruction is decoded in ID stage
- ALU operand register numbers in ID stage are given by
 - IF/ID. Rs1, IF/ID. Rs2
- Load-use hazard when
 - ID/EX.MemRead and
((ID/EX. Rd = IF/ID. Rs1) or
(ID/EX. Rd = IF/ID. Rs2))
- If detected, stall and insert bubble

How to Stall the Pipeline

- Force control values in ID/EX register to 0
 - EX, MEM and WB do nop (no-operation)
- Prevent update of PC and IF/ID register
 - Using instruction is decoded again
 - Following instruction is fetched again
 - 1-cycle stall allows MEM to read data for 1d
 - Can subsequently forward to EX stage

Datapath with Hazard Detection



Stalls and Performance

- Stalls reduce performance
 - But are required to get correct results
- Compiler can arrange code to avoid hazards and stalls
 - Requires knowledge of the pipeline structure

Summary of Data Hazards

- EX hazard & MEM hazard
 - forwarding
- Double hazard
 - Revise the forwarding condition
- Load-use hazard
 - One stall is needed, except for forwarding


Control Hazards



Control Hazard

Flow of execution depends on previous instruction

Problem: The conditional branch instruction

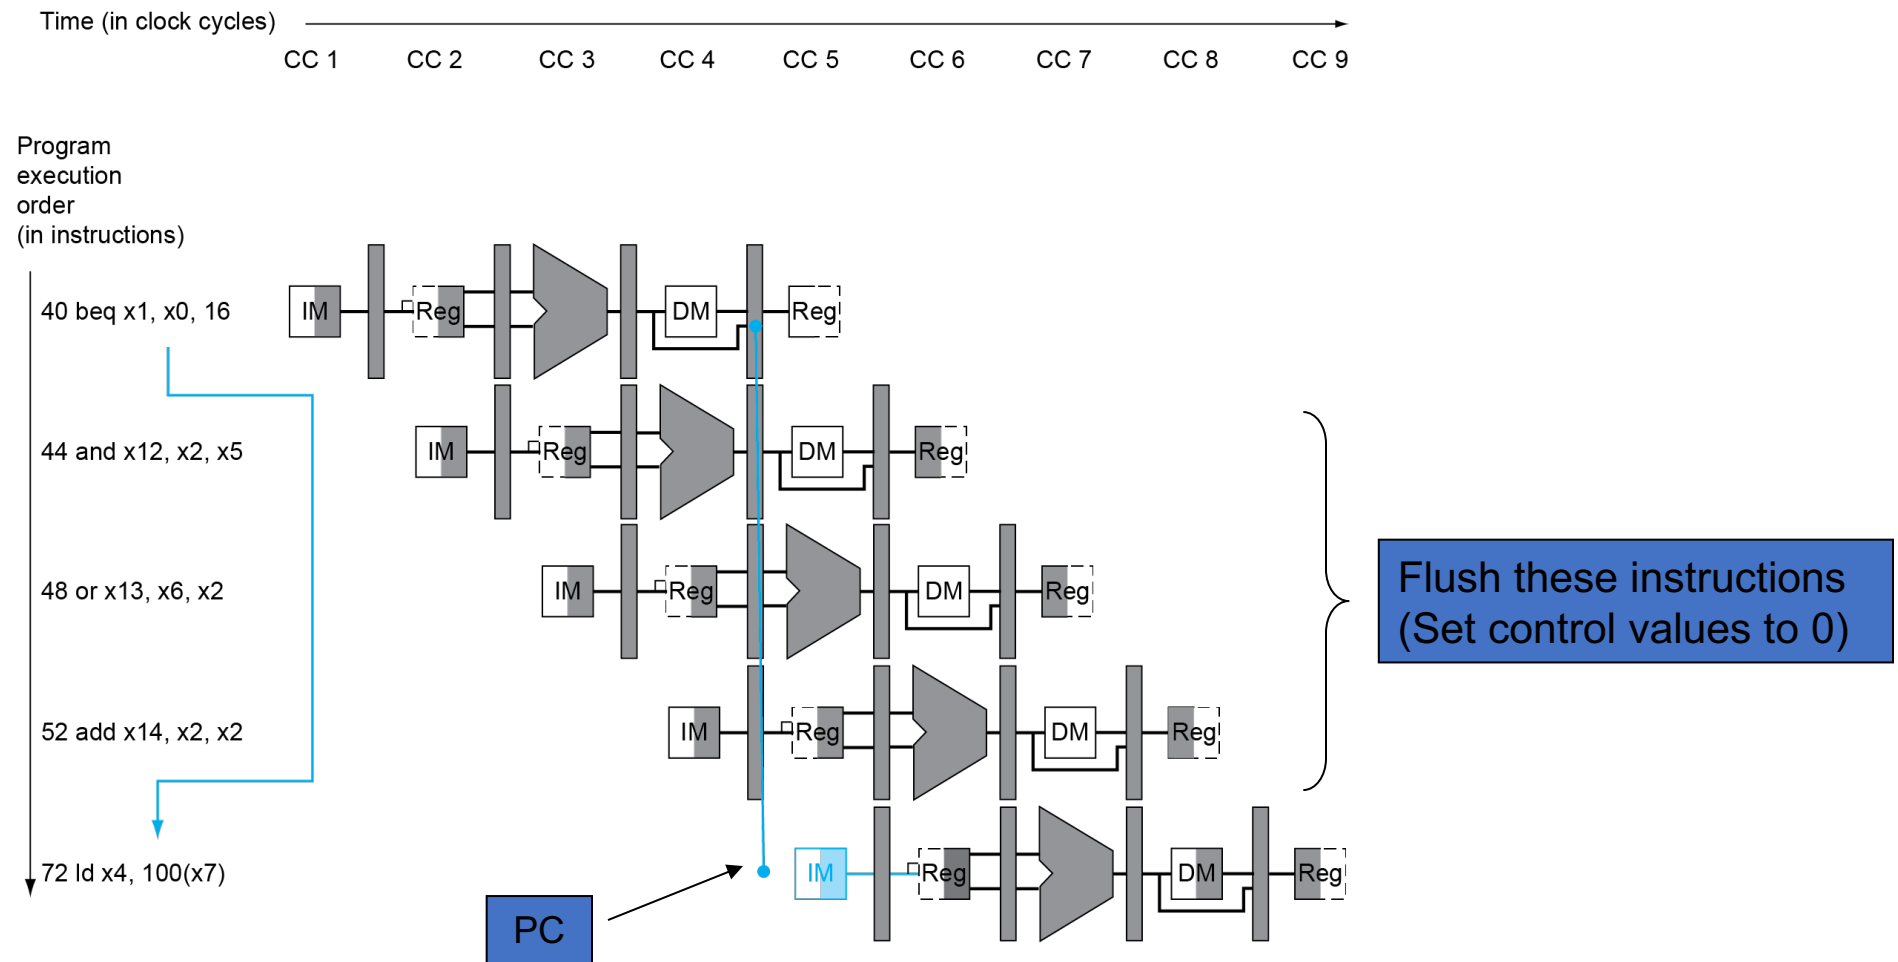


the pipeline cannot possibly know
what the next instruction should be



Branch Hazards

- If branch outcome determined in MEM



Control Hazards

- Branch determines flow of control
 - Fetching next instruction depends on branch outcome
 - Pipelining can't always fetch correct instruction
 - Still working on ID stage of branch
- Wait until branch outcome determined before fetching next instruction
 - The penalty is significant
 - How many clock cycles does the stall waste?

3 clock cycles! Almost downgrading to single-cycle CPU!




Stall on Branch

Branch taken

Branch	IF	ID	EX	MEM	WB					
Object		IF	stall	stall	IF	ID	EX	MEM	WB	
Object+1						IF	ID	EX	MEM	WB
Object+2							IF	ID	EX	MEM
Object+3								IF	ID	EX

Branch untaken

Branch	IF	ID	EX	MEM	WB					
subsequent		IF	ID	EX	MEM	WB				
subsequent+1			IF	ID	EX	MEM	WB			
subsequent+2				IF	ID	EX	MEM	WB		
subsequent+3					IF	ID	EX	MEM	WB	



How to Reduce Stall

- In RISC-V pipelining
 - Need to **compare registers** and **compute target** early in the pipelining
 - Add hardware to do it in ID stage

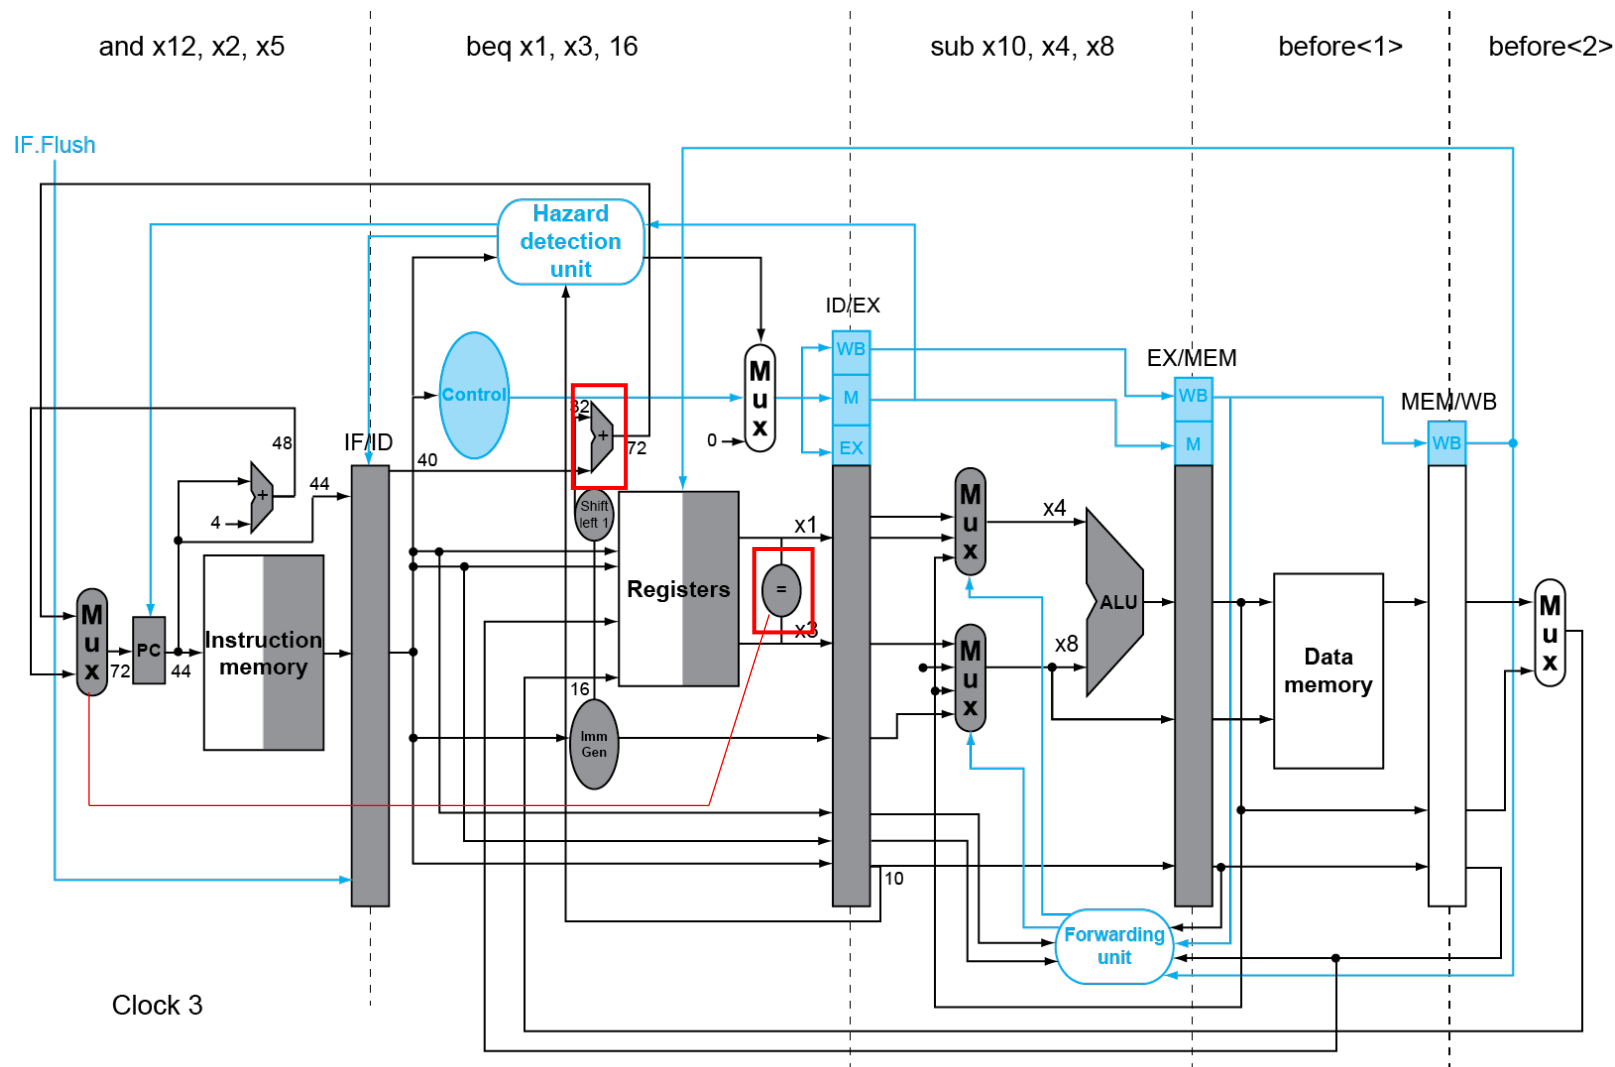
What is the hardware?

- Key processes in branch instructions
 - Compute the branch target address
 - Judge if the branch success

Which stages do they happen?

- Move hardware to determine outcome to ID stage
 - Target address adder
 - Register comparator

Forwarding Branch to Earlier Stage

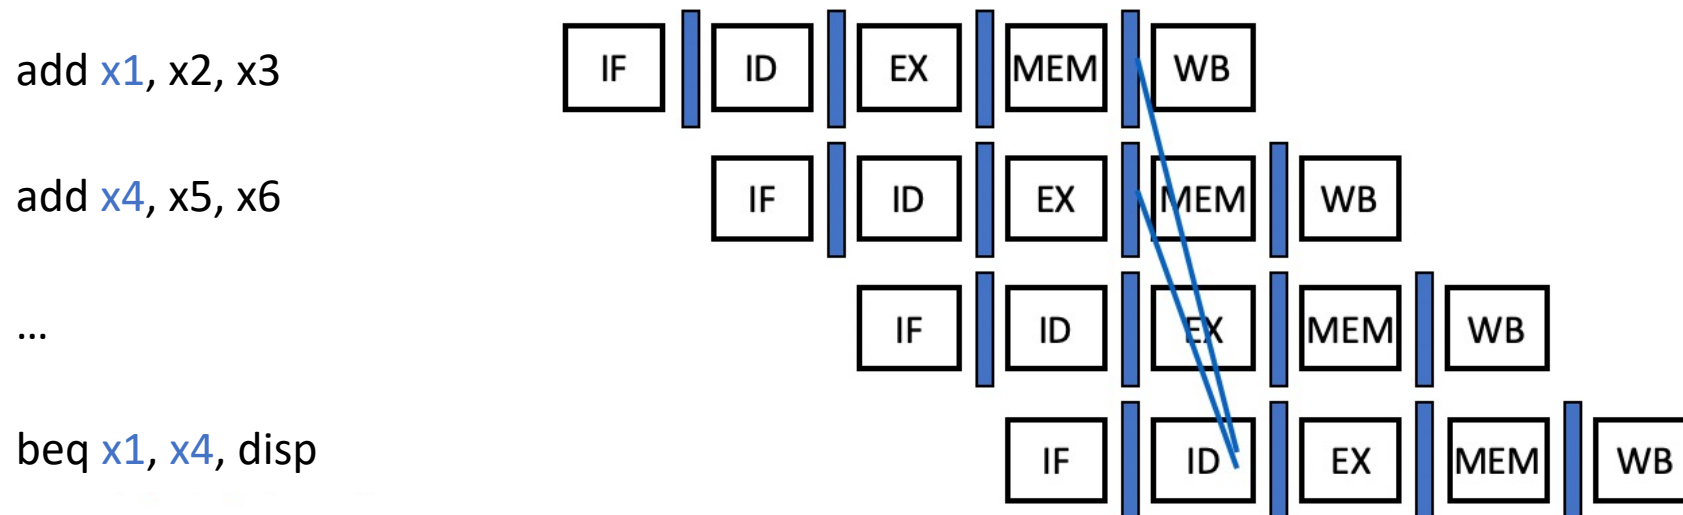


Anything wrong?



Data Hazards for Branches

- If a comparison register is a destination of 2nd or 3rd preceding ALU instruction



- Can resolve using forwarding



Data Hazards for Branches


- If a comparison register is a destination of preceding ALU instruction or 2nd preceding load instruction
 - Need 1 stall cycle

lw **x1**, addr

add **x4**, x5, x6

beq **stalled**

beq **x1**, **x4**, disp



Data Hazards for Branches


- If a comparison register is a destination of immediately preceding load instruction
 - Need 2 stall cycles

lw **x1**, addr


beq **stalled**

beq **stalled**

beq **x1**, **x0**, disp



Forwarding Branch to Earlier Stage



Stall on Branch with Optimized Solution



No!

Branch still causes a stall

Branch	IF	ID	EX	MEM	WB					
subsequent		IF	IF	ID	EX	MEM	WB			
subsequent+1				IF	ID	EX	MEM	WB		
subsequent+2					IF	ID	EX	MEM	WB	
subsequent+3						IF	ID	EX	MEM	WB



Branch Prediction

- Longer pipelines can't readily determine branch outcome early
 - Stall penalty becomes unacceptable
- Predict outcome of branch
 - Only stall if prediction is wrong
- In RISC-V pipeline
 - Can predict branches not taken
 - Fetch instruction after branch, with no delay



Reducing Branch Delay


- Predict branch taken
- Predict branch untaken
- Delay Branch



Predict branch

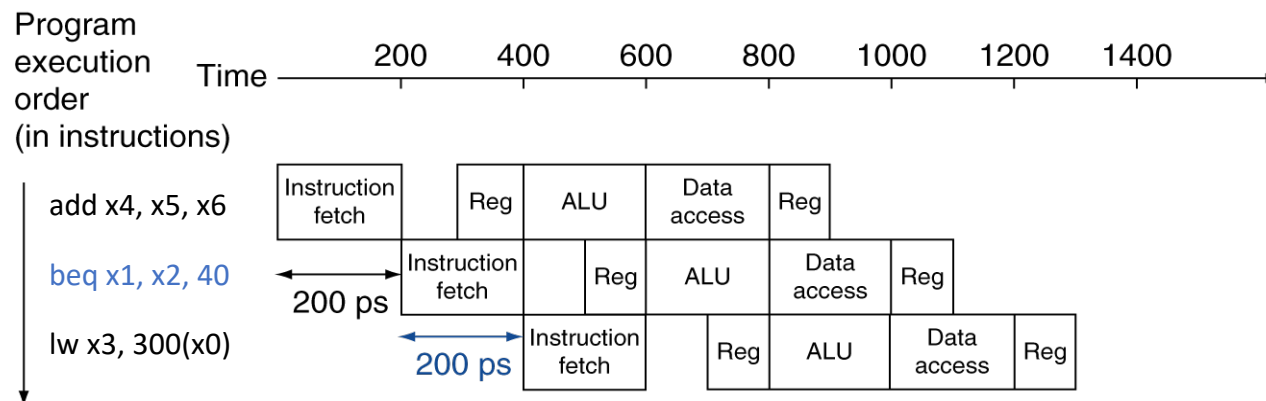
Branch i (Taken)	IF	ID	EX	MEM	WB				
i+1		IF	stall	stall	stall	stall			
Object j			IF	ID	EX	MEM	WB		
Object j+1				IF	ID	EX	MEM	WB	
Object j+2					IF	ID	EX	MEM	WB

Branch i (Untaken)	IF	ID	EX	MEM	WB				
i+1		IF	ID	EX	MEM	WB			
i+2			IF	ID	EX	MEM	WB		
i+3				IF	ID	EX	MEM	WB	
i+4					IF	ID	EX	MEM	WB

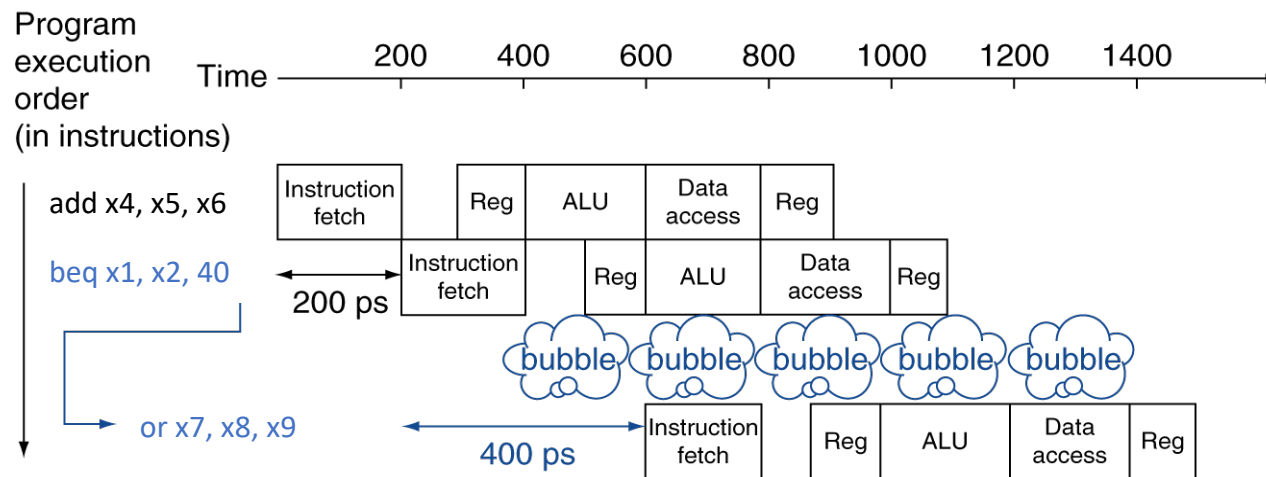


RISC-V with Predict Not Taken

Prediction
correct



Prediction
incorrect



How to Reduce Branch Delay

- Example: branch taken

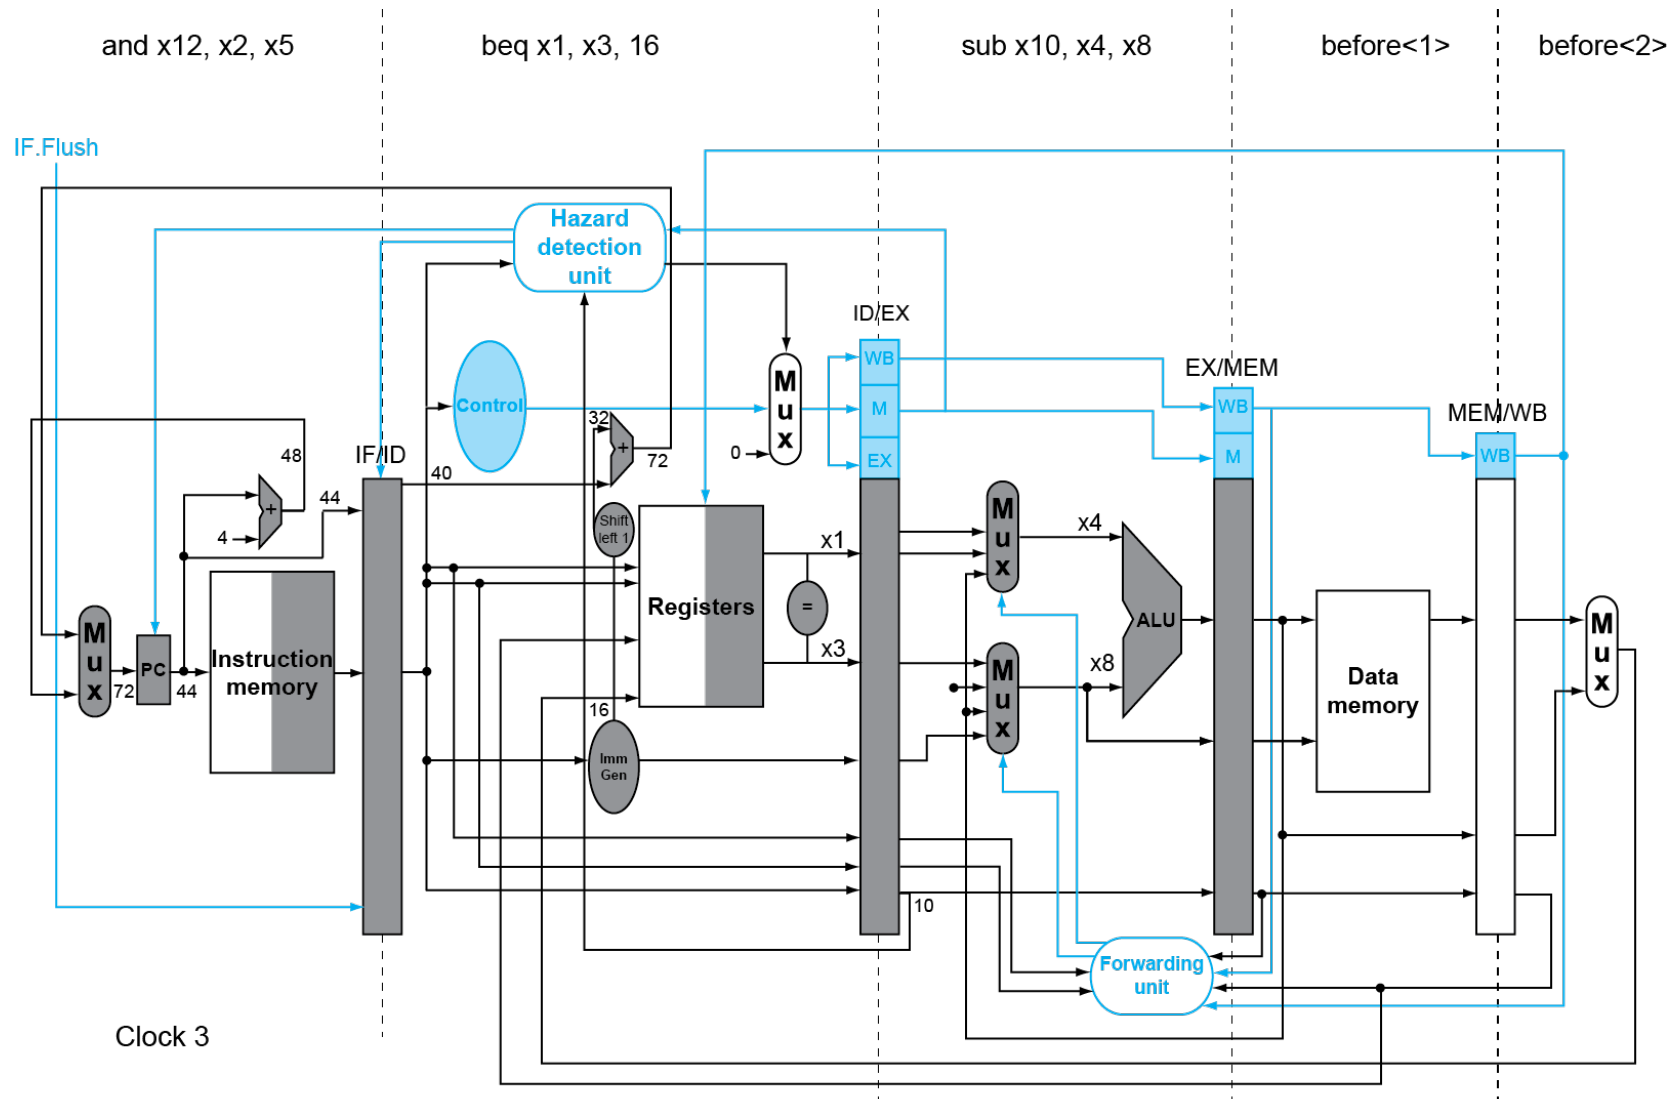
```

36:  sub   x10, x4, x8
40:  beq   x1,  x3, 16    // PC-relative branch
                          // to 40+16*2=72
44:  and   x12, x2, x5
48:  or    x13, x2, x6
52:  add   x14, x4, x2
56:  sub   x15, x6, x7

72:  ld    x4, 50(x7)

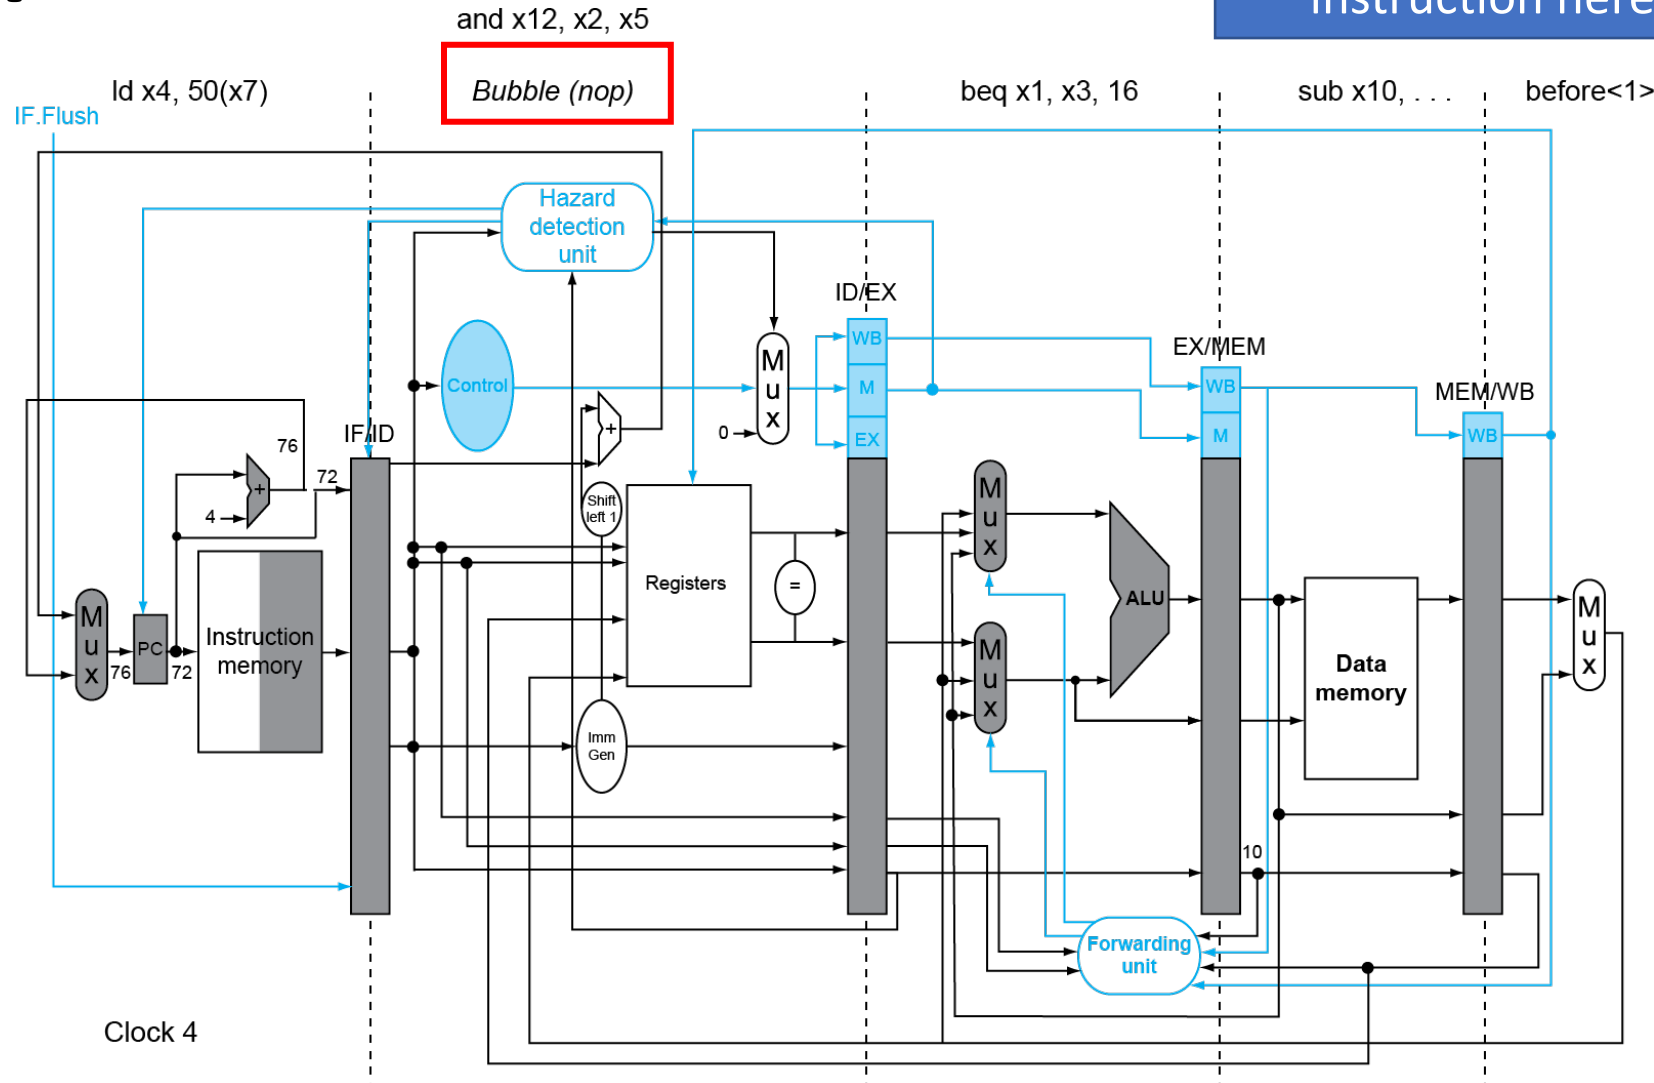
```


Example: Branch Taken



Example: Branch Taken

What is the original instruction here?



Reducing Branch Delay


- Predict branch success
- Predict branch failure
- Delay Branch



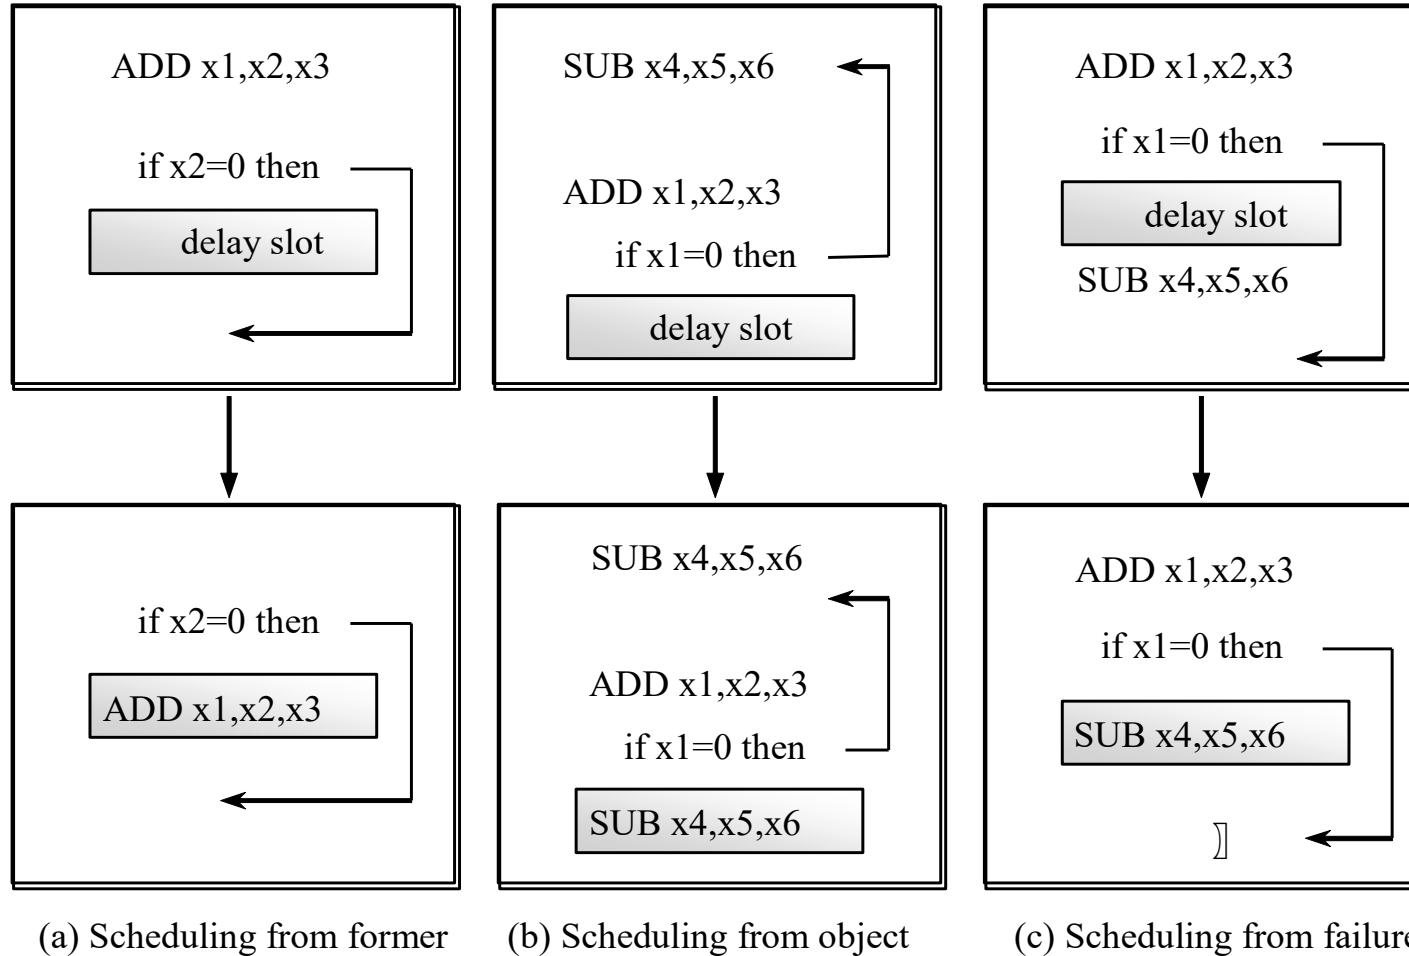
Pipelining with a branch delay slot

Branch	Branch i	IF	ID	EX	MEM	WB				
	Delay slot i+1		IF	ID	EX	MEM	WB			
Failure	i+2			IF	ID	EX	MEM	WB		
	i+3				IF	ID	EX	MEM	WB	
	i+4					IF	ID	EX	MEM	WB

Branch	Branch i	IF	ID	EX	MEM	WB				
	Delay slot i+1		IF	ID	EX	MEM	WB			
Success	Object j			IF	ID	EX	MEM	WB		
	Object j+1				IF	ID	EX	MEM	WB	
	Object j+2					IF	ID	EX	MEM	WB




Code Scheduling



Code Scheduling

Branch	Branch i	IF	ID	EX	MEM	WB				
	Delay slot i+1		IF	idle	idle	idle	idle			
Failure	i+2			IF	ID	EX	MEM	WB		
	i+3				IF	ID	EX	MEM	WB	
	i+4					IF	ID	EX	MEM	WB

Branch	Branch i	IF	ID	EX	MEM	WB				
	Delay slot i+1		IF	ID	EX	MEM	WB			
Success	Object j			IF	ID	EX	MEM	WB		
	Object j+1				IF	ID	EX	MEM	WB	
	Object j+2					IF	ID	EX	MEM	WB



Question: Is delay slot a really good design?

Review

- “A **RISC-V ISA** is defined as a base integer ISA, which must be present in any implementation, plus optional extensions to the base ISA.
- The base integer ISAs are very similar to that of the early RISC processors except **with no branch delay slots** and with support for optional variable-length instruction encodings. ”

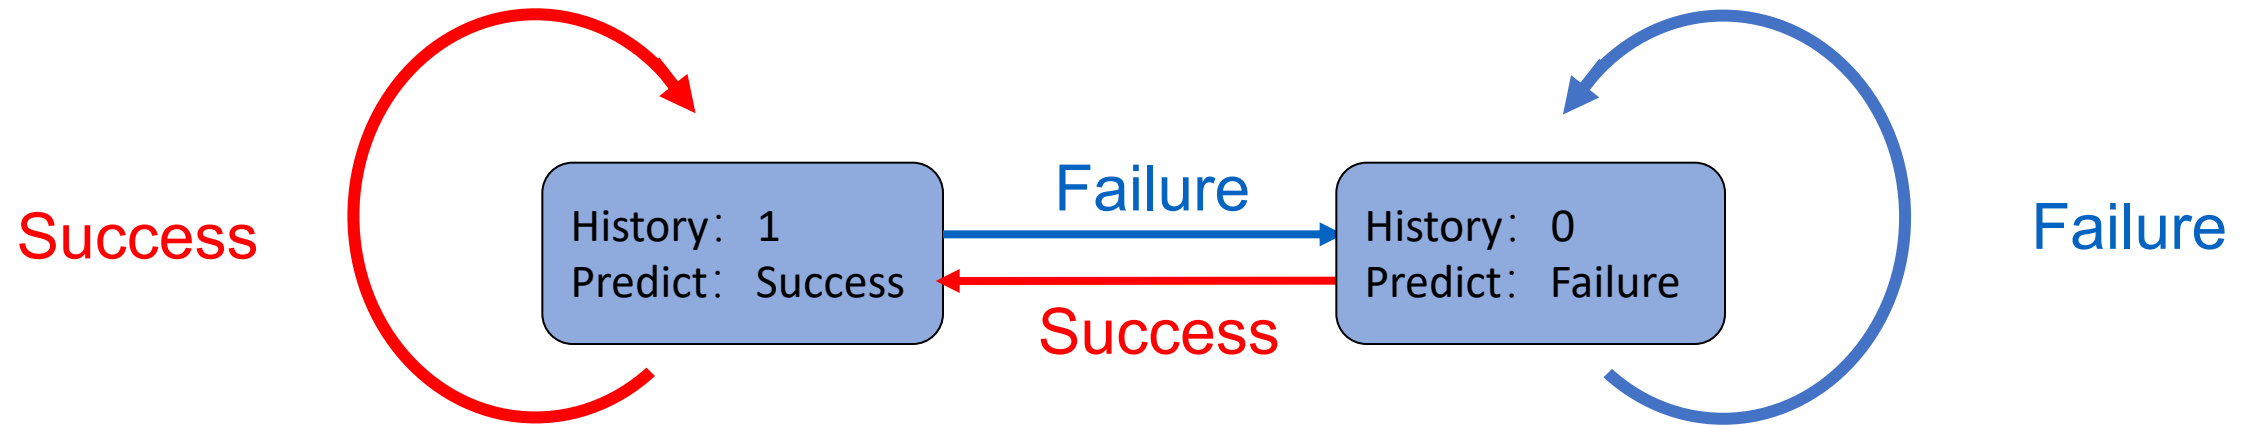
——The RISC-V Instruction Set Manual Volume I



Dynamic Branch Prediction


- In deeper and superscalar pipelines, branch penalty is more significant
- Use dynamic prediction
 - Branch prediction buffer (aka branch history table)
 - Indexed by recent branch instruction addresses
 - Stores outcome (taken/not taken)
 - To execute a branch
 - Check table, expect the same outcome
 - Start fetching from fall-through or target
 - If wrong, flush pipeline and flip prediction

Branch History Table (BHT)



1-Bit Predictor: Shortcoming


- Inner loop branches mispredicted twice!



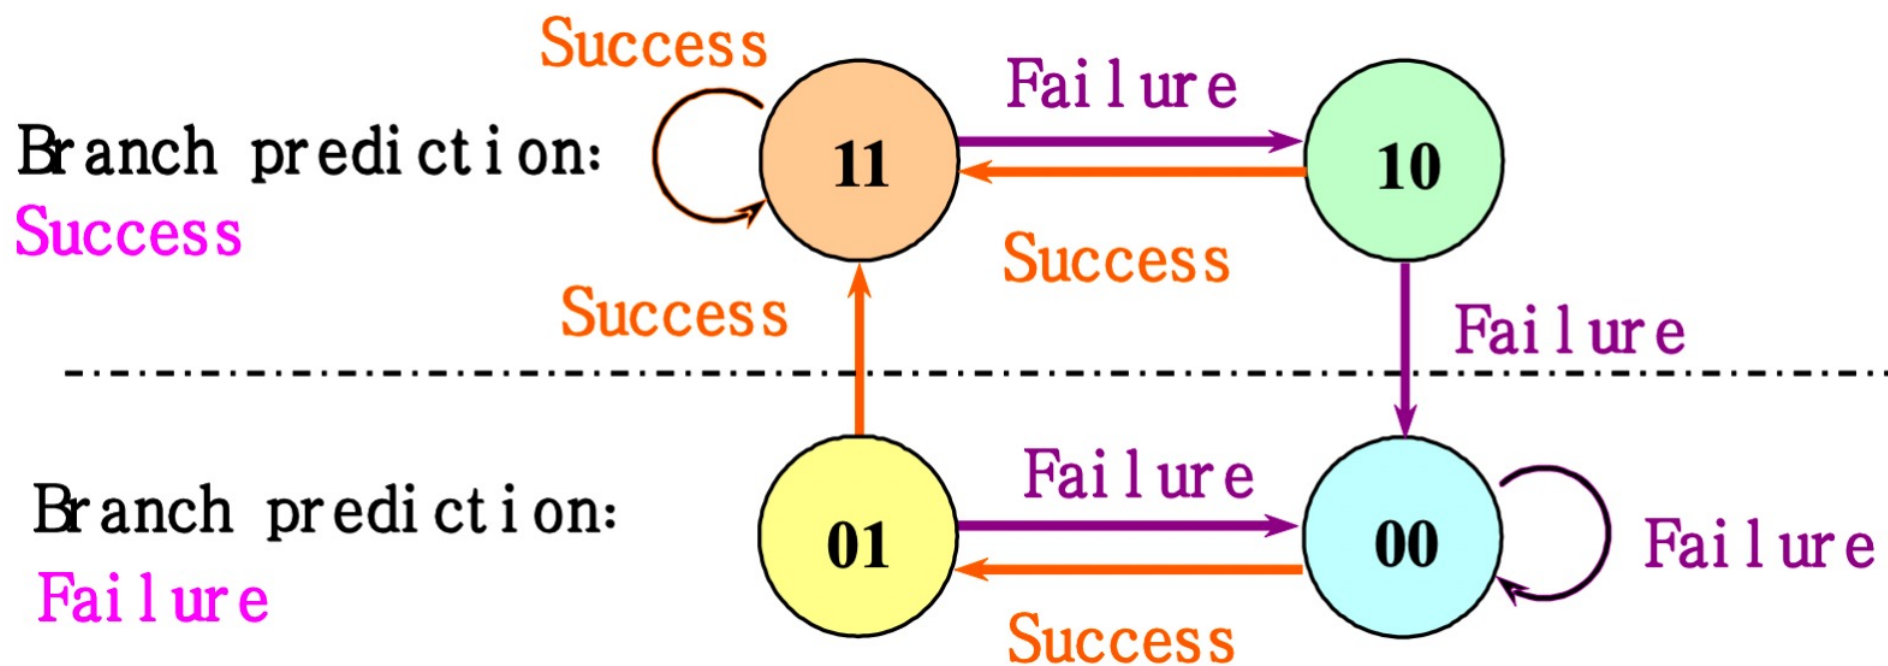
- Mispredict as taken on last iteration of inner loop
- Then mispredict as not taken on first iteration of inner loop next time around

2-Bit Predictor

- Only change prediction on two successive mispredictions




Branch History Table



2-Bit Predictor: Example

- Inner loop branches mispredicted only once!



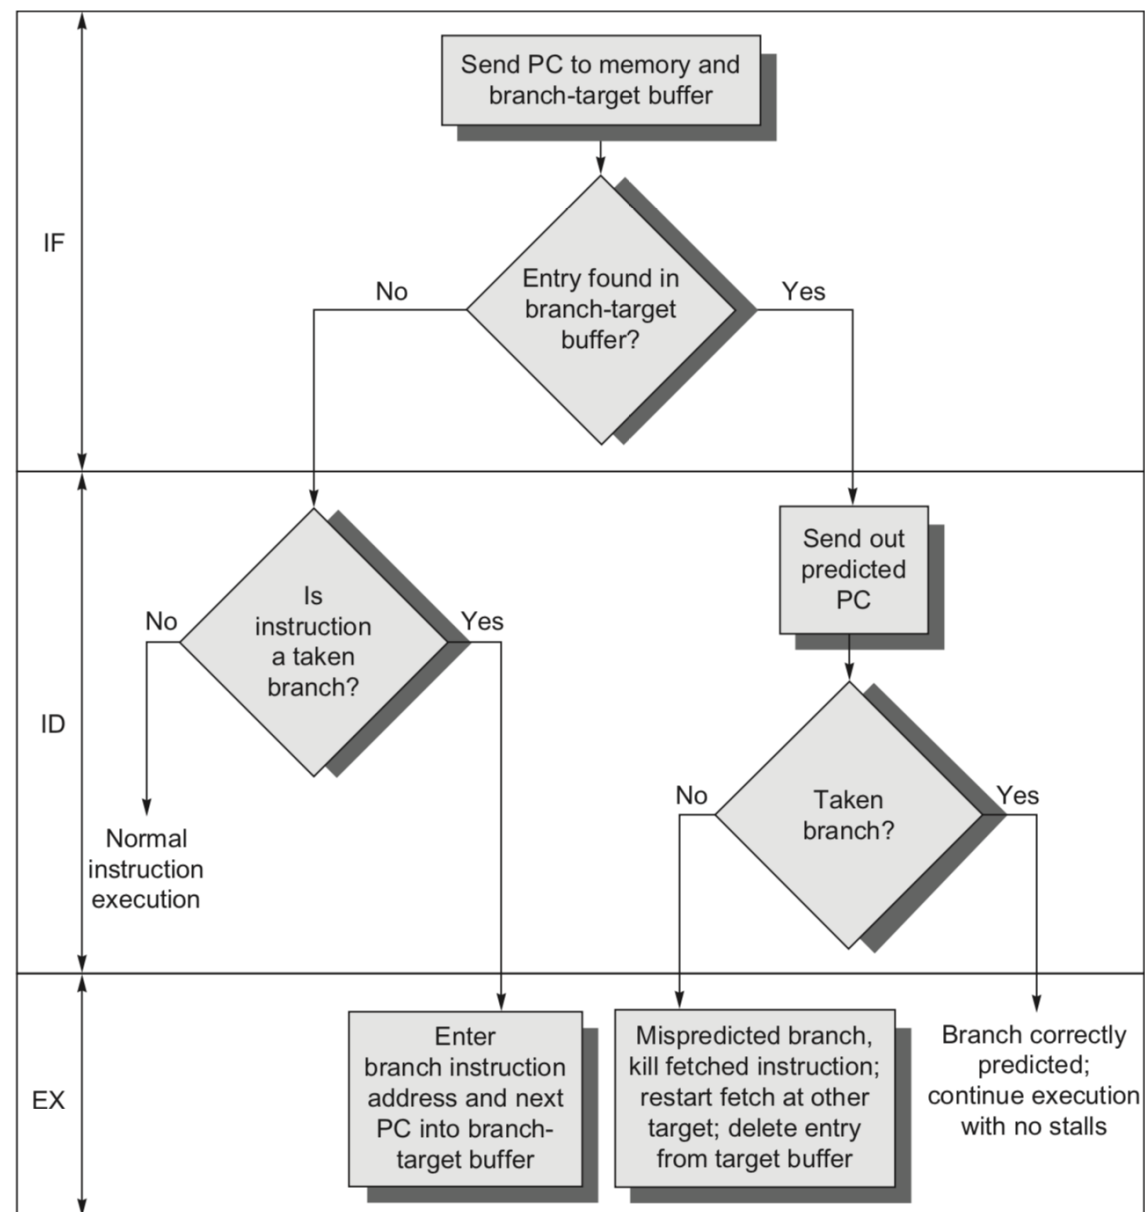

- Only mispredict as taken on last iteration of inner loop

Advanced Techniques for Instruction Delivery and Speculation

- Increasing Instruction Fetch Bandwidth
 - Branch-Target Buffers
- Specialized Branch Predictors: Predicting Procedure Returns, Indirect Jumps, and Loop Branches
 - Integrated Instruction Fetch Units

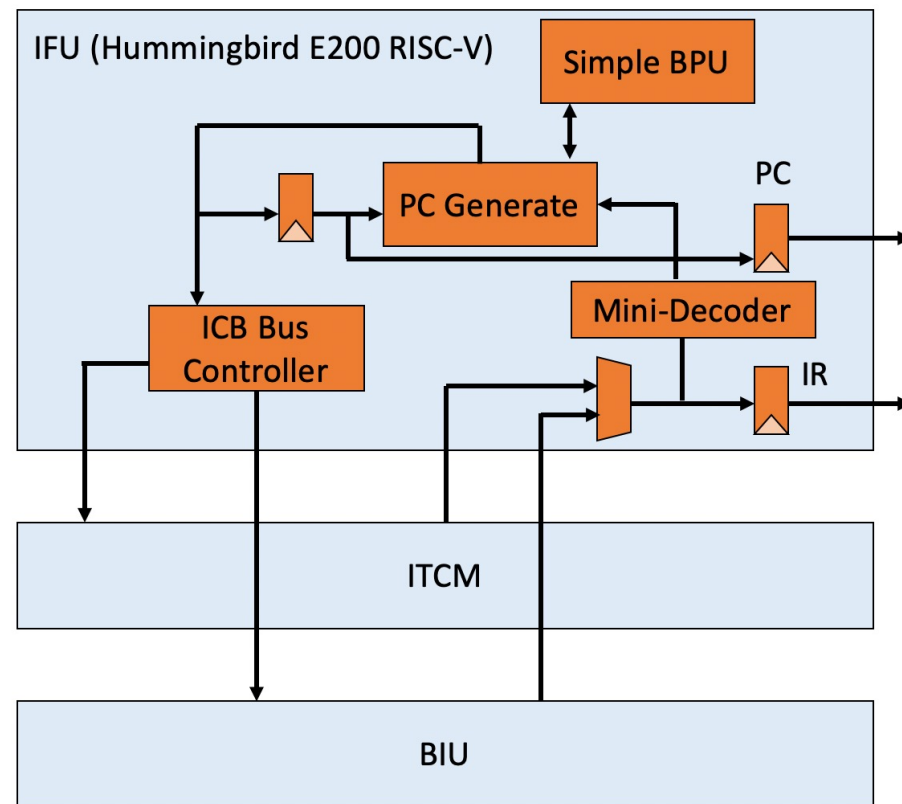


Branch-Target Buffers



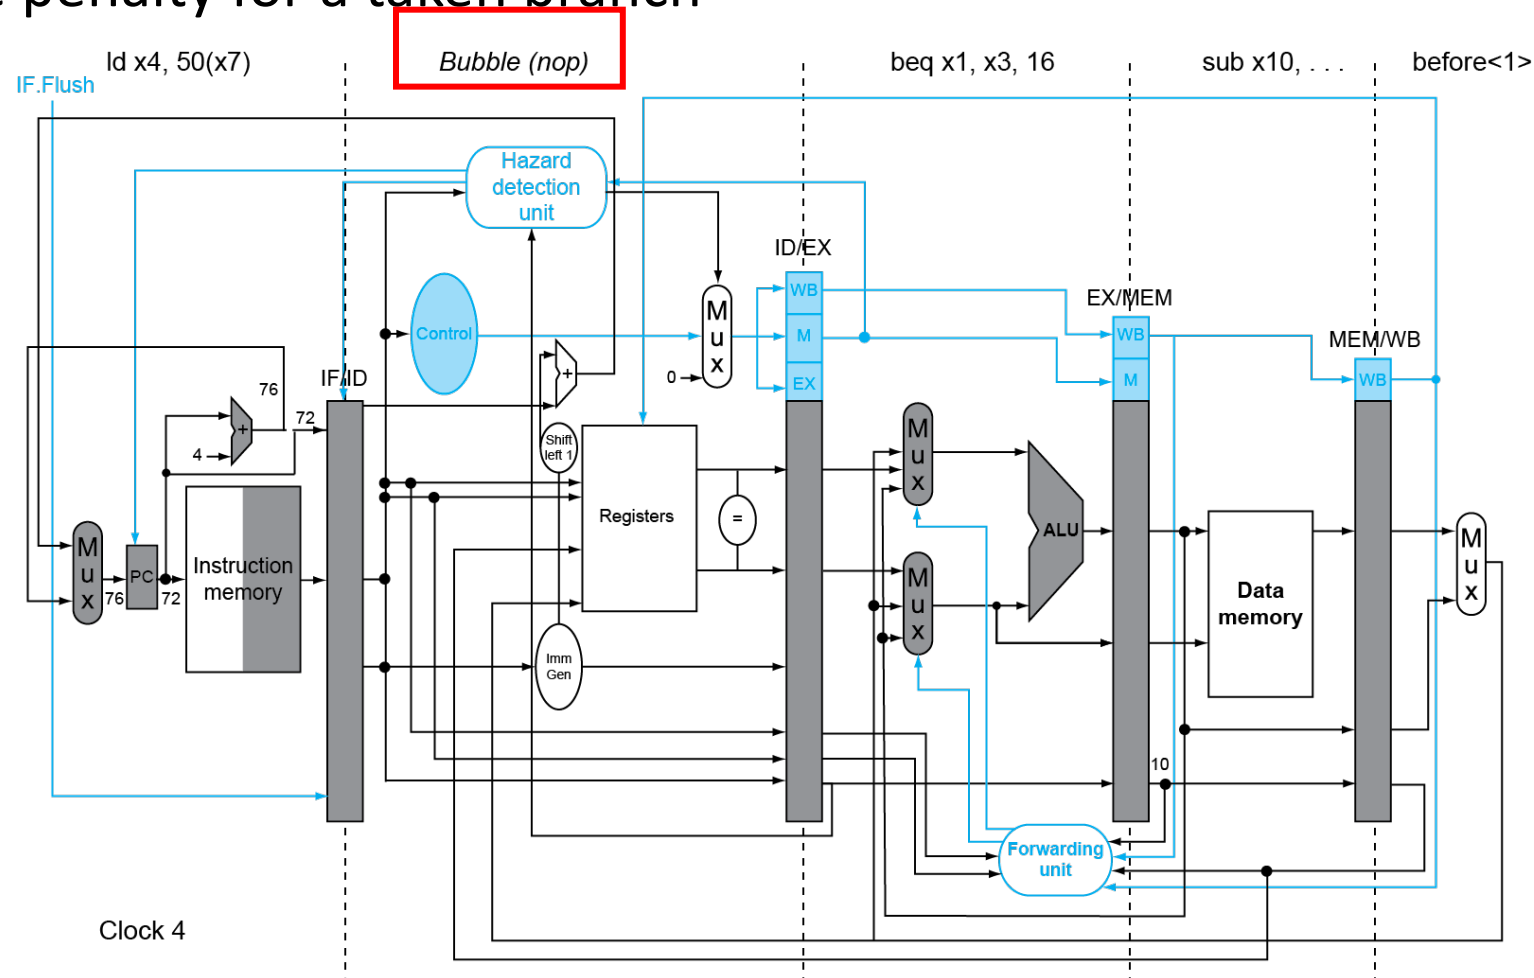
Integrated Instruction Fetch Units

- An integrated instruction fetch unit that integrates several functions:
 - Integrated branch prediction
 - Instruction prefetch
 - Instruction memory access and buffering
- Instruction fetch as a simple single pipe stage given the complexities of multiple issue is no longer valid



Calculating the Branch Target

- Even with predictor, still need to calculate the target address
 - 1-cycle penalty for a taken branch



Calculating the Branch Target

- Even with predictor, still need to calculate the target address
 - 1-cycle penalty for a taken branch
- **Branch target buffer**
 - Cache of target addresses
 - Indexed by PC when instruction fetched
 - If hit and instruction is branch predicted taken, can fetch target immediately

Branch-Target Buffer/Branch-Target Cache

Benefit

- Get instructions at branch target faster
- It can provide multiple instructions at the branch target once, which is necessary for the multi processor
- branch folding
 - It is possible to achieve unconditional branching without delay, or sometimes conditional branching without delay

